

OPerating Systems Laboratory

OPSLab

PART II

Bash Command Line

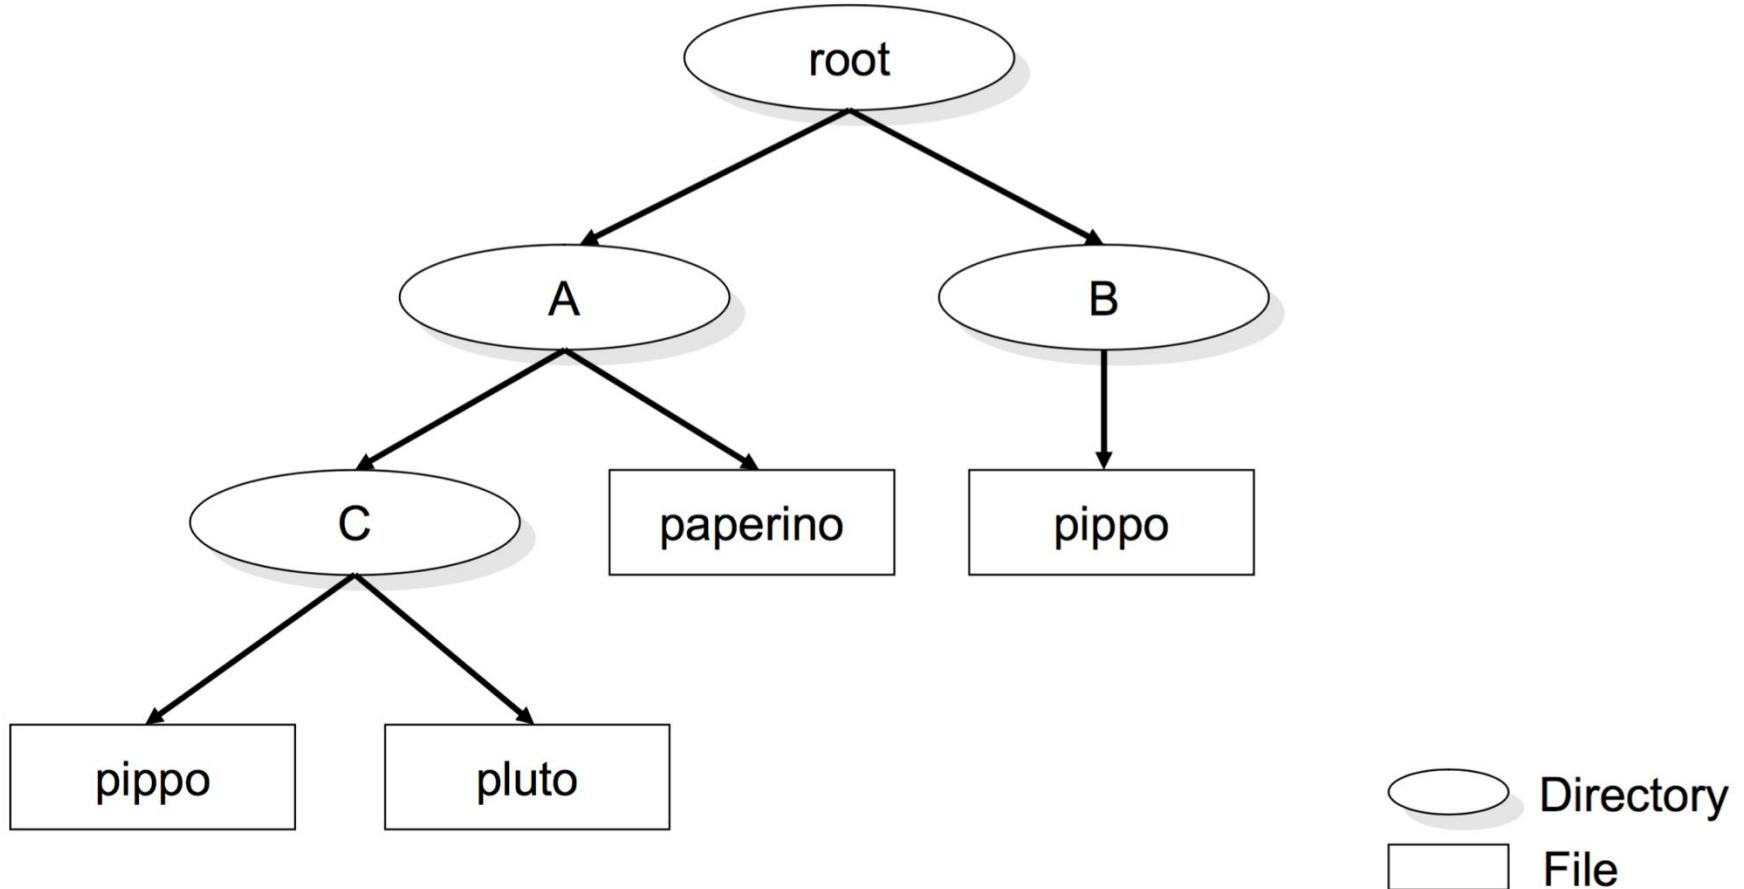
Prof. Marco Autili

University of L'Aquila

marco.autili@univaq.it

<http://people.disim.univaq.it/marco.autili/>

File System

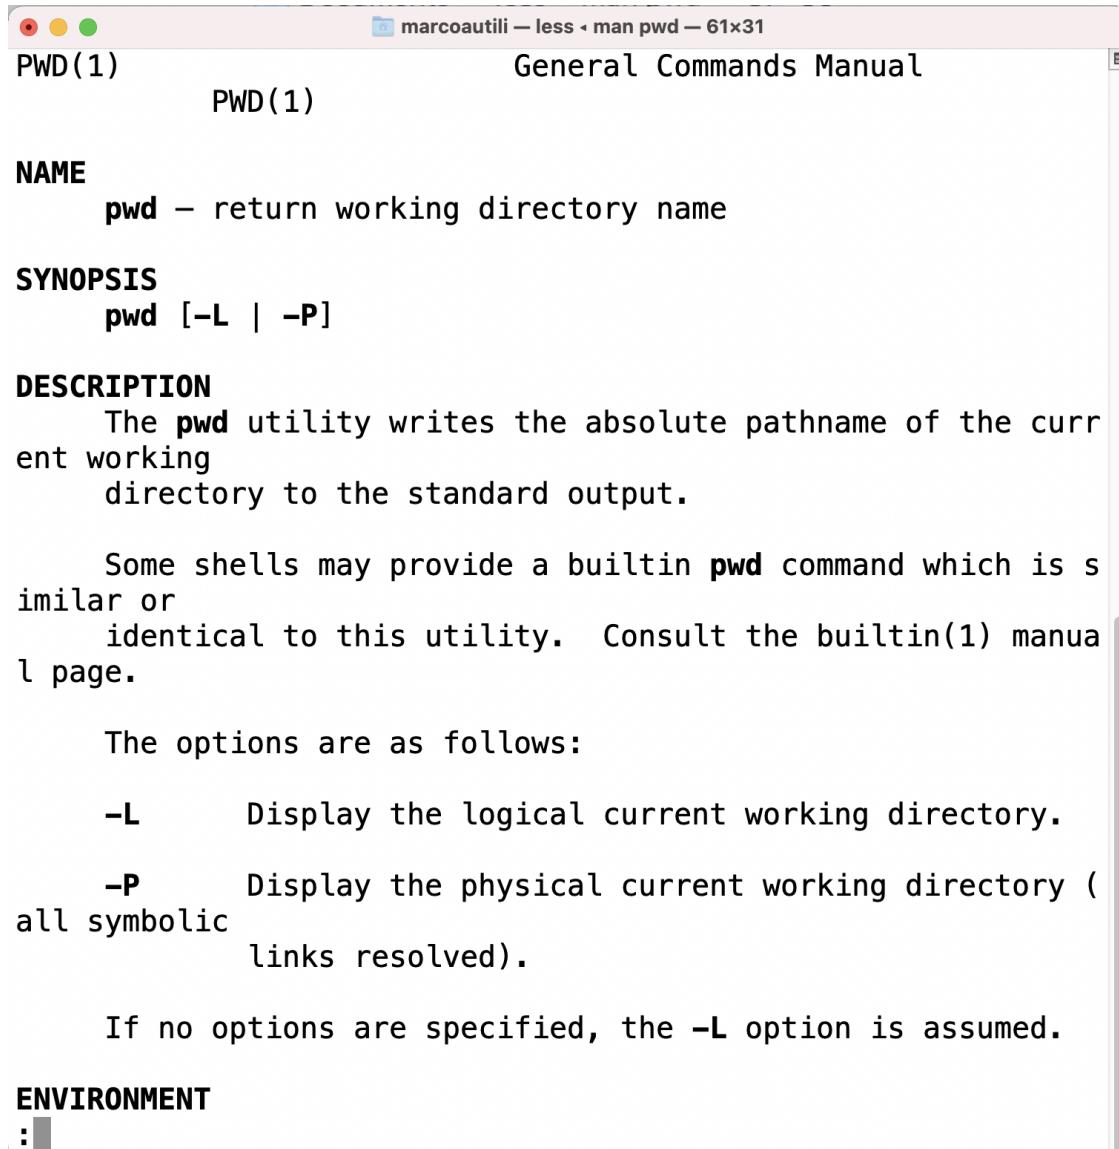


Manual Pages

- `man <command name>`

- `man pwd`
- `man 1 pwd`
- `man s1 pwd`

- press `q` to exit the manual pages



The screenshot shows a terminal window with the title "marcoautili — less - man pwd — 61x31". The window displays the man page for the `pwd` command. The page is titled "PWD(1)" and belongs to the "General Commands Manual".

NAME
`pwd` — return working directory name

SYNOPSIS
`pwd [-L | -P]`

DESCRIPTION
The `pwd` utility writes the absolute pathname of the current working directory to the standard output.

Some shells may provide a builtin `pwd` command which is similar or identical to this utility. Consult the `builtin(1)` manual page.

The options are as follows:

- L** Display the logical current working directory.
- P** Display the physical current working directory (all symbolic links resolved).

If no options are specified, the `-L` option is assumed.

ENVIRONMENT
:

Searching Manual Pages

Search within all the manual pages

- man -k <search term>

Search within specific manual pages

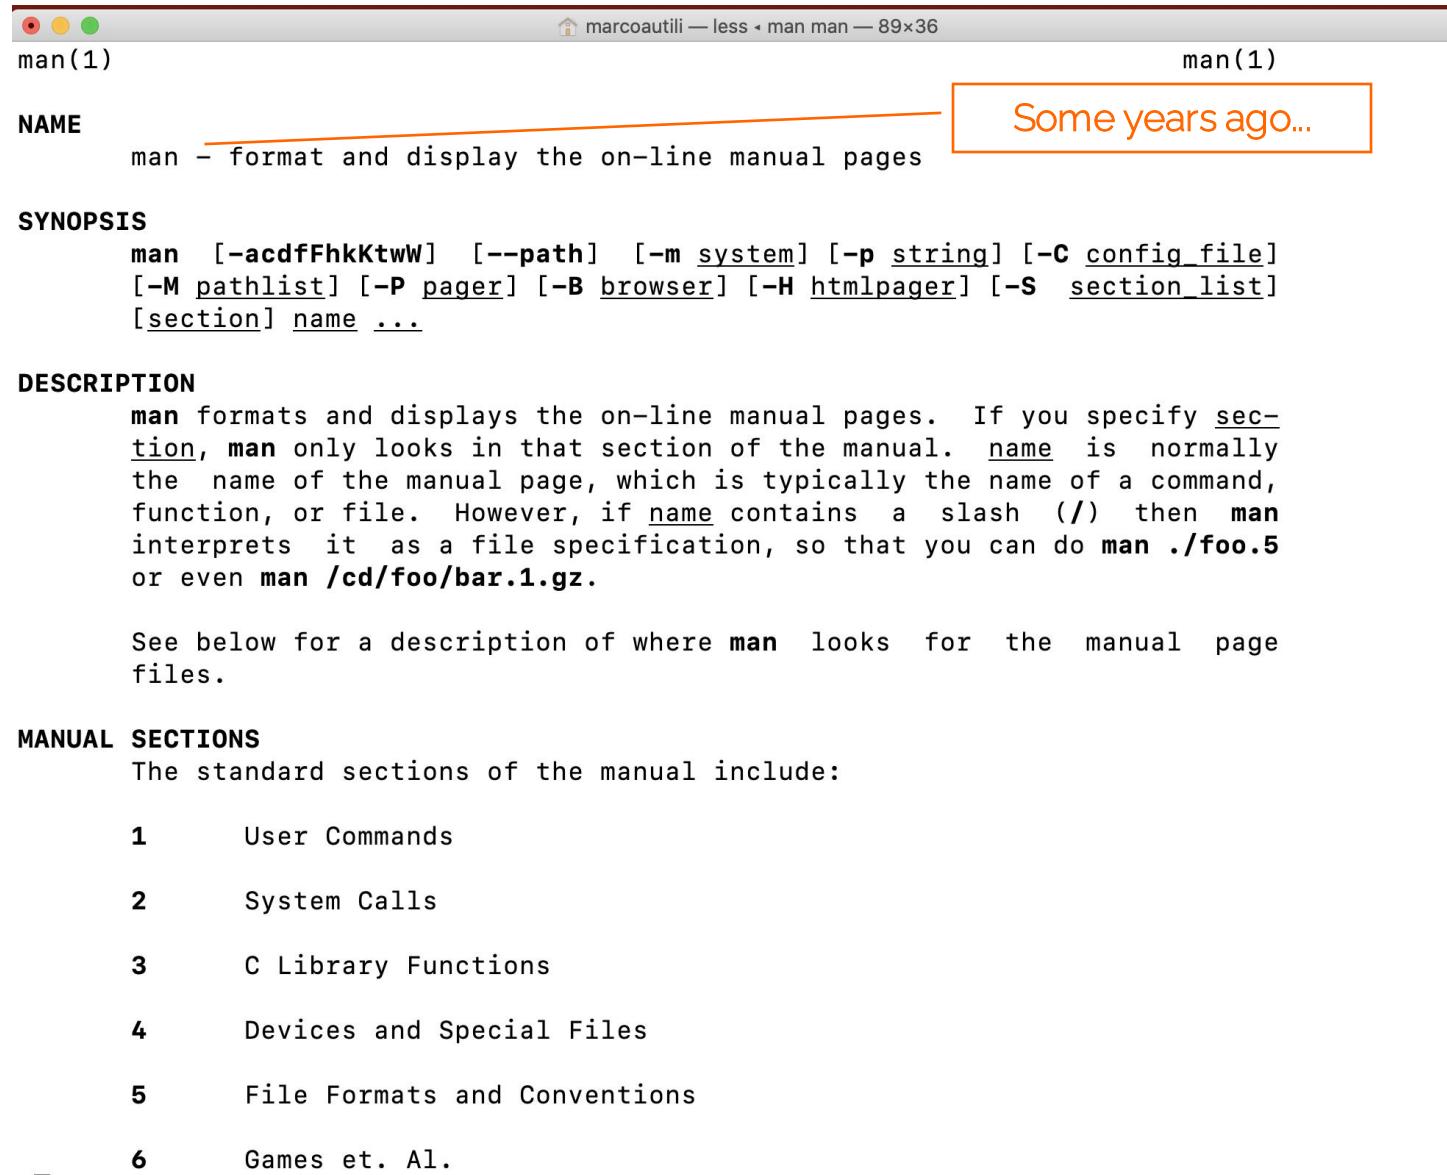
- while on the page, type forward slash '/' followed by the **term** you want to search for, then hit 'enter'
- press 'n' to go to the next occurrence

```
Marcos-MacBook-Pro:Documents marcoautili$  
Marcos-MacBook-Pro:Documents marcoautili$  
Marcos-MacBook-Pro:Documents marcoautili$ man -k pluto  
pluto: nothing appropriate  
Marcos-MacBook-Pro:Documents marcoautili$  
Marcos-MacBook-Pro:Documents marcoautili$  
Marcos-MacBook-Pro:Documents marcoautili$ man -k order  
bindtags(ntcl)           - Determine which bindings apply to a window, and order of evaluation  
lastcomm(1)              - show last commands executed in reverse order  
lower(ntcl)              - Change a window's position in the stacking order  
lreverse(ntcl)           - Reverse the order of a list  
operator(7)              - C operator precedence and order of evaluation  
raise(ntcl)              - Change a window's position in the stacking order  
Marcos-MacBook-Pro:Documents marcoautili$
```

Want to know how man works?

Simply type

- `man man`



The screenshot shows a terminal window with the title bar "marcoautili — less - man man — 89x36". The window contains the man(1) manual page. A red arrow points from the word "man" in the "NAME" section to the text "Some years ago...".

```
man(1)                                         man(1)

NAME
    man - format and display the on-line manual pages

SYNOPSIS
    man [-acdfFhkKtwW] [--path] [-m system] [-p string] [-C config_file]
        [-M pathlist] [-P pager] [-B browser] [-H htmlpager] [-S section_list]
        [section] name ...

DESCRIPTION
    man formats and displays the on-line manual pages. If you specify section, man only looks in that section of the manual. name is normally the name of the manual page, which is typically the name of a command, function, or file. However, if name contains a slash (/) then man interprets it as a file specification, so that you can do man ./foo.5 or even man /cd/foo/bar.1.gz.

    See below for a description of where man looks for the manual page files.

MANUAL SECTIONS
    The standard sections of the manual include:

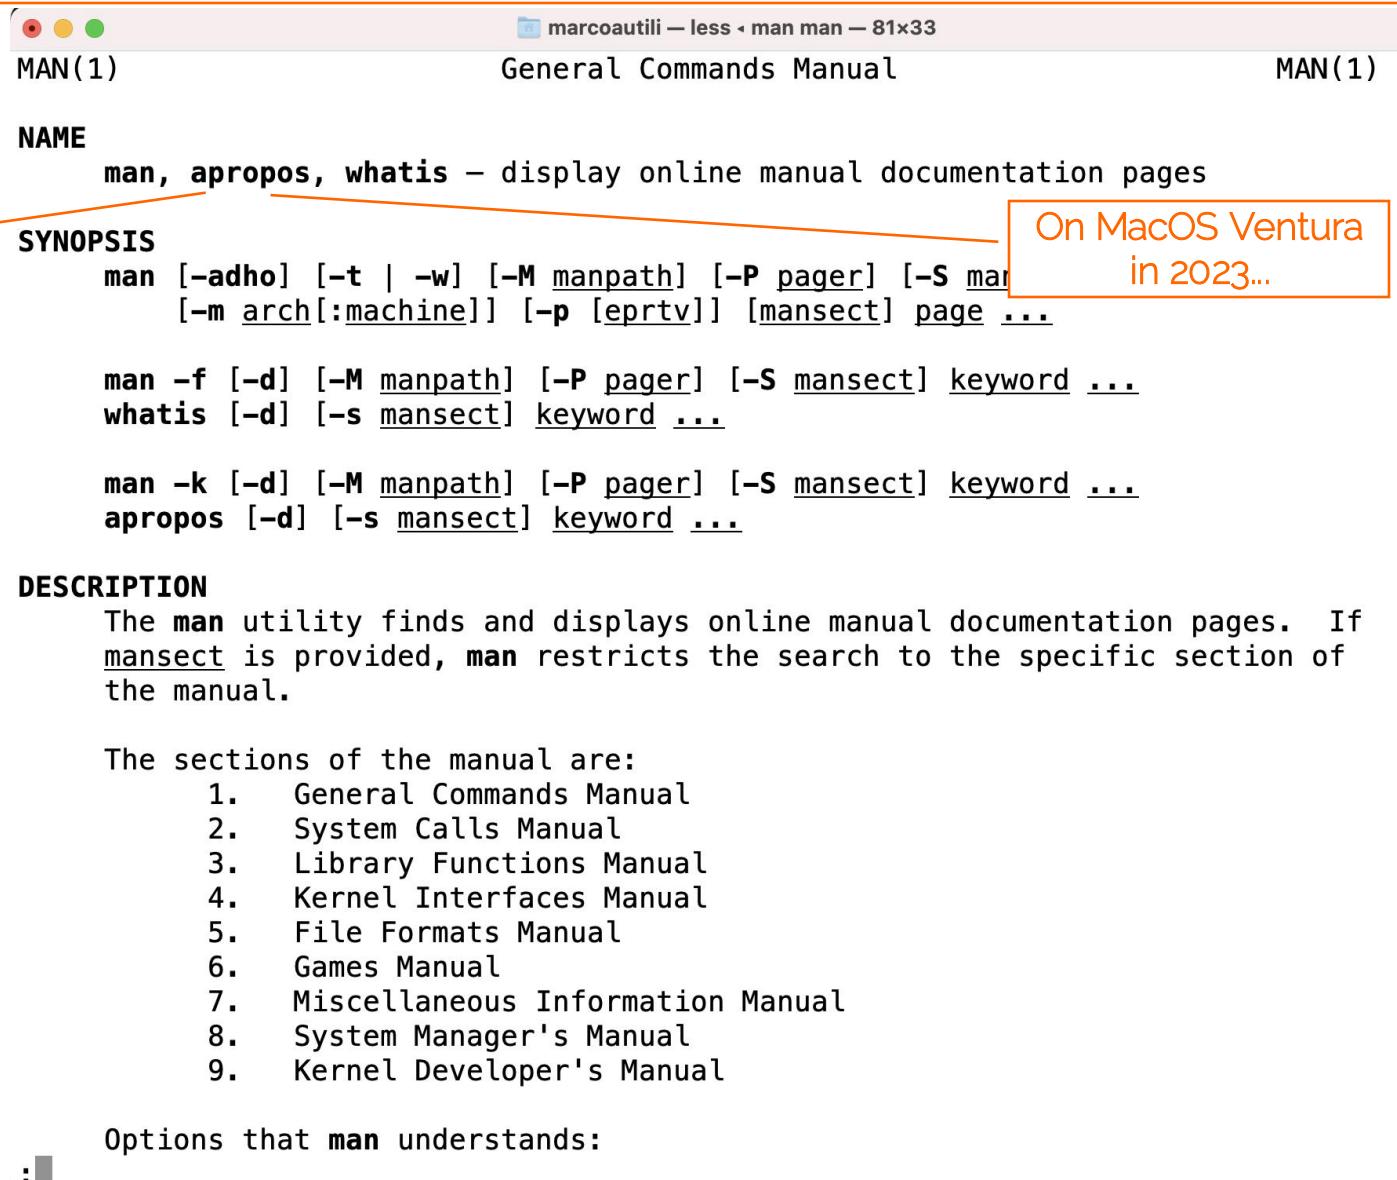
        1      User Commands
        2      System Calls
        3      C Library Functions
        4      Devices and Special Files
        5      File Formats and Conventions
        6      Games et. Al.
```

Want to know how man works?

Simply type

- `man man`

Often a wrapper for
the
`man -k` command



The screenshot shows a terminal window titled "marcoautili — less < man man — 81x33". The window displays the "MAN(1)" manual page for the "man" command. The page is divided into sections: NAME, SYNOPSIS, DESCRIPTION, and SEE ALSO. The SYNOPSIS section contains several command-line options and their descriptions. A callout box labeled "On MacOS Ventura in 2023..." points to the SYNOPSIS section.

NAME
`man, apropos, whatis` – display online manual documentation pages

SYNOPSIS

```
man [-adho] [-t | -w] [-M manpath] [-P pager] [-S manpage]
     [-m arch[:machine]] [-p [eprtv]] [mansect] page ...
man -f [-d] [-M manpath] [-P pager] [-S mansect] keyword ...
whatis [-d] [-s mansect] keyword ...
man -k [-d] [-M manpath] [-P pager] [-S mansect] keyword ...
apropos [-d] [-s mansect] keyword ...
```

DESCRIPTION
The `man` utility finds and displays online manual documentation pages. If `mansect` is provided, `man` restricts the search to the specific section of the manual.

The sections of the manual are:

1. General Commands Manual
2. System Calls Manual
3. Library Functions Manual
4. Kernel Interfaces Manual
5. File Formats Manual
6. Games Manual
7. Miscellaneous Information Manual
8. System Manager's Manual
9. Kernel Developer's Manual

Options that `man` understands:

Searching Manual Pages

Search within all the manual pages

- man -k <search term>

Search within specific manual pages

- while on the page, type forward slash '/' followed by the **term** you want to search for, then hit '**enter**'
- press '**n**' to go to the next occurrence

```
[marcoautili@iMac-di-User ~ %  
[marcoautili@iMac-di-User ~ % man -k pluto  
makewhatis: /Library/TeX/texbin/man: Not a directory  
pluto: nothing appropriate  
[marcoautili@iMac-di-User ~ % man -k order  
makewhatis: /Library/TeX/texbin/man: Not a directory  
bindtags(ntcl)           - Determine which bindings apply to a window, and order of evaluation  
lastcomm(1)              - show last commands executed in reverse order  
lower(ntcl)              - Change a window's position in the stacking order  
lreverse(ntcl)           - Reverse the order of a list  
operator(7)              - C operator precedence and order of evaluation  
raise(ntcl)              - Change a window's position in the stacking order  
reversetemplated(8)      - daemon that processes user content in order to detect events  
suggestd(8)              - daemon that processes user content in order to detect contacts, events, named entities, etc  
git-rev-list(1)           - Lists commit objects in reverse chronological order  
lorder(1)                - list dependencies for object files  
[marcoautili@iMac-di-User ~ %
```

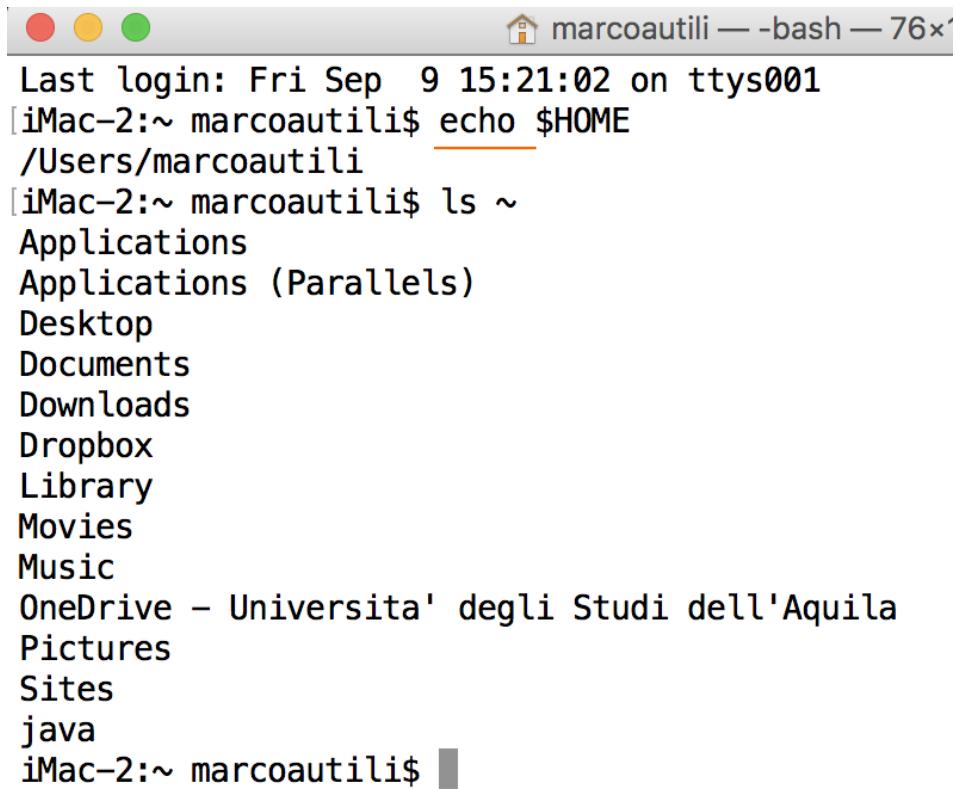
On MacOS Ventura in 2023...

More on apropos

- [https://en.wikipedia.org/wiki/Apropos_\(Unix\)](https://en.wikipedia.org/wiki/Apropos_(Unix))

Path shortcuts

- `~` (tilde) - shortcut for the home directory
 - e.g., `/Users/marcoautili/Documents` is equivalent to `~/Documents`
 - try `echo ~`
- `.` (dot) - reference to the current directory
- `..` (dotdot) - reference to the parent directory (it can use this several times, e.g., `ls ../../`)



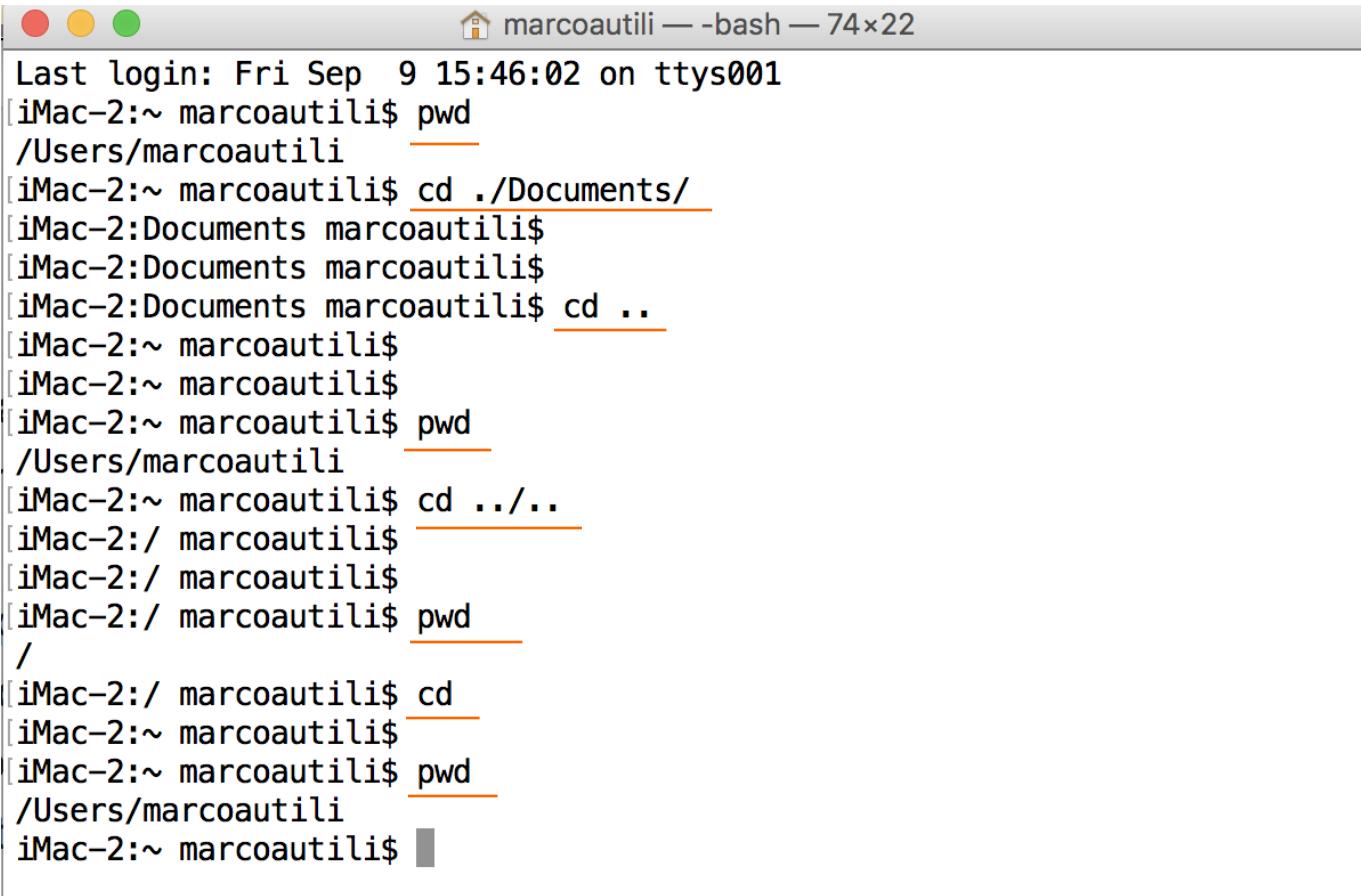
A screenshot of a macOS terminal window titled "marcoautili — bash — 76x1". The window shows the user's home directory structure. The command `echo $HOME` was run, outputting `/Users/marcoautili`. Then, the `ls ~` command was run, listing the contents of the user's home directory: Applications, Applications (Parallels), Desktop, Documents, Downloads, Dropbox, Library, Movies, Music, OneDrive – Universita' degli Studi dell'Aquila, Pictures, Sites, and java.

```
Last login: Fri Sep  9 15:21:02 on ttys001
[iMac-2:~ marcoautili$ echo $HOME
/Users/marcoautili
[iMac-2:~ marcoautili$ ls ~
Applications
Applications (Parallels)
Desktop
Documents
Downloads
Dropbox
Library
Movies
Music
OneDrive – Universita' degli Studi dell'Aquila
Pictures
Sites
java
[iMac-2:~ marcoautili$ ]
```

Navigating the file system

cd [location]

- it takes you back to your home directory when called without any arguments



```
Last login: Fri Sep  9 15:46:02 on ttys001
[iMac-2:~ marcoautili$ pwd
/Users/marcoautili
[iMac-2:~ marcoautili$ cd ./Documents/
[iMac-2:Documents marcoautili$
[iMac-2:Documents marcoautili$ cd ..
[iMac-2:~ marcoautili$ cd ..
[iMac-2:~ marcoautili$ pwd
/Users/marcoautili
[iMac-2:~ marcoautili$ cd ../../
[iMac-2:/ marcoautili$ cd ..
[iMac-2:/ marcoautili$ pwd
/
[iMac-2:/ marcoautili$ cd ..
[iMac-2:~ marcoautili$ pwd
/Users/marcoautili
[iMac-2:~ marcoautili$ ]
```

Listing directory contents

- Remember **Absolute** and **Relative Paths**
- **pwd** (stands for Print WOrking Directory)
- **ls [options] [location]** (stands for list)
 - **ls** (listing of the current directory)
 - **ls -l** (long listing of the current directory)
 - **ls /etc** (listing of the /etc directory in the root /)
 - **ls -l /etc** (long listing of the /etc directory)
 - **ls documents** or **ls ./documents** (listing of the documents directory in the current directory)
 - **ls -la** (include directory entries whose names begin with a dot "." - i.e., **hidden files or directories**, see next slide)
 - **ls -1** (list one entry per line)
 - **ls -1t** (list one entry per line and sort by time modified - most recently modified first- before sorting the operands by lexicographical order)
 - see also **ls -all** and compare it with **ls -al** and **ls -la**

→| Tab Completion helps us typing command by offering the shell's quoting and escape mechanism

Hidden files and directories

A file or directory whose name begins with a “.” (full stop) is considered to be hidden

```
marcoautili@iMac-di-User ~ %
marcoautili@iMac-di-User ~ %
marcoautili@iMac-di-User ~ %
marcoautili@iMac-di-User ~ % cd Documents
marcoautili@iMac-di-User Documents %
marcoautili@iMac-di-User Documents %
marcoautili@iMac-di-User Documents % ls -l
total 3744
drwxr-xr-x  3 marcoautili  staff      96 Sep 10  2018 Parallels
drwxr-xr-x  3 marcoautili  staff      96 Jul 28  2015 RDC Connections
-rw-r--r--@ 1 marcoautili  staff    7424 Dec 12  2018 Terminal Saved Output.txt
-rw-r--r--@ 1 marcoautili  staff     178 May 23  2015 Untitled.tex.rtf
drwxr-xr-x@ 2 marcoautili  staff      64 Oct 28  2019 WebEx
-rw-r--r--@ 1 marcoautili  staff     402 Mar 14  2018 desktop.ini
-rw-r--r--@ 1 marcoautili  staff  1899478 Dec  5  2019 website.zip
drwxr-xr-x  5 marcoautili  staff     160 Sep 10  2018 workspace
marcoautili@iMac-di-User Documents %
marcoautili@iMac-di-User Documents % ls -la
total 3760
drwx-----+ 13 marcoautili  staff     416 Dec  5  2019 .
drwxr-xr-x+ 56 marcoautili  staff   1792 Oct  5  16:30 ..
-rw-r--r--@ 1 marcoautili  staff    6148 Nov 12  2017 .DS_Store
-rw-----  1 marcoautili  staff      0 Sep 10  2018 .localized
drwxr-xr-x  2 marcoautili  staff      64 Oct 30  2015 .parallels-vm-directory
drwxr-xr-x  3 marcoautili  staff      96 Sep 10  2018 Parallels
drwxr-xr-x  3 marcoautili  staff     96 Jul 28  2015 RDC Connections
-rw-r--r--@ 1 marcoautili  staff    7424 Dec 12  2018 Terminal Saved Output.txt
-rw-r--r--@ 1 marcoautili  staff     178 May 23  2015 Untitled.tex.rtf
drwxr-xr-x@ 2 marcoautili  staff      64 Oct 28  2019 WebEx
-rw-r--r--@ 1 marcoautili  staff     402 Mar 14  2018 desktop.ini
-rw-r--r--@ 1 marcoautili  staff  1899478 Dec  5  2019 website.zip
drwxr-xr-x  5 marcoautili  staff     160 Sep 10  2018 workspace
marcoautili@iMac-di-User Documents %
```

Fields displayed by the ls command

- First character – e.g., file (-), directory (d), soft link (l), etc.
(soft links are also called symbolic links)
- Next 9 characters - permissions
- The next field - number of hard links of the file, when referring to a file, or the number of contained directory entries, when referring to a directory (more on that later on)
- The next field - owner name (root in this case)
- The next field - owner group (e.g., admin)
- The next field - file size
- The next field - file modification time
- The last field - name of the file or directory or link

```
iMac-2:/ marcoautili$ pwd  
/  
iMac-2:/ marcoautili$ ls -l  
total 45  
drwxrwxr-x+ 76 root admin 2584 Sep  5 12:58 Applications  
drwxr-xr-x+ 67 root wheel 2278 Dec 18 2015 Library  
drwxr-xr-x@  2 root wheel  68 Aug 24 2015 Network  
drwxr-xr-x@  4 root wheel 136 Sep  5 13:01 System  
drwxr-xr-x   7 root admin 238 Dec 17 2015 Users  
drwxrwxrwt@  4 root admin 136 Sep  8 16:18 Volumes  
drwxr-xr-x@ 39 root wheel 1326 Aug 25 02:03 bin  
drwxrwxr-t@  2 root admin 68 Aug 24 2015 cores  
dr-xr-xr-x   3 root wheel 4173 Sep  6 10:17 dev  
lrwxr-xr-x@  1 root wheel 11 Oct 18 2015 etc -> private/etc  
dr-xr-xr-x   2 root wheel 1 Sep  9 14:49 home  
-rw-r--r--@  1 root wheel 313 Aug 23 2015 installer.failurerequests  
dr-xr-xr-x   2 root wheel 1 Sep  9 14:49 net  
drwxr-xr-x   3 root wheel 102 Dec  2 2015 opt  
drwxr-xr-x@  6 root wheel 204 Oct 18 2015 private  
drwxr-xr-x@ 59 root wheel 2006 Aug 25 02:03 sbin  
lrwxr-xr-x@  1 root wheel 11 Oct 18 2015 tmp -> private/tmp  
drwxr-xr-x@ 12 root wheel 408 Dec 23 2015 usr  
lrwxr-xr-x@  1 root wheel 11 Oct 18 2015 var -> private/var  
iMac-2:/ marcoautili$
```

File permissions (three sets of characters, three times, indicating permissions for owner, group and other)

r = readable

w = writable

x = executable

Permissions for a directory

- Note that **read** permission for a directory and **execute** permission for a directory mean different things
- **Read permission** lets us read the directory, obtaining a list of all the filenames in the directory
- **Execute permission** lets us pass through the directory when it is a component of a pathname that we are trying to access
- We need to search the directory to look for a specific filename

File permissions (three sets of characters, three times, indicating permissions for owner, group and other)

r = readable **w = writable** **x = executable**

The **wheel** group

- In computing, the term **wheel** refers to a user account **with a wheel bit**, a system setting that provides additional special system privileges that empower a user to execute restricted commands that ordinary user accounts cannot access.
- The term is **derived from the slang phrase "big wheel"**, referring to a person with great power or influence.
- It was first used in this context with regard to the TENEX operating system, later distributed under the name TOPS-20 in the 1960s and early 1970s.
- The term was adopted by Unix users in the 1980s, due to the movement of operating system developers and users from TENEX/TOPS-20 to Unix.
- Modern Unix implementations generally **include a security protocol** that requires a user to be a member of the **wheel user privileges group** in order to gain **superuser** access to a machine by using the **su command**.
- Modern Unix systems use user groups to control access privileges. The **wheel group** is a special user group used on some Unix systems to control access to the **su command**, which **allows a user to masquerade as another user (usually the super user)**

An anticipation on soft links

With a symbolic link, the actual contents of the file (i.e., the data blocks) **store the name of the file that the symbolic link points to**

For example, the filename in the directory entry is the three-character string **var** and the **11 bytes of data in the file** are **private/var** (one byte for char)

More on that on PART 4

```
iMac-2:/ marcoautili$ pwd
/
iMac-2:/ marcoautili$ ls -l
total 45
drwxrwxr-x+ 76 root admin 2584 Sep  5 12:58 Applications
drwxr-xr-x+ 67 root wheel 2278 Dec 18 2015 Library
drwxr-xr-x@  2 root wheel   68 Aug 24 2015 Network
drwxr-xr-x@  4 root wheel  136 Sep  5 13:01 System
drwxr-xr-x   7 root admin  238 Dec 17 2015 Users
drwxrwxrwt@  4 root admin  136 Sep  8 16:18 Volumes
drwxr-xr-x@ 39 root wheel 1326 Aug 25 02:03 bin
drwxrwxr-t@  2 root admin   68 Aug 24 2015 cores
dr-xr-xr-x   3 root wheel 4173 Sep  6 10:17 dev
lrwxr-xr-x@  1 root wheel   11 Oct 18 2015 etc -> private/etc
dr-xr-xr-x   2 root wheel   1 Sep  9 14:49 home
-rw-r--r--@  1 root wheel  313 Aug 23 2015 installer.failurerequests
dr-xr-xr-x   2 root wheel   1 Sep  9 14:49 net
drwxr-xr-x   3 root wheel  102 Dec  2 2015 opt
drwxr-xr-x@  6 root wheel  204 Oct 18 2015 private
drwxr-xr-x@ 59 root wheel 2006 Aug 25 02:03 sbin
lrwxr-xr-x@  1 root wheel   11 Oct 18 2015 tmp -> private/tmp
drwxr-xr-x@ 12 root wheel  408 Dec 23 2015 usr
lrwxr-xr-x@  1 root wheel   11 Oct 18 2015 var -> private/var
iMac-2:/ marcoautili$
```

```
[drwxr-xr-x@  9 root wheel  288 Oct 25 18:22 usr
[lrwxr-xr-x@  1 root wheel   11 Nov 11 19:07 var -> private/var
[valentina-vai:/ marcoautili$ 
valentina-vai:/ marcoautili$ 
valentina-vai:/ marcoautili$ readlink var
private/var
valentina-vai:/ marcoautili$
```

Filesystem block

The block size specifies the size that the filesystem uses to read and write data

- Larger block sizes will help improve disk I/O performance when using large files, such as databases. This happens because the disk can read or write data for a longer period of time before having to search for the next block.
- On the downside, if you are going to have a lot of smaller files on that filesystem, like the /etc, there the potential for a lot of wasted disk space.
- For example, if you set your block size to 4096, or 4K, and you create a file that is 256 bytes in size, it will still consume 4K of space on your hard drive. For one file that may seem trivial, but when your filesystem contains hundreds or thousands of files, this can add up.
- Block size can also affect the maximum supported file size on some filesystems.
 - This is because many modern filesystem are limited not by block size or file size, but by the number of blocks. Therefore, you would be using a "block size * max # of blocks = max file size" formula.

* <http://www.tldp.org/LDP/sag/html/filesystems.html>

Back to the size of files...

The screenshot shows a Mac OS X file info window for the file "DISCLAIMER authors and publisher.txt" and a terminal session on an iMac running macOS.

File Info Window:

- File name: DISCLAIMER authors and publisher.txt
- Type: TEXT
- Kind: Plain Text
- Size: 1.086 bytes (4 KB on disk)
- Last modified: Yesterday 14:06
- File size: 1 KB (highlighted by an orange arrow)

Terminal Session:

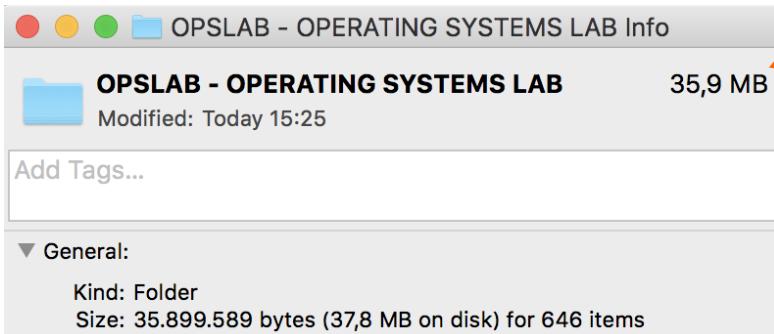
```
iMac-2:src marcoautili$ ls -la
total 104
drwxr-xr-x@ 16 marcoautili  staff  544 Sep  8 16:52 .
drwxr-xr-x@  5 marcoautili  staff  170 Sep  9 17:35 ..
-rw-r--r--@  1 marcoautili  staff  6148 Sep  9 17:57 .DS_Store
-rw-r--r--@  1 marcoautili  staff  1086 Sep  8 14:06 DISCLAIMER authors and publisher.txt
-rw-r--r--@  1 marcoautili  staff  656 Apr 27 2013 DISCLAIMER teacher.txt
-rw-r--r--@  1 marcoautili  staff  528 Apr 27 2013 Make.defines.freebsd
-rw-r--r--@  1 marcoautili  staff  526 Jul 22 2012 Make.defines.linux
-rw-r--r--@  1 marcoautili  staff  533 Sep  8 17:00 Make.defines.macos
-rw-r--r--@  1 marcoautili  staff  575 Apr 27 2013 Make.defines.solaris
-rw-r--r--@  1 marcoautili  staff  41 Apr 27 2013 Make.libapue.inc
-rw-r--r--@  1 marcoautili  staff  352 Sep  8 16:53 Makefile
-rw-r--r--@  1 marcoautili  staff  224 Sep  8 14:12 README.txt
drwxr-xr-x@ 12 marcoautili  staff  408 Sep  8 17:29 chap1-intro
drwxr-xr-x@  3 marcoautili  staff  102 Mar 20 2014 include
drwxr-xr-x@ 38 marcoautili  staff  1292 Sep  8 17:29 lib
-rwxr-xr-x@  1 marcoautili  staff  467 Jul  5 2012 systype.sh

[iMac-2:src marcoautili$]
[iMac-2:src marcoautili$]
[iMac-2:src marcoautili$ du DISCLAIMER\ authors\ and\ publisher.txt
8  DISCLAIMER authors and publisher.txt
[iMac-2:src marcoautili$]
[iMac-2:src marcoautili$]
[iMac-2:src marcoautili$ du -k DISCLAIMER\ authors\ and\ publisher.txt
4  DISCLAIMER authors and publisher.txt
[iMac-2:src marcoautili$]
[iMac-2:src marcoautili$]
[iMac-2:src marcoautili$ du -h DISCLAIMER\ authors\ and\ publisher.txt
4.0K  DISCLAIMER authors and publisher.txt
[iMac-2:src marcoautili$]
```

Annotations:

- A callout box points from the "Allocation space on disk" text to the file size "1 KB" in the file info window.
- A callout box points from the "Allocation space on disk" text to the file size "1.086 bytes (4 KB on disk)" in the file info window.
- A callout box points from the "Allocation space on disk" text to the file size "1086" in the terminal output.
- A callout box points from the "Allocation space on disk" text to the file size "4.0K" in the terminal output.
- A callout box points from the "Human-readable format" text to the file size "4.0K" in the terminal output.

Back to the size of files...



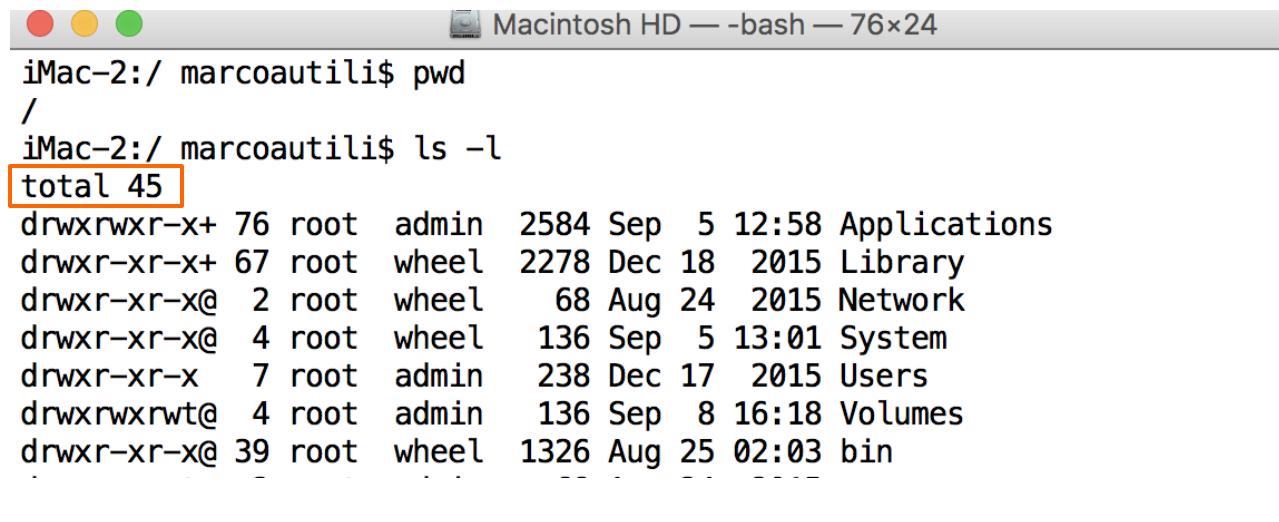
A **directory** is just a file
that contains directory
entries
(back on this later)

This is the size of space that is used to store the meta information (i.e., **the directory entries**) for the current directory

```
iMac-2:OPSLAB - OPERATING SYSTEMS LAB marcoautili$ ls -la
total 480
drwxr-xr-x@ 10 marcoautili staff      340 Sep  9 15:25 .
drwxr-xr-x@  8 marcoautili staff      272 Sep  8 22:49 ..
-rw-r--r--@  1 marcoautili staff     8196 Sep  9 19:21 .DS_Store
drwxr-xr-x@  5 marcoautili staff      170 Sep  7 13:46 EXAMS
drwxr-xr-x@  4 marcoautili staff      136 Sep  9 15:01 LEZIONI
drwxr-xr-x@  5 marcoautili staff      170 Sep  9 17:35 OPSLab-teaching-material-private
drwxr-xr-x@ 11 marcoautili staff     374 Sep  9 17:57 REFERENCE TEXTBOOKs
-rw-r--r--@  1 marcoautili staff   228550 Sep  7 13:22 SCHEDA-TUNING.pdf
drwxr-xr-x@  5 marcoautili staff      170 Sep  7 13:50 SLIDES
-rw-r--r--@  1 marcoautili staff      39 Jul 15 10:29 USEFUL LINKs.txt
iMac-2:OPSLAB - OPERATING SYSTEMS LAB marcoautili$ du -sh .
36M .
iMac-2:OPSLAB - OPERATING SYSTEMS LAB marcoautili$
```

That's the size of the
directory content!

Total



```
iMac-2:/ marcoautili$ pwd
/
iMac-2:/ marcoautili$ ls -l
total 45
drwxrwxr-x+ 76 root admin 2584 Sep  5 12:58 Applications
drwxr-xr-x+ 67 root wheel 2278 Dec 18 2015 Library
drwxr-xr-x@  2 root wheel   68 Aug 24 2015 Network
drwxr-xr-x@  4 root wheel  136 Sep  5 13:01 System
drwxr-xr-x   7 root admin  238 Dec 17 2015 Users
drwxrwxrwt@  4 root admin  136 Sep  8 16:18 Volumes
drwxr-xr-x@ 39 root wheel 1326 Aug 25 02:03 bin
```



```
marcoautili — less < man ls — 115x37
```

The Long Format

If the **-l** option is given, the following information is displayed for each file: file mode, number of links, owner name, group name, number of bytes in the file, abbreviated month, day-of-month file was last modified, hour file last modified, minute file last modified, and the pathname. In addition, for each directory whose contents are displayed, the total number of 512-byte blocks used by the files in the directory is displayed on a line by itself, immediately before the information for the files in the directory. If the file or directory has extended attributes, the permissions field printed by the **-l** option is followed by a '@' character. Otherwise, if the file or directory has extended security information (such as an access control list), the permissions field printed by the **-l** option is followed by a '+' character. If the **-%** option is given, a '%' character follows the permissions field for dataless files and directories, possibly replacing the '@' or '+' character.

More on that later

Total

If the output is to a terminal, the total number of blocks for all the files is output on a line before the listing

- **ls -s** Display the number of file system blocks actually used by each file, in units of 512 bytes, where partial units are rounded up to the next integer value
- The environment variable **BLOCKSIZE** (when set) overrides the unit size of 512 bytes

N.B.: Only files are counted (not directories)

Multiples of 8 are used since the read and write unit in my system is 4096
(i.e., $8 * 512 = 4096$ bytes)

```
● ● ● marcoautili — bash — 85x26
iMac-di-User:~ marcoautili$ ls -ls
total 24
0 drwxr-xr-x 7 marcoautili staff 224 Sep 10 2018 Applications
0 drwxr-xr-x@ 5 marcoautili staff 160 May 2 15:35 Applications (Parallels)
0 drwx-----@ 58 marcoautili staff 1856 Oct 1 17:11 Desktop
0 drwx-----+ 11 marcoautili staff 352 Dec 12 2018 Documents
0 drwx-----+ 164 marcoautili staff 5248 Oct 1 16:57 Downloads
0 drwx-----@ 18 marcoautili staff 576 Jul 17 15:42 Dropbox
0 drwx-----@ 78 marcoautili staff 2496 Apr 4 14:12 Library
0 drwx-----@ 7 marcoautili staff 224 Sep 10 2018 Movies
0 drwx-----@ 6 marcoautili staff 192 Sep 10 2018 Music
0 drwx-----@ 5 marcoautili staff 160 Jan 20 2016 OneDrive - Universita' degli
Studi dell'Aquila
0 drwx----- 2 marcoautili staff 64 May 2 15:24 Parallels
0 drwx-----+ 10 marcoautili staff 320 Sep 10 2018 Pictures
0 drwxr-xr-x+ 6 marcoautili staff 192 May 23 2015 Public
0 drwxr-xr-x@ 3 marcoautili staff 96 Sep 10 2018 Pydio
0 drwxr-xr-x+ 3 marcoautili staff 96 Sep 10 2018 Sites
0 drwxr-xr-x 7 marcoautili staff 224 Sep 24 2017 SoapUI-Tutorials
0 drwxr-xr-x 3 marcoautili staff 96 Apr 19 2016 TEMP
0 drwxr-xr-x 3 marcoautili staff 96 Oct 23 2018 Wondershare
8 -rw-r--r-- 1 marcoautili staff 453 Oct 8 2017 default-soapui-workspace.xml
0 drwxr-xr-x 3 marcoautili staff 96 Sep 27 2017 eclipse-workspace
8 -rw-r--r-- 1 marcoautili staff 129 Sep 29 2016 pippo.txt
8 -rw-r--r-- 1 marcoautili staff 2686 Oct 8 2017 soapui-settings.xml
iMac-di-User:~ marcoautili$ 

iMac-di-User:SoapUI-Tutorials marcoautili$ ls -ls
total 320
16 -rw-r--r-- 1 marcoautili staff 7268 Nov 30 2016 Sample-Authentication-Project
24 -rw-r--r-- 1 marcoautili staff 9456 Nov 30 2016 Sample-REST-Project-soapui-pr
280 -rw-r--r-- 1 marcoautili staff 142945 Nov 30 2016 Sample-SOAP-Project-soapui-pr
0 drwxr-xr-x 4 marcoautili staff 128 Sep 24 2017 WSDL-WADL
0 drwxr-xr-x 8 marcoautili staff 256 Sep 24 2017 popular-apis
iMac-di-User:SoapUI-Tutorials marcoautili$ 
iMac-di-User:SoapUI-Tutorials marcoautili$ 
```

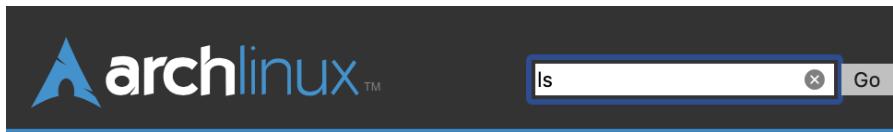
Similarities and differences

MANPAGES

Index About Manpages FAQ Service Information

debian / bullseye / coreutils / ls(1)

```
-s, --size
    print the allocated size of each file, in blocks
-k, --kibibytes
    default to 1024-byte blocks for disk usage; used only with -s and
    per directory totals
```



```
-s, --size
    print the allocated size of each file, in blocks
-k, --kibibytes
    default to 1024-byte blocks for file system usage; used only with -s and per directory
    totals
```

A screenshot of a terminal window on macOS. The title bar says "Desktop — less - man ls — 84x38". The search bar contains the text "ls". The main area shows the man page for ls(1). A red arrow points from the word "NAME" in the table of contents to the "NAME" section of the man page, which is highlighted with an orange box. The "NAME" section contains the text "ls -- list directory contents".

LS(1)	BSD General Commands Manual	LS(1)
NAME	From my macOS	
ls -- list directory contents		

-s Display the number of file system blocks actually used by each file, in units of 512 bytes, where partial units are rounded up to the next integer value. If the output is to a terminal, a total sum for all the file sizes is output on a line before the listing. The environment variable BLOCKSIZE overrides the unit size of 512 bytes.

Total (more on ...)

- Starting from the early 2015, the MacBook comes with drives that use native advanced format 4096 bytes per sector drives
- Prior Macs use 512 bytes emulation even though the physical sectors are 4K
- Also, the Master File Table (MFT – more on file table on PART 4) on pre-2015 Mac drives use file record segments in 1024-byte units, whereas the 2015 Mac's drives use file record segments in 4096-byte segments
- For several years, hard drive manufacturers have been transitioning drives from 512 bytes sector size to 4096 bytes sector size. Hard drives have had physical sector sizes of 4096 bytes for many years, but the controllers in these drives present logical sector size as 512 bytes in emulation (512e) to support older computer hardware and operating systems
- After the 2015 change, Mac drive controllers present native 4096 bytes (advanced format) logical sectors to the operating system rather than 512 bytes

Total (more on ...)

- Pre-2015 Macs generally show the Device Block Size (or logical sector size) as 512 bytes and 2015 Macs as 4096 bytes.

- pre-2015_Mac\$ diskutil info / | grep 'Block Size'

Device Block Size: 512 Bytes

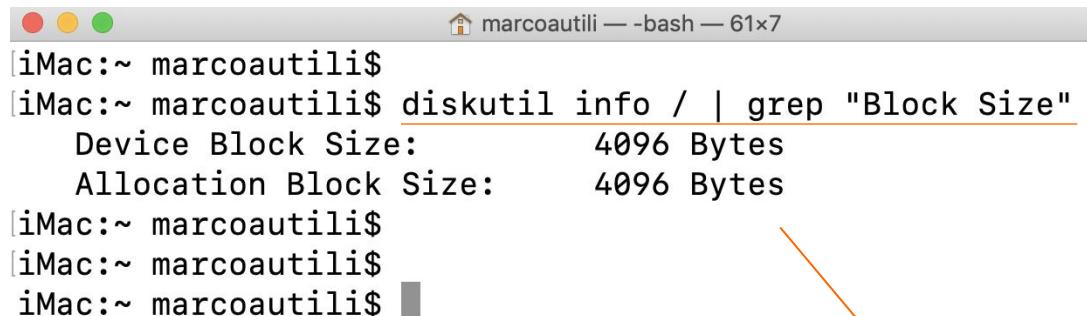
Allocation Block Size: 4096 Bytes

- 2015_Mac\$ diskutil info / | grep 'Block Size'

Device Block Size: 4096 Bytes

Allocation Block Size: 4096 Bytes

- The **Device Block Size** is the raw block size used by the hard drive controller hardware and cannot be changed
- The **Allocation Block Size** is used by the file system and is set when formatting and partitioning drives



```
[iMac:~ marcoautili$ diskutil info / | grep "Block Size"
Device Block Size:        4096 Bytes
Allocation Block Size:    4096 Bytes
[iMac:~ marcoautili$]
```

A screenshot of a macOS terminal window titled "marcoautili — bash — 61x7". The window shows the command "diskutil info / | grep "Block Size"" being run. The output displays two lines: "Device Block Size: 4096 Bytes" and "Allocation Block Size: 4096 Bytes". A red arrow points from the text "Allocation Block Size: 4096 Bytes" in the list above to the "Allocation Block Size: 4096 Bytes" line in the terminal output.

Screenshot of my Mac after the transition to 4096 bytes (October 2019)

Total (more on ...)

Screenshot of my Mac prior the transition to 4096 bytes in 2015

```
iMac-2:/ marcoautili$ pwd  
/  
iMac-2:/ marcoautili$ ls -l  
total 45  
drwxrwxr-x+ 76 root admin 2584 Sep 0 drwxr-xr-x 7 marcoautili staff 224 Sep 10 2018 Applications  
drwxr-xr-x+ 67 root wheel 2278 Dec 0 drwxr-xr-x@ 5 marcoautili staff 160 May 2 15:35 Applications (Parallels)  
drwxr-xr-x@ 2 root wheel 68 Aug 0 drwx-----@ 58 marcoautili staff 1856 Oct 1 17:11 Desktop  
drwxr-xr-x@ 4 root wheel 136 Sep 0 drwx-----+ 11 marcoautili staff 352 Dec 12 2018 Documents  
drwxr-xr-x 7 root admin 238 Dec 0 drwx-----@ 164 marcoautili staff 5248 Oct 1 16:57 Downloads  
drwxrwxrwt@ 4 root admin 136 Sep 0 drwx-----@ 18 marcoautili staff 576 Jul 17 15:42 Dropbox  
drwxr-xr-x@ 39 root wheel 1326 Aug 0 drwx-----@ 78 marcoautili staff 2496 Apr 4 14:12 Library  
drwxrwxr-t@ 2 root admin 68 Aug 0 drwx-----@ 7 marcoautili staff 224 Sep 10 2018 Movies  
dr-xr-xr-x 3 root wheel 4173 Sep 0 drwx-----@ 6 marcoautili staff 192 Sep 10 2018 Music  
lrwxr-xr-x@ 1 root wheel 11 Oct Studi dell'Aquila  
dr-xr-xr-x 2 root wheel 1 Sep 0 drwx----- 2 marcoautili staff 64 May 2 15:24 Parallels  
-rw-r--r--@ 1 root wheel 313 Aug 0 drwx-----+ 10 marcoautili staff 320 Sep 10 2018 Pictures  
dr-xr-xr-x 2 root wheel 1 Sep 0 drwxr-xr-x+ 6 marcoautili staff 192 May 23 2015 Public  
drwxr-xr-x 3 root wheel 102 Dec 0 drwxr-xr-x@ 3 marcoautili staff 96 Sep 10 2018 Pydio  
drwxr-xr-x@ 6 root wheel 204 Oct 0 drwxr-xr-x+ 3 marcoautili staff 96 Sep 10 2018 Sites  
drwxr-xr-x@ 59 root wheel 2006 Aug 0 drwxr-xr-x 7 marcoautili staff 224 Sep 24 2017 SoapUI-Tutorials  
lrwxr-xr-x@ 1 root wheel 11 Oct 0 drwxr-xr-x 3 marcoautili staff 96 Apr 19 2016 TEMP  
drwxr-xr-x@ 12 root wheel 408 Dec 0 drwxr-xr-x 3 marcoautili staff 96 Oct 23 2018 Wondershare  
lrwxr-xr-x@ 1 root wheel 11 Oct 8 -rw-r--r-- 1 marcoautili staff 453 Oct 8 2017 default-soapui-workspace.xml  
iMac-2:/ marcoautili$
```

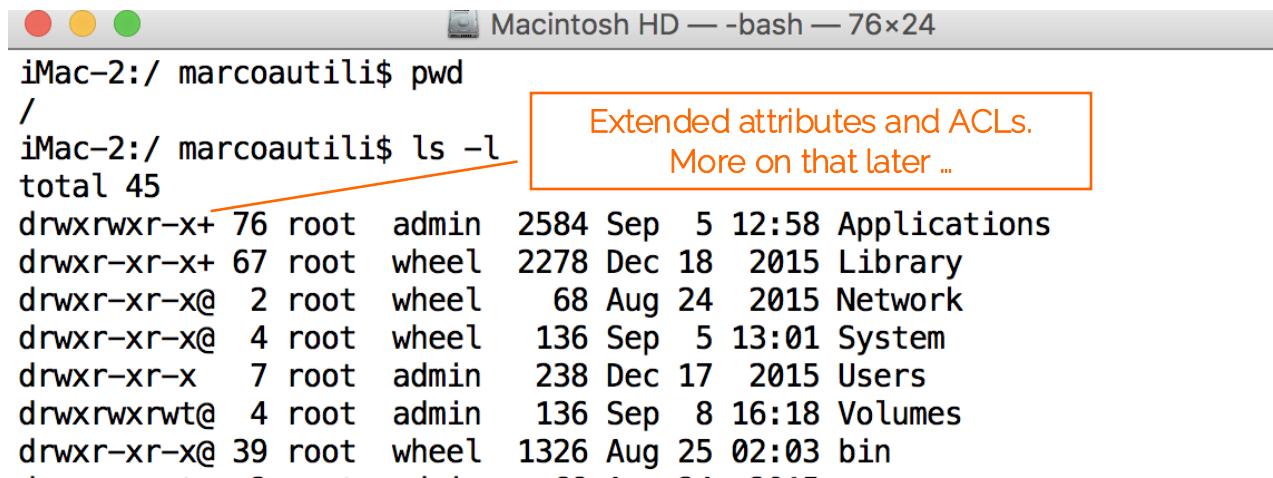
Screenshot of my Mac after the transition to 4096 bytes (October 2019)

```
iMac-di-User:~ marcoautili$ ls -ls  
total 24  
0 drwxr-xr-x 7 marcoautili staff 224 Sep 10 2018 Applications  
0 drwxr-xr-x@ 5 marcoautili staff 160 May 2 15:35 Applications (Parallels)  
0 drwx-----@ 58 marcoautili staff 1856 Oct 1 17:11 Desktop  
0 drwx-----+ 11 marcoautili staff 352 Dec 12 2018 Documents  
0 drwx-----+ 164 marcoautili staff 5248 Oct 1 16:57 Downloads  
0 drwx-----@ 18 marcoautili staff 576 Jul 17 15:42 Dropbox  
0 drwx-----@ 78 marcoautili staff 2496 Apr 4 14:12 Library  
0 drwx-----@ 7 marcoautili staff 224 Sep 10 2018 Movies  
0 drwx-----@ 6 marcoautili staff 192 Sep 10 2018 Music  
11 Oct Studi dell'Aquila  
0 drwx----- 2 marcoautili staff 64 May 2 15:24 Parallels  
313 Aug 0 drwx-----+ 10 marcoautili staff 320 Sep 10 2018 Pictures  
0 drwxr-xr-x+ 6 marcoautili staff 192 May 23 2015 Public  
0 drwxr-xr-x@ 3 marcoautili staff 96 Sep 10 2018 Pydio  
0 drwxr-xr-x+ 3 marcoautili staff 96 Sep 10 2018 Sites  
0 drwxr-xr-x 7 marcoautili staff 224 Sep 24 2017 SoapUI-Tutorials  
0 drwxr-xr-x 3 marcoautili staff 96 Apr 19 2016 TEMP  
0 drwxr-xr-x 3 marcoautili staff 96 Oct 23 2018 Wondershare  
8 -rw-r--r-- 1 marcoautili staff 453 Oct 8 2017 default-soapui-workspace.xml  
0 drwxr-xr-x 3 marcoautili staff 96 Sep 27 2017 eclipse-workspace  
8 -rw-r--r-- 1 marcoautili staff 129 Sep 29 2016 pippo.txt  
8 -rw-r--r-- 1 marcoautili staff 2686 Oct 8 2017 soapui-settings.xml  
iMac-di-User:~ marcoautili$
```

Now it is clear why "45" is not a multiple of 8...

Extended attributes and ACL

- Extended attributes
- Access Control List



```
iMac-2:/ marcoautili$ pwd  
/  
iMac-2:/ marcoautili$ ls -l  
total 45  
drwxrwxr-x+ 76 root admin 2584 Sep  5 12:58 Applications  
drwxr-xr-x+ 67 root wheel 2278 Dec 18 2015 Library  
drwxr-xr-x@  2 root wheel   68 Aug 24 2015 Network  
drwxr-xr-x@  4 root wheel  136 Sep  5 13:01 System  
drwxr-xr-x  7 root admin  238 Dec 17 2015 Users  
drwxrwxrwt@ 4 root admin  136 Sep  8 16:18 Volumes  
drwxr-xr-x@ 39 root wheel 1326 Aug 25 02:03 bin
```



The Long Format

If the **-l** option is given, the following information is displayed for each file: file mode, number of links, owner name, group name, number of bytes in the file, abbreviated month, day-of-month file was last modified, hour file last modified, minute file last modified, and the pathname. In addition, for each directory whose contents are displayed, the total number of 512-byte blocks used by the files in the directory is displayed on a line by itself, immediately before the information for the files in the directory. If the file or directory has extended attributes, the permissions field printed by the **-l** option is followed by a '@' character. Otherwise, if the file or directory has extended security information (such as an access control list), the permissions field printed by the **-l** option is followed by a '+' character. If the **-%** option is given, a '%' character follows the permissions field for dataless files and directories, possibly replacing the '@' or '+' character.

https://en.wikipedia.org/wiki/Access-control_list

Users and groups

The `id` utility displays the user and group names and numeric IDs, **of the calling process**

- `id -u` display the User ID (UID)
- `id -un` display the username
- `id -g` display the primary (a.k.a. effective) Group ID (GID)
- `id -gn` display the primary Group Name
- `id <login name or user ID>` display user and group IDs of that user
- `id -G` display the IDs of all the groups
- `id -Gn` display the names of all the groups

The `whoami` utility has been obsoleted by the `id` utility, and is equivalent to `id -un`

```
bash-3.2$ id marcoautili
uid=501(marcoautili) gid=20(staff) groups=20(staff),12(everyone),61(logicalaccounts),79(_appserverusr),80(admin),81(_appserveradm),98(_lpadmin),33(_appstore),100(_lpoperator),204(_developer),250(_analyticsusers),395(com.apple.access_ftp),398(com.apple.access_screensharing),399(com.apple.access_ssh),400(com.apple.access_remote_ae)
bash-3.2$
```

On macOS, by default, your primary group is `staff`

more on that in Part 4

More on the Directory Service

dscacheutil

stands for Directory Service Cache Utility

Performs various operations against the Directory Service cache including gathering statistics, initiating lookups, inspection, cache flush, etc.

This tool replaces most of the functionality of the `lookupd` tool previously available

The Directory Service manages information about users and resources such printers and servers ([Under macOS the Directory Service Architecture is called Open Directory](#))

Under [UNIX-like or Linux](#) systems you are probably used to modify files such `/etc/passwd` and `/etc/group` to create and modify users and groups

```
iMac-2:~ marcoautili$ dscacheutil -q group -a name staff  
name: staff  
password: *  
gid: 20  
users: root  
  
iMac-2:~ marcoautili$ iMac-2:~ marcoautili$ dscacheutil -q group -a name wheel  
name: wheel  
password: *  
gid: 0  
users: root  
  
dscl -- Directory Service Command Line utility (see next slide)  
  
iMac-2:~ marcoautili$ dscl -q group -a name wheel  
Cannot open remote host, error: DSOpenDirServiceErr  
iMac-2:~ marcoautili$ iMac-2:~ marcoautili$ dscacheutil -q group -a name wheel  
name: wheel  
password: *  
gid: 0  
users: root  
  
iMac-2:~ marcoautili$ iMac-2:~ marcoautili$ dscacheutil -q group -a name admin  
name: admin  
password: *  
gid: 80  
users: root marcoautili  
  
iMac-2:~ marcoautili$
```

This is not part of the course

More on the Directory Service

```
iMac:~ marcoautili$ dscl . list /groups
_iamavisd
_analyticsd
_analyticsusers
_appleevents
_applepay
_appowner
_appserveradm
_appserverusr
_appstore
_ard
_assetcache
_astris
_atsserver
_calendar
_captiveagent
_ces
_clamav
_cmiodalassistants
_coreaudiod
```

List all groups

Directory Service Command Line utility

dscl is a general-purpose utility for operating on Directory Service directory nodes. Its commands allow one to create, read, and manage Directory Service data.

```
iMac-2:~ marcoautili$ dscacheutil -q user -a name root
name: root
password: *
uid: 0
gid: 0
dir: /var/root
shell: /bin/sh
gecos: System Administrator

iMac-2:~ marcoautili$ dscacheutil -q user
name: _amavisd
password: *
uid: 83
gid: 83
dir: /var/virusmails
shell: /usr/bin/false
gecos: AMaViS Daemon

name: _appleevents
password: *
uid: 55
gid: 55
dir: /var/empty
shell: /usr/bin/false
gecos: AppleEvents Daemon
```

Lookup a user

Lookup all users

```
marcoautili@iMac ~ %
marcoautili@iMac ~ % dscl
Entering interactive mode... (type "help" for commands)
>
>
>
```

If invoked without any commands, dscl runs in an interactive mode, reading commands from standard input

This is not part of the course

```
name: _appowner
password: *
uid: 87
gid: 87
dir: /var/empty
shell: /usr/bin/false
gecos: Application Owner

name: _appserver
password: *
uid: 79
gid: 79
```

Homework Activities

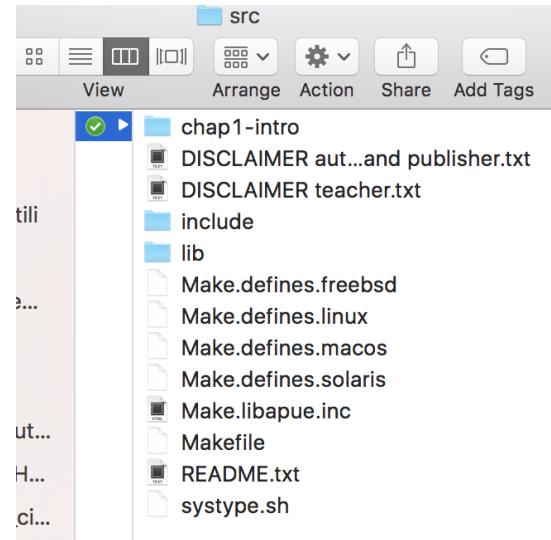
Use the commands `cd` and `ls` to navigate your system and explore your directories by using different methods and different relative and absolute paths

For example

- `/etc` - the place where config files for the system are stored
- `/var/log` - the place where log files for various system programs are stored (You may not have the right permissions...)
- `/bin` - the place where many programs are contained
- `/usr/bin` - system programs are also stored here

Everything is a File

- A directory is a file
- Mouse, keyboard and monitor are files
 - the keyboard is a **read-only** file!
 - the monitor is **write-only** file!
- **Windows** uses the extension (.txt, .png) of a file to understand what type of file it is
- **Linux/Unix-like** systems are said to be "extensionless" systems, i.e., the system looks inside the file to understand what type of file it is
 - The command **file [path]** is of help



```
iMac-2:src marcoautili$ file Make.libapue.inc
Make.libapue.inc: ASCII text
iMac-2:src marcoautili$
iMac-2:src marcoautili$
iMac-2:src marcoautili$
iMac-2:src marcoautili$
iMac-2:src marcoautili$ file systype.sh
systype.sh: ASCII English text
iMac-2:src marcoautili$
iMac-2:src marcoautili$
iMac-2:src marcoautili$
iMac-2:src marcoautili$
iMac-2:src marcoautili$
```

Case Sensitive VS Case Insensitive

Historically,

- UNIX-like and Linux systems are **Case Sensitive** (e.g., `pippo.txt` is distinct from `PIPPO.txt`)
- Window systems is **Case Insensitive**
- Linux is **Case Sensitive** also with command line options
 - `ls` command
 - `ls -S` Sort files by size
 - `ls -s` (as already said) Display the number of file system blocks actually used by each file, in units of 512 bytes, where partial units are rounded up to the next integer value. If the output is to a terminal, a total sum for all the file sizes is output on a line before the listing. The environment variable `BLOCKSIZE` (when set) overrides the unit size of 512 bytes
- macOS systems (which are Unix-based) typically have **case-insensitive file systems** (i.e., `pippo.txt` is the same as `PIPPO.txt`) so on such systems commands external to the shell are in fact treated case-insensitively. But builtins, like `cd`, remain case-sensitive (more on builtin commands later)

Again, many similarities together with many differences!

Special chars in file and directory names

Spaces in file and directory names are valid... but be careful!

Two way to deal with

1. Quotes (single or double not a problem for now...)

- cd 'Applications (Parallels)'/
- cd "Applications (Parallels)"/

2. Escape character

- cd Applications\\(Parallels\\)/
 - Note that three chars have been escaped with \, i.e., space, (and)

Remember

→| Tab Completion automatically escapes chars for you!

Long hand and short-hand version

```
Marcos-MacBook-Pro:Documents marcoautili$ ls -a
.
.. .localized RDC Connections
.DS_Store .parallels-vm-directory Untitled.tex.rtf
Parallels desktop.ini

Marcos-MacBook-Pro:Documents marcoautili$ ls -all
total 32
drwx-----+ 9 marcoautili staff 306 Aug 5 14:19 .
drwxr-xr-x++ 33 marcoautili staff 1122 Sep 12 14:55 ..
-rw-r--r--@ 1 marcoautili staff 6148 Dec 1 2015 .DS_Store
-rw----- 1 marcoautili staff 0 May 7 2015 .localized
drwxr-xr-x 2 marcoautili staff 68 Oct 30 2015 .parallels-vm-directory
drwxr-xr-x 3 marcoautili staff 102 Jul 8 09:54 Parallels
drwxr-xr-x 3 marcoautili staff 102 Jul 28 2015 RDC Connections
-rw-r--r--@ 1 marcoautili staff 178 May 23 2015 Untitled.tex.rtf
[-rw-r--r--@ 1 marcoautili staff 402 Aug 5 14:19 desktop.ini
Marcos-MacBook-Pro:Documents marcoautili$ ls -la
total 32
drwx-----+ 9 marcoautili staff 306 Aug 5 14:19 .
drwxr-xr-x++ 33 marcoautili staff 1122 Sep 12 14:55 ..
-rw-r--r--@ 1 marcoautili staff 6148 Dec 1 2015 .DS_Store
-rw----- 1 marcoautili staff 0 May 7 2015 .localized
drwxr-xr-x 2 marcoautili staff 68 Oct 30 2015 .parallels-vm-directory
drwxr-xr-x 3 marcoautili staff 102 Jul 8 09:54 Parallels
drwxr-xr-x 3 marcoautili staff 102 Jul 28 2015 RDC Connections
-rw-r--r--@ 1 marcoautili staff 178 May 23 2015 Untitled.tex.rtf
-rw-r--r--@ 1 marcoautili staff 402 Aug 5 14:19 desktop.ini
Marcos-MacBook-Pro:Documents marcoautili$ ls -alh
total 32
drwx-----+ 9 marcoautili staff 306B Aug 5 14:19 .
drwxr-xr-x++ 33 marcoautili staff 1.1K Sep 12 14:55 ..
-rw-r--r--@ 1 marcoautili staff 6.0K Dec 1 2015 .DS_Store
-rw----- 1 marcoautili staff 0B May 7 2015 .localized
drwxr-xr-x 2 marcoautili staff 68B Oct 30 2015 .parallels-vm-directory
drwxr-xr-x 3 marcoautili staff 102B Jul 8 09:54 Parallels
drwxr-xr-x 3 marcoautili staff 102B Jul 28 2015 RDC Connections
-rw-r--r--@ 1 marcoautili staff 178B May 23 2015 Untitled.tex.rtf
-rw-r--r--@ 1 marcoautili staff 402B Aug 5 14:19 desktop.ini
Marcos-MacBook-Pro:Documents marcoautili$ ls --all
ls: illegal option -- -
usage: ls [-ABCDEFGHJKLMNOPRSTUWabcdefghijklmnoprstuwx1] [file ...]
Marcos-MacBook-Pro:Documents marcoautili$
```

Under MAC OSX, it is equivalent to `ls -al` or `ls -la`

When used with the `-l` option, use unit suffixes: Byte, Kilobyte, Megabyte, Gigabyte, Terabyte and Petabyte in order to reduce the number of digits to three or less using base 2 for sizes

Try it under Linux...

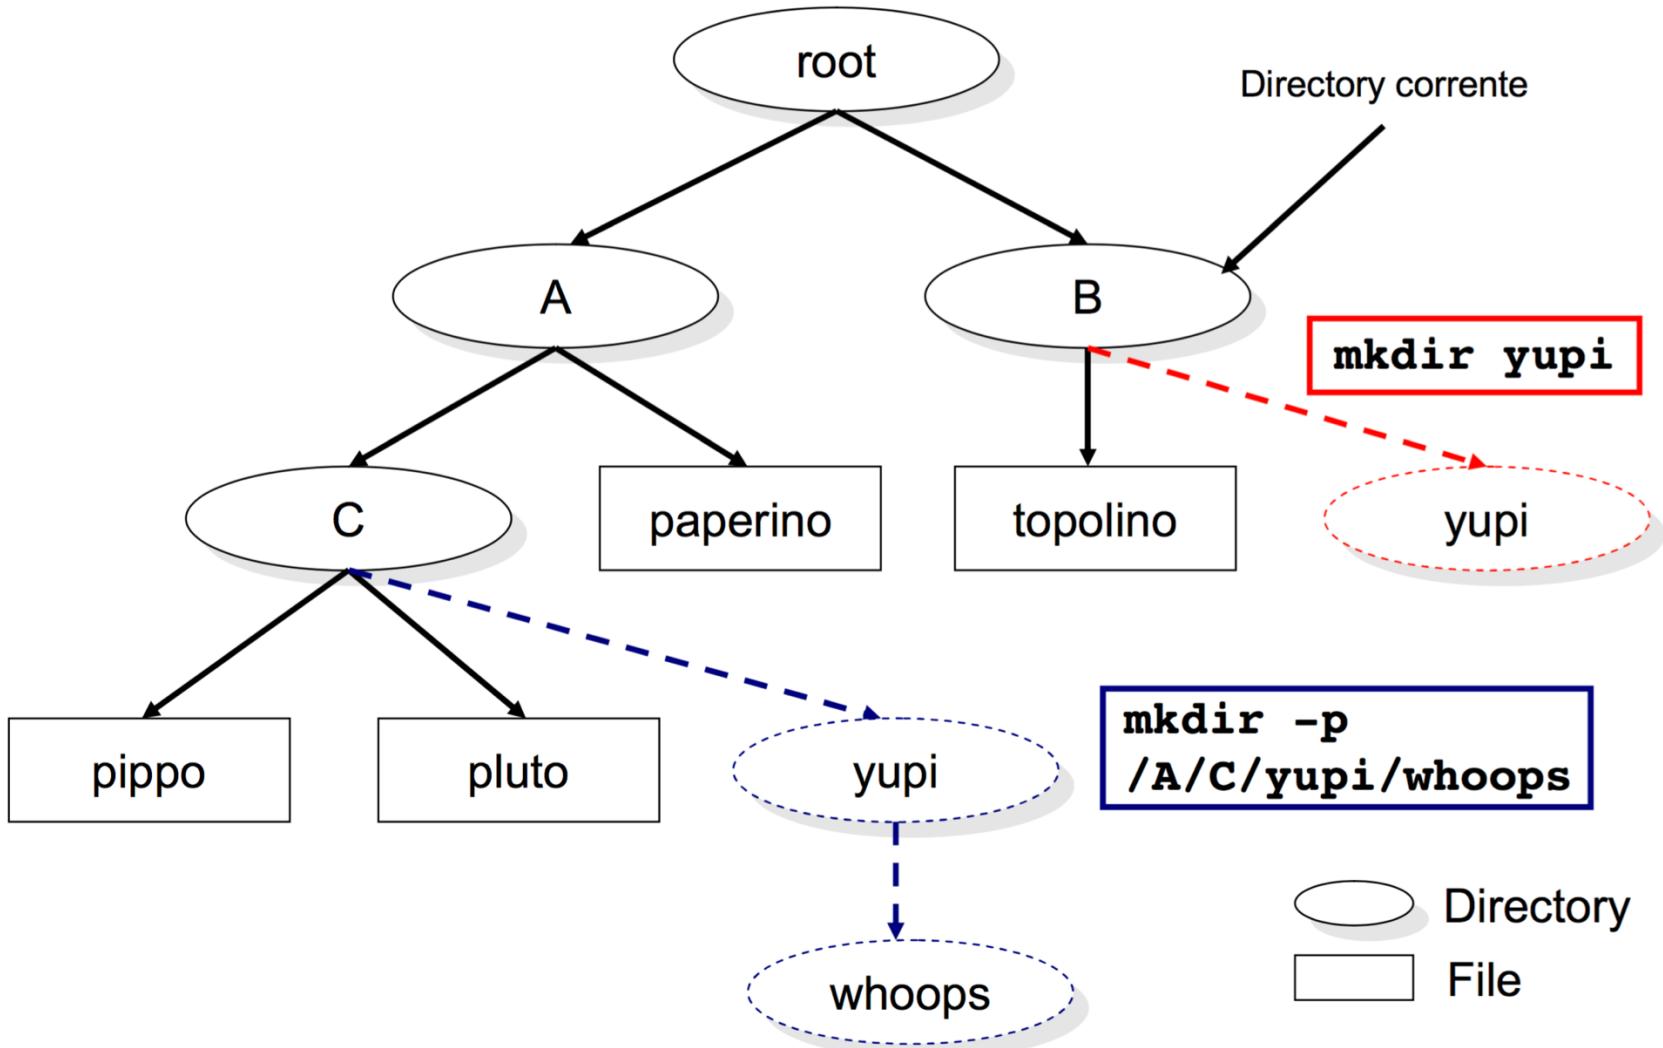
\$PATH

- When we type a command on the terminal, the system searches for the implementation of the command in a predefined set of directories
 - The predefined set of directories to be searched is specified by the variable **\$PATH**
- The system will not look inside the current directory, that's why when we want to run our own script from within the current directory, we need to use “**./**”
 - **./myfirstscript.sh**

More on that in PART III

```
[vincet:~ marcoautili$ echo $PATH
/usr/local/bin:/usr/bin:/bin:/usr/sbin:/sbin:/usr/texbin
[vincet:~ marcoautili$
[vincet:~ marcoautili$ cd /bin
[vincet:bin marcoautili$
[vincet:bin marcoautili$ ls -la
total 5168
drwxr-xr-x@ 39 root wheel 1326 Aug 3 11:09 .
drwxr-xr-x 34 root wheel 1224 Sep 21 13:24 ..
-rw xr-xr-x 1 root wheel 22464 Dec 3 2015 [
-rw xr-xr-x 1 root wheel 628496 Dec 3 2015 bash
-rw xr-xr-x 1 root wheel 23520 Dec 3 2015 cat
-rw xr-xr-x 1 root wheel 33904 Jul 9 04:52 chmod
-rw xr-xr-x 1 root wheel 28832 Jul 9 04:53 cp
-rw xr-xr-x 1 root wheel 378624 Mar 12 2016 csh
-rw xr-xr-x 1 root wheel 28368 Dec 3 2015 date
-rw xr-xr-x 1 root wheel 31856 Jul 9 04:52 dd
-rw xr-xr-x 1 root wheel 27440 Jul 9 04:52 df
-rw xr-xr-x 1 root wheel 18144 Dec 3 2015 domainname
-rw xr-xr-x 1 root wheel 18032 Dec 3 2015 echo
-rw xr-xr-x 1 root wheel 53872 Dec 3 2015 ed
-rw xr-xr-x 1 root wheel 22992 Dec 3 2015 expr
-rw xr-xr-x 1 root wheel 18192 Dec 3 2015 hostname
-rw xr-xr-x 1 root wheel 18528 Dec 3 2015 kill
-rw xr-xr-x 1 root wheel 1394432 Mar 12 2016 ksh
-rw xr-xr-x 1 root wheel 124048 Jul 9 04:52 launchctl
-rw xr-xr-x 1 root wheel 18944 Jul 9 04:52 link
-rw xr-xr-x 1 root wheel 18944 Jul 9 04:52 ln
-rw xr-xr-x 1 root wheel 38512 Jul 9 04:52 ls
-rw xr-xr-x 1 root wheel 18496 Jul 9 04:53 mkdir
-rw xr-xr-x 1 root wheel 24144 Jul 9 04:52 mv
-rw xr-xr-x 1 root wheel 110800 Jul 9 04:52 pax
-rw sr-xr-x 1 root wheel 51008 Dec 3 2015 ps
-rw xr-xr-x 1 root wheel 18176 Dec 3 2015 pwd
-rw xr-xr-x 1 root wheel 29520 Dec 3 2015 rcp
-rw xr-xr-x 1 root wheel 23744 Jul 9 04:52 rm
-rw xr-xr-x 1 root wheel 18064 Jul 9 04:52 rmdir
-rw xr-xr-x 1 root wheel 632672 Dec 3 2015 sh
-rw xr-xr-x 1 root wheel 17984 Dec 3 2015 sleep
-rw xr-xr-x 1 root wheel 32048 Dec 3 2015 stty
-rw xr-xr-x 1 root wheel 42320 Jul 9 04:52 sync
-rw xr-xr-x 1 root wheel 378624 Mar 12 2016 tcsh
-rw xr-xr-x 1 root wheel 22464 Dec 3 2015 test
-rw xr-xr-x 1 root wheel 23744 Jul 9 04:52 unlink
-rw xr-xr-x 1 root wheel 18080 Jul 9 04:52 wait4path
-rw xr-xr-x 1 root wheel 573600 Dec 3 2015 zsh
[vincet:bin marcoautili$ ]
```

Manipulating files and directories



mkdir & rmdir

mkdir [options] <DirectoryPath>

- mkdir pluto OR mkdir ./pluto
- mkdir /Users/marcoautili/Documents/pluto
- mkdir ../pluto
- mkdir ~/pluto

mkdir -p <DirectoryPath>

Create intermediate
directories as required

```
Marcos-MacBook-Pro:Desktop marcoautili$  
Marcos-MacBook-Pro:Desktop marcoautili$ mkdir pluto  
Marcos-MacBook-Pro:Desktop marcoautili$ mkdir ./pluto  
mkdir: ./pluto: File exists  
Marcos-MacBook-Pro:Desktop marcoautili$  
Marcos-MacBook-Pro:Desktop marcoautili$  
Marcos-MacBook-Pro:Desktop marcoautili$ cd pluto/  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
Marcos-MacBook-Pro:pluto marcoautili$ mkdir -p pippo/paperino  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
plippo  
Marcos-MacBook-Pro:pluto marcoautili$ cd pippo  
Marcos-MacBook-Pro:pippo marcoautili$ ls  
paperino  
Marcos-MacBook-Pro:pippo marcoautili$
```

```
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$ rmdir -p pippo/  
rmdir: pippo/: Directory not empty  
Marcos-MacBook-Pro:pluto marcoautili$ rmdir -p pippo/paperino/  
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
Marcos-MacBook-Pro:pluto marcoautili$
```

rmdir [options] <DirectoryPath>

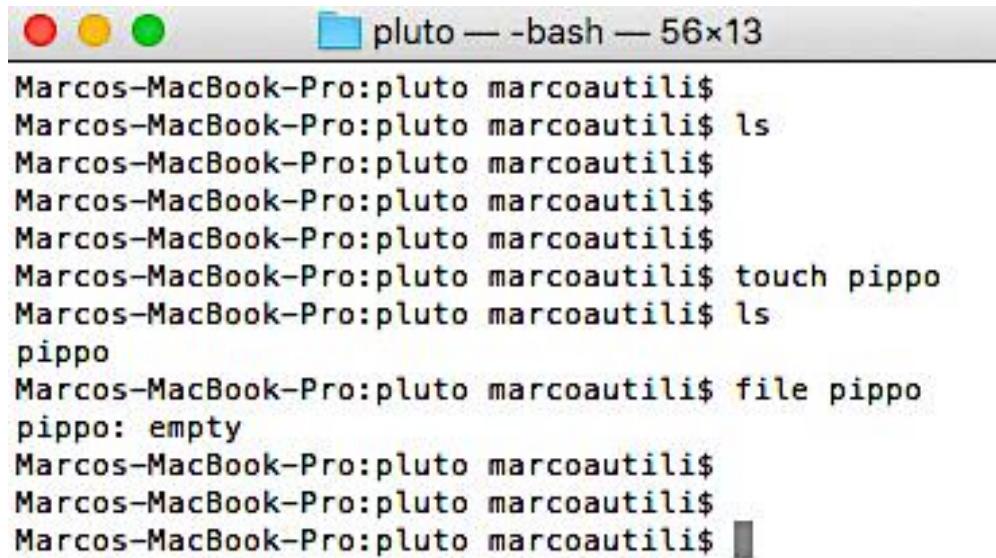
Remove a directory
if it is empty

See also later slides

Creating an empty file

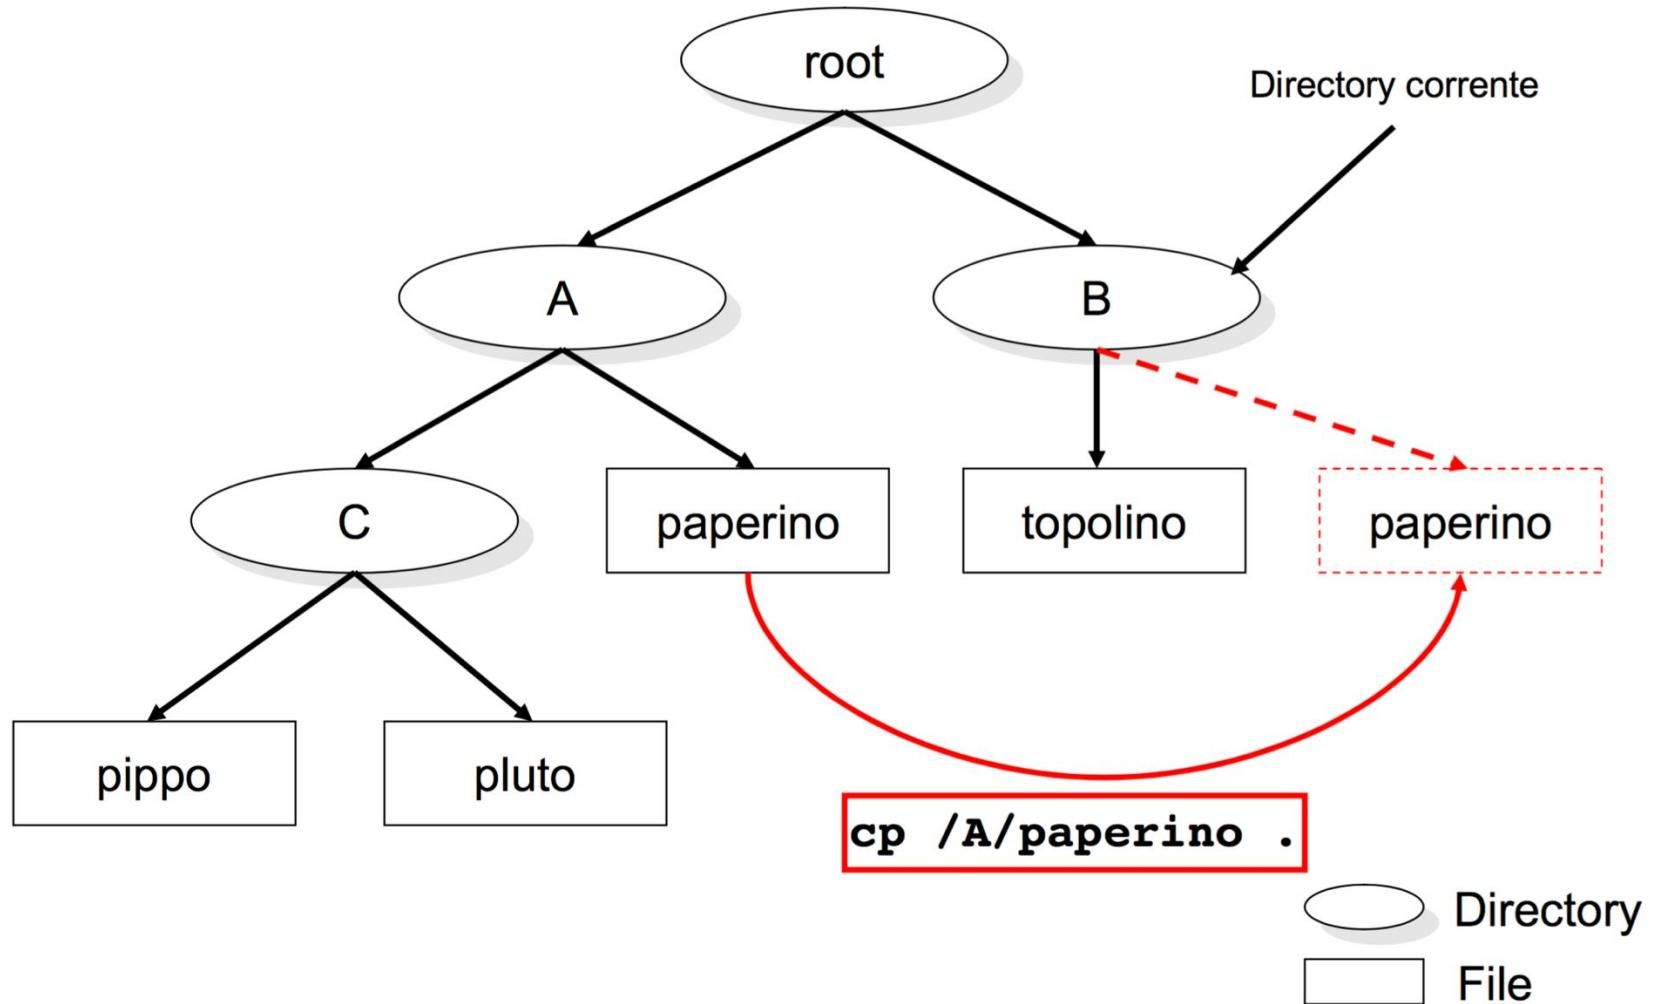
touch [options] <filename>

- change file access and modification times
 - It is not for creating files, but it creates an empty file if the specified file <filename> does not exist

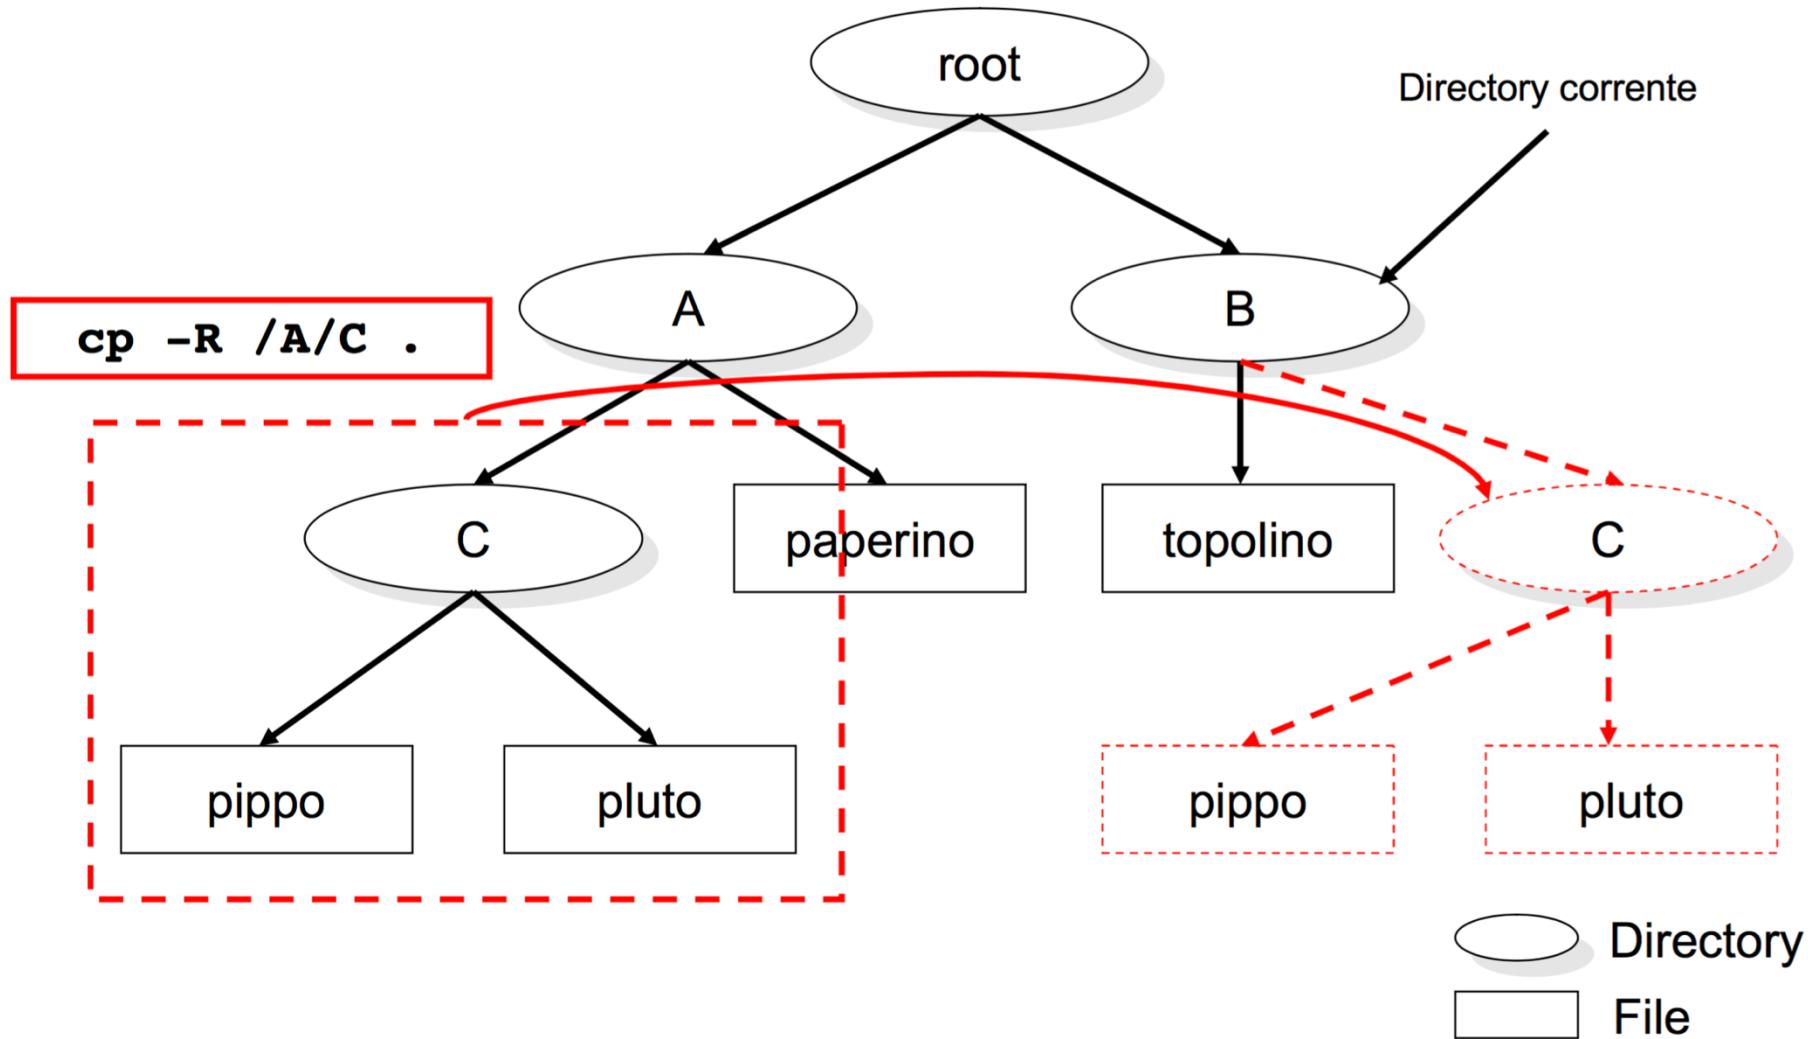


```
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$ touch pippo  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
pippo  
Marcos-MacBook-Pro:pluto marcoautili$ file pippo  
pippo: empty  
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$
```

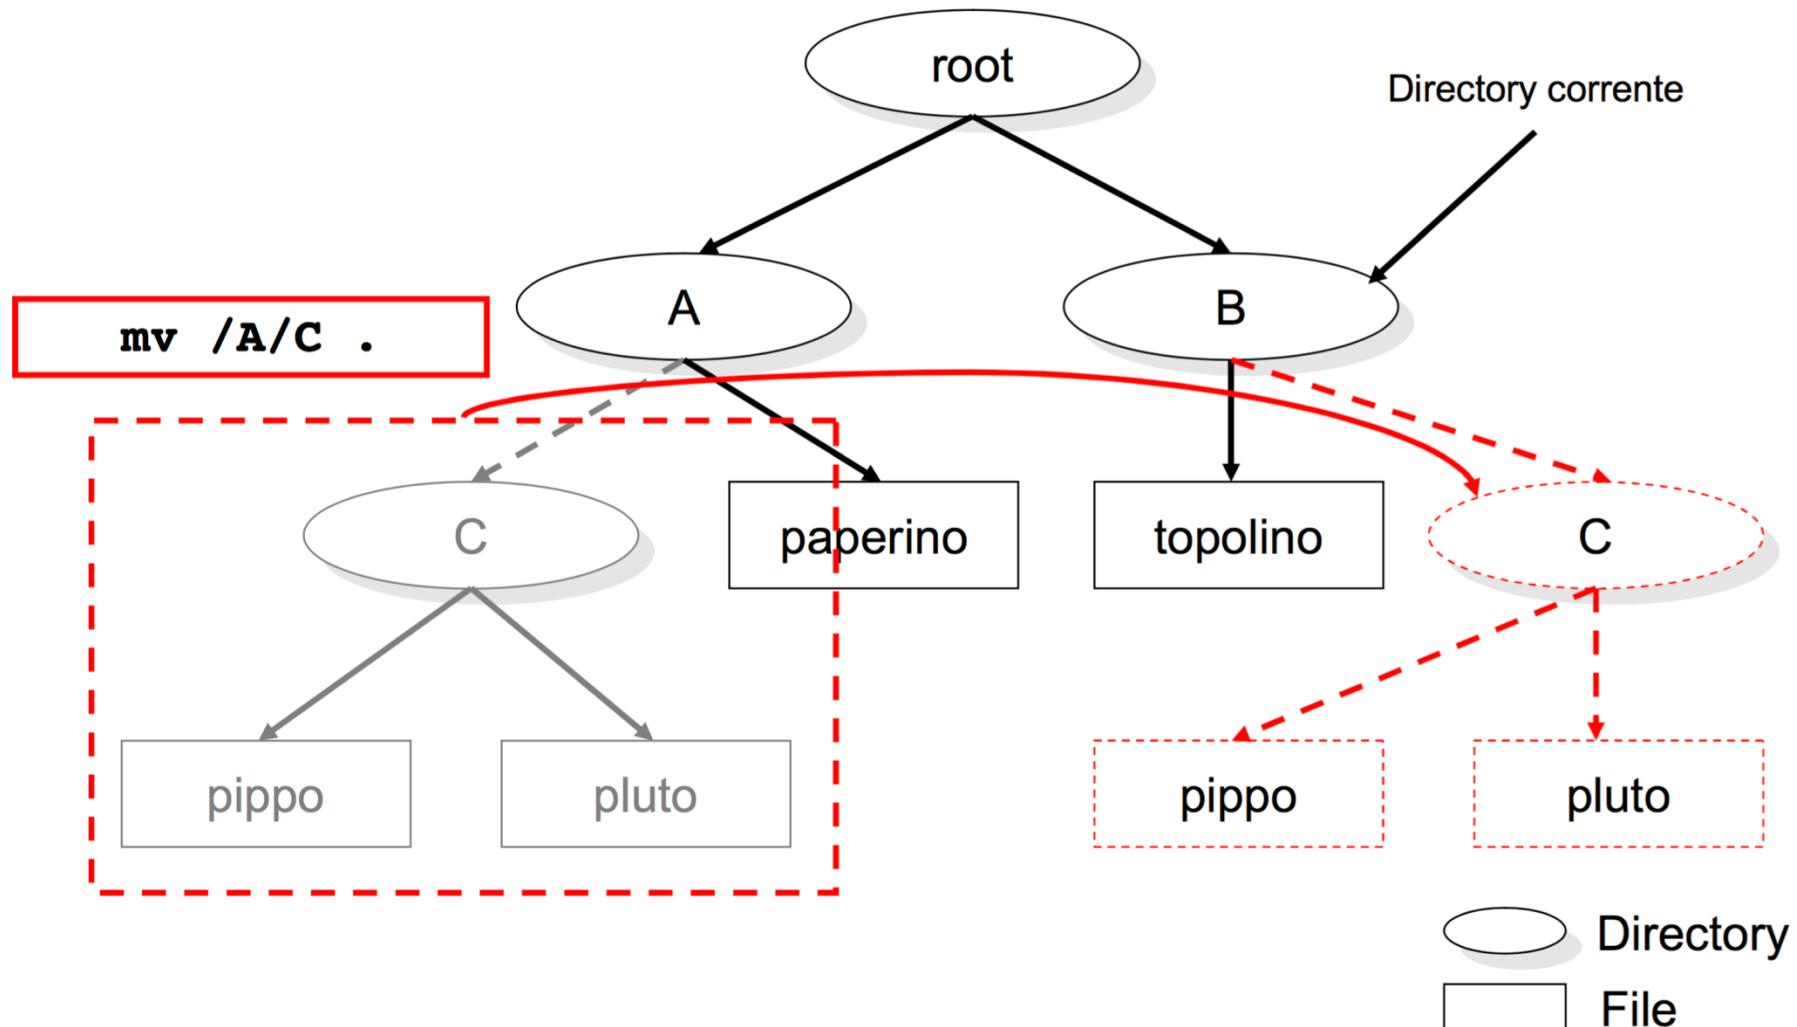
Copying



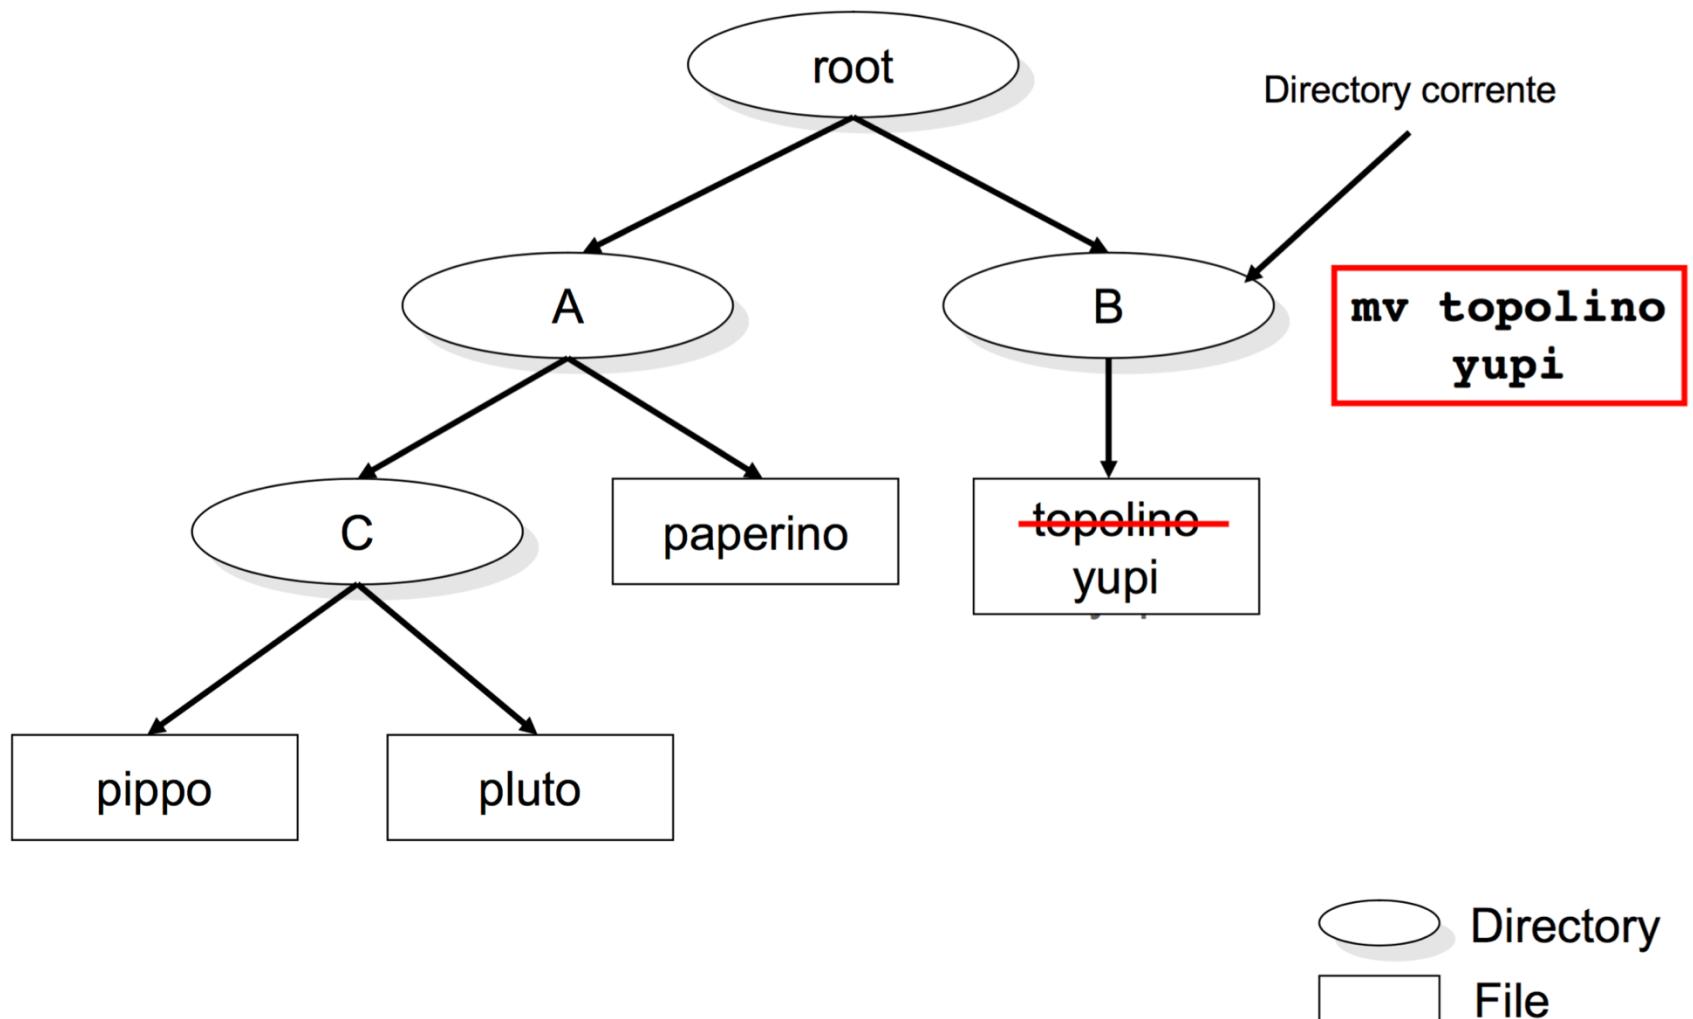
Copying



Moving



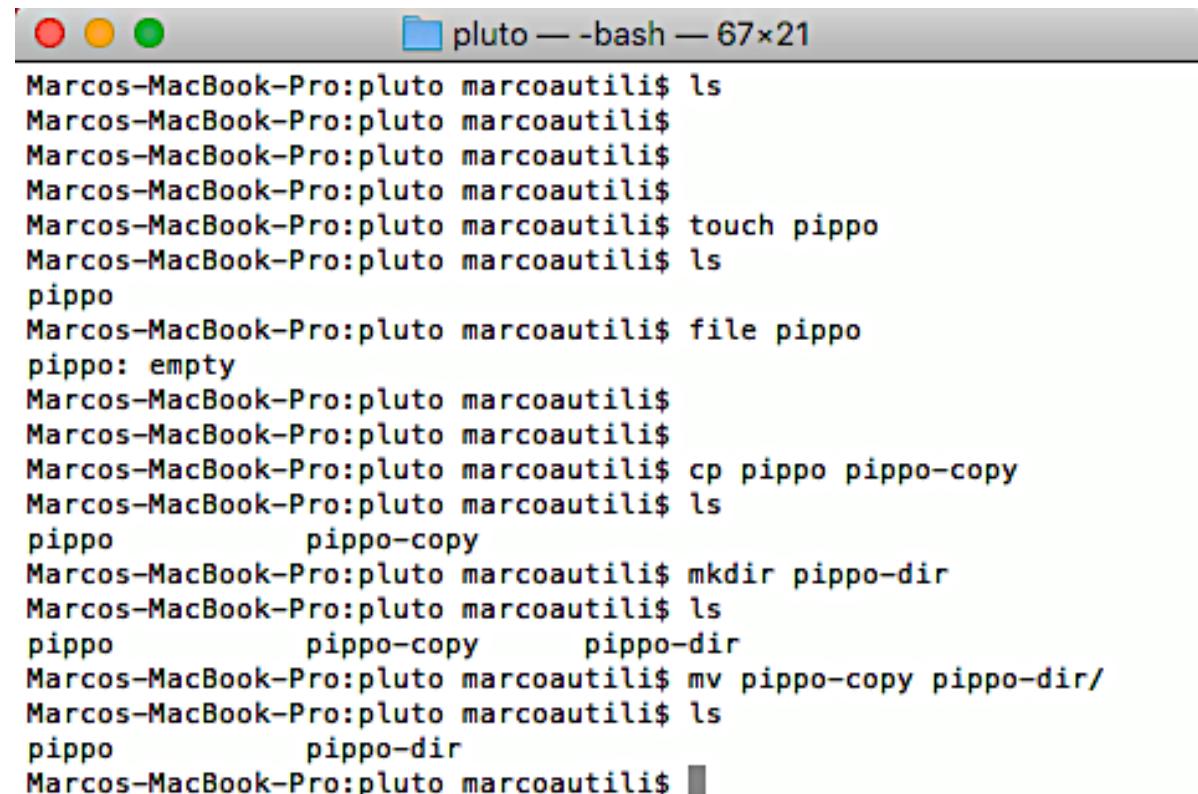
Moving (when used for renaming)



cp & mv

- cp [options] <sourcePath>
<destinationPath>

- many many options... see
the dedicated man page
- use the **-r** option (i.e.,
recursive) to copy
directories



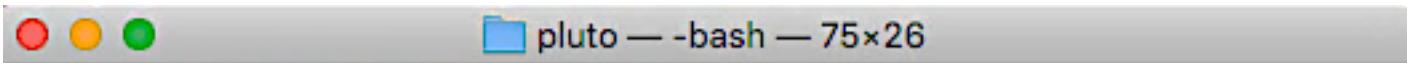
```
Marcos-MacBook-Pro:pluto marcoautili$ ls
Marcos-MacBook-Pro:pluto marcoautili$ 
Marcos-MacBook-Pro:pluto marcoautili$ 
Marcos-MacBook-Pro:pluto marcoautili$ 
Marcos-MacBook-Pro:pluto marcoautili$ touch pippo
Marcos-MacBook-Pro:pluto marcoautili$ ls
pippo
Marcos-MacBook-Pro:pluto marcoautili$ file pippo
pippo: empty
Marcos-MacBook-Pro:pluto marcoautili$ 
Marcos-MacBook-Pro:pluto marcoautili$ 
Marcos-MacBook-Pro:pluto marcoautili$ cp pippo pippo-copy
Marcos-MacBook-Pro:pluto marcoautili$ ls
pippo      pippo-copy
Marcos-MacBook-Pro:pluto marcoautili$ mkdir pippo-dir
Marcos-MacBook-Pro:pluto marcoautili$ ls
pippo      pippo-copy      pippo-dir
Marcos-MacBook-Pro:pluto marcoautili$ mv pippo-copy pippo-dir/
Marcos-MacBook-Pro:pluto marcoautili$ ls
pippo      pippo-dir
Marcos-MacBook-Pro:pluto marcoautili$ 
```

- mv [options] <source>
<destination>

- directories can be moved
without the **-r** option

Renaming

- mv [options] <source> <destination>
 - Note that it is not for renaming but it can be used for renaming both files and directories

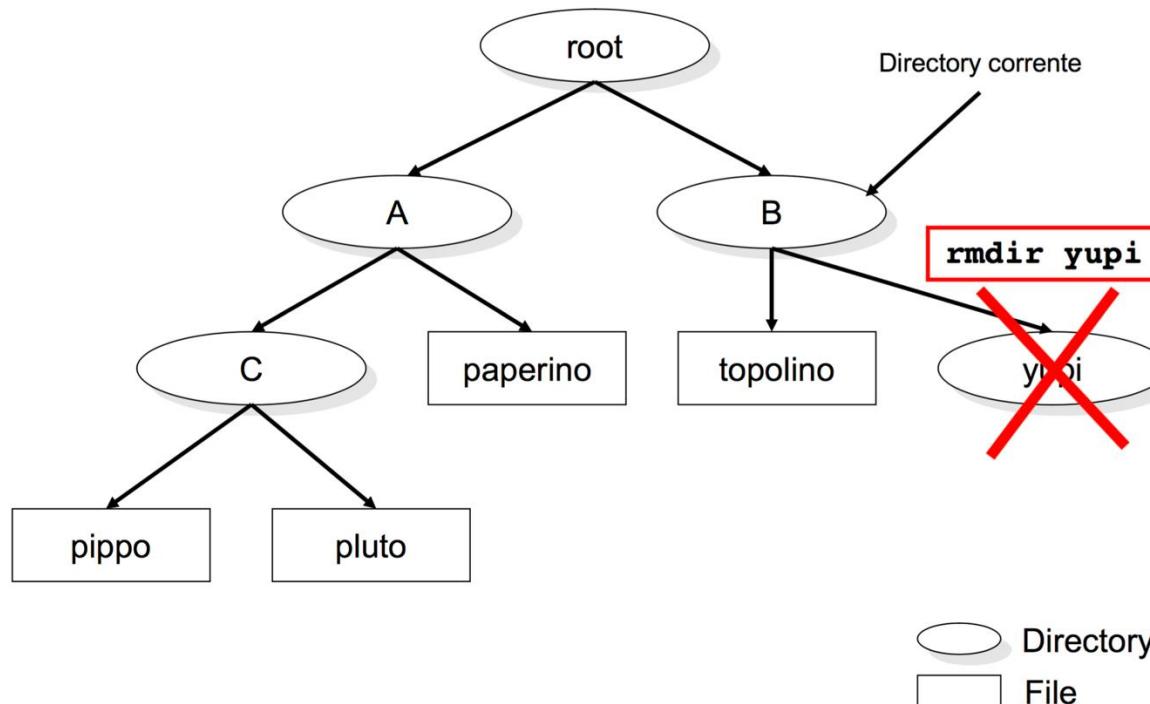


```
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$ pwd  
/Users/marcoautili/Desktop/pluto  
Marcos-MacBook-Pro:pluto marcoautili$ touch pippo.txt  
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
pippo.txt  
Marcos-MacBook-Pro:pluto marcoautili$ mv pippo.txt pippo-bis.txt  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
pippo-bis.txt  
Marcos-MacBook-Pro:pluto marcoautili$ mkdir pippo-dir  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
pippo-bis.txt pippo-dir  
Marcos-MacBook-Pro:pluto marcoautili$ touch pippo-dir/paperino.txt  
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$ ls pippo-dir/  
paperino.txt  
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
pippo-bis.txt pippo-dir  
Marcos-MacBook-Pro:pluto marcoautili$ mv pippo-dir/ pippo-dir-bis  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
pippo-bis.txt pippo-dir-bis  
Marcos-MacBook-Pro:pluto marcoautili$ ls pippo-dir-bis/  
paperino.txt  
Marcos-MacBook-Pro:pluto marcoautili$
```

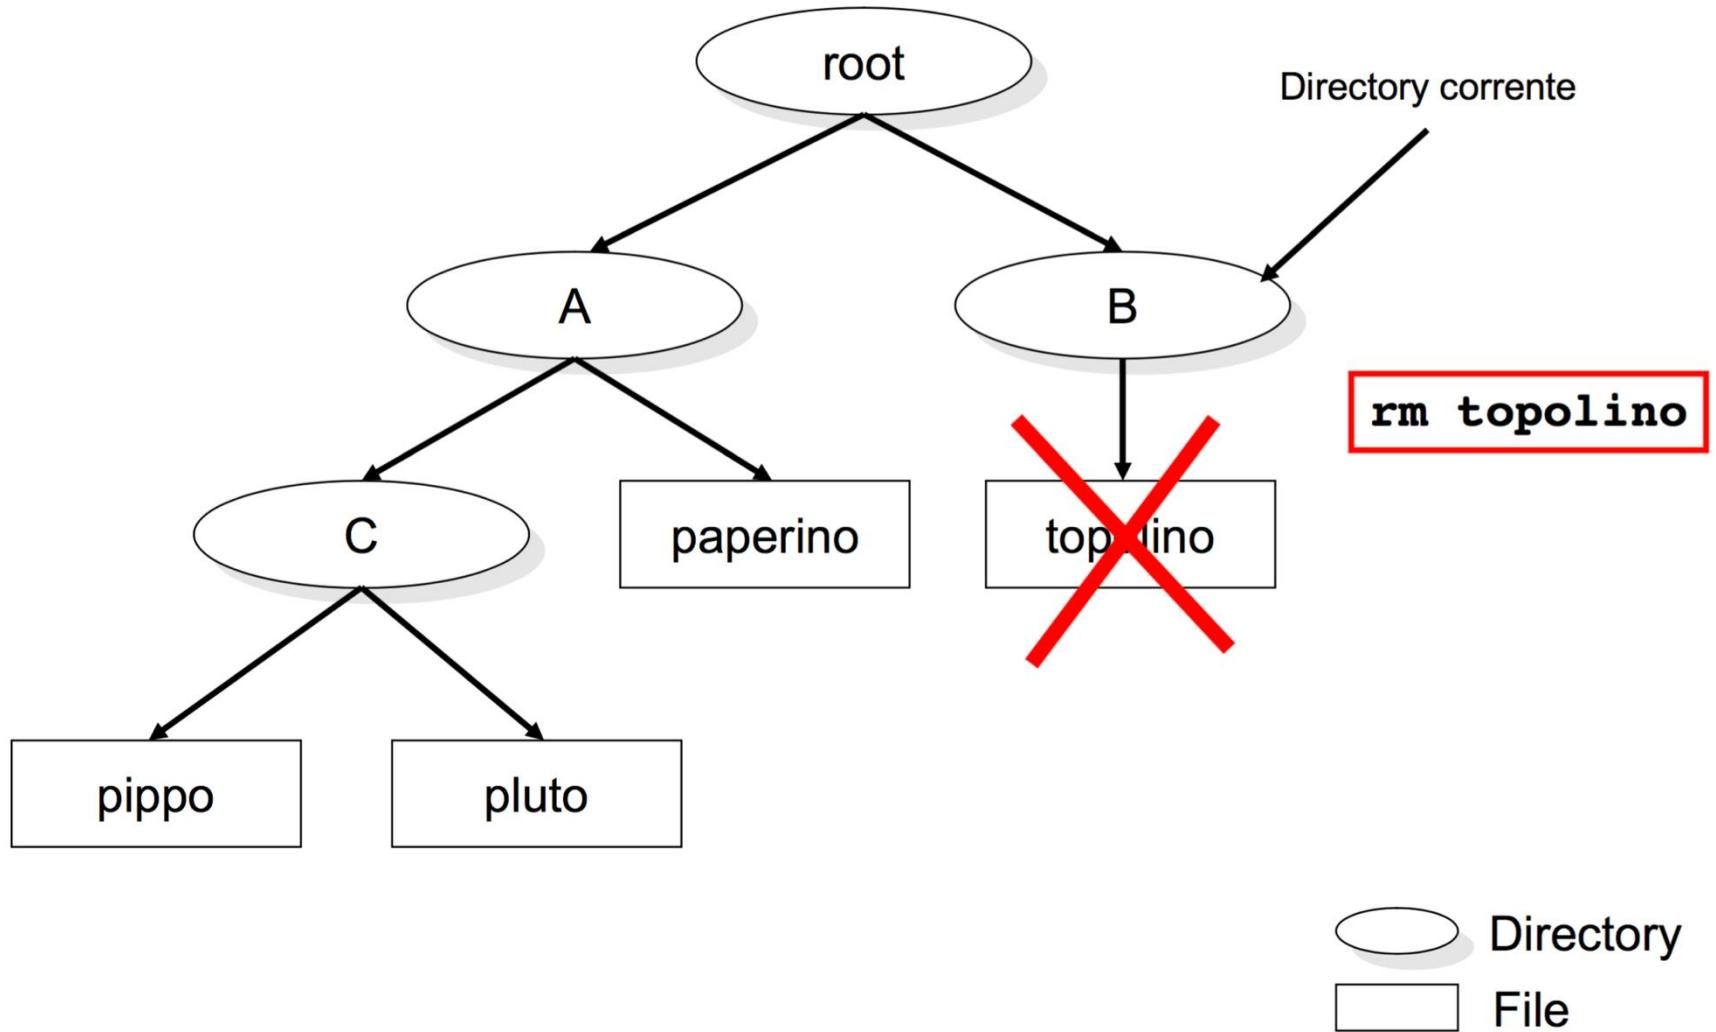
Removing empty directories

`rmdir [-p] <pathname>`

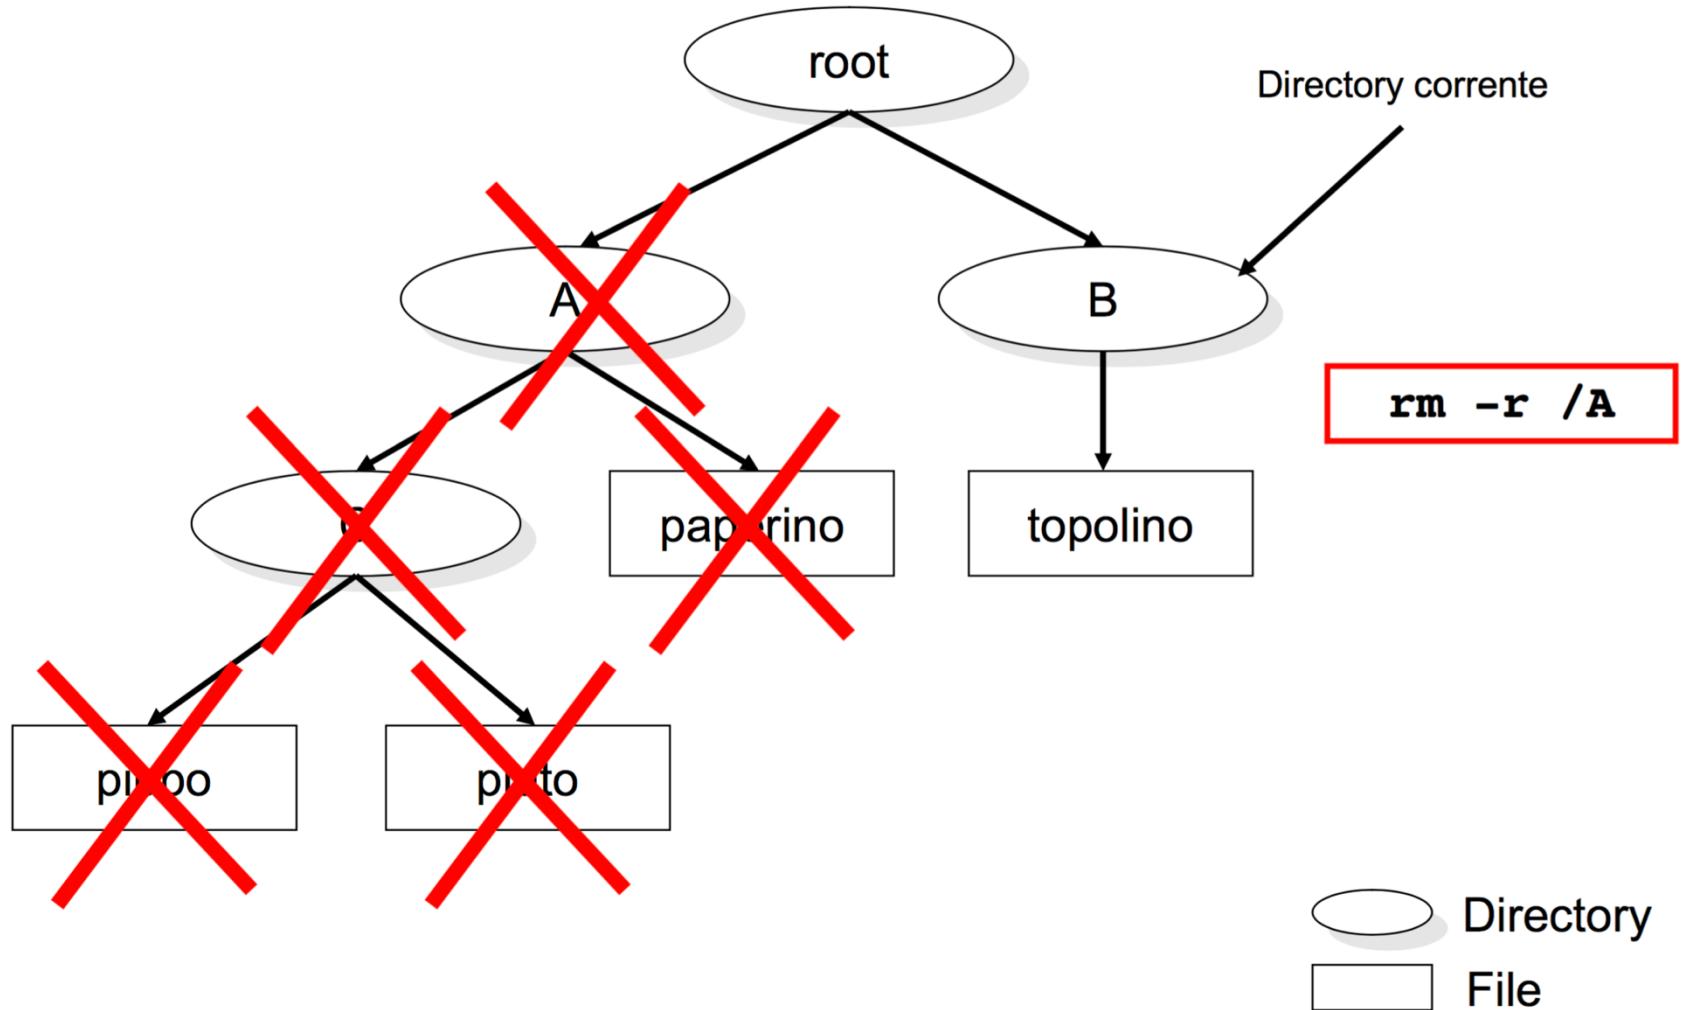
- remove **empty** directories
- The **-p** option tries to remove all the directories specified by the pathname,
 - e.g., `rmdir -p a/c` is equivalent to `rmdir a/c a`



Removing files



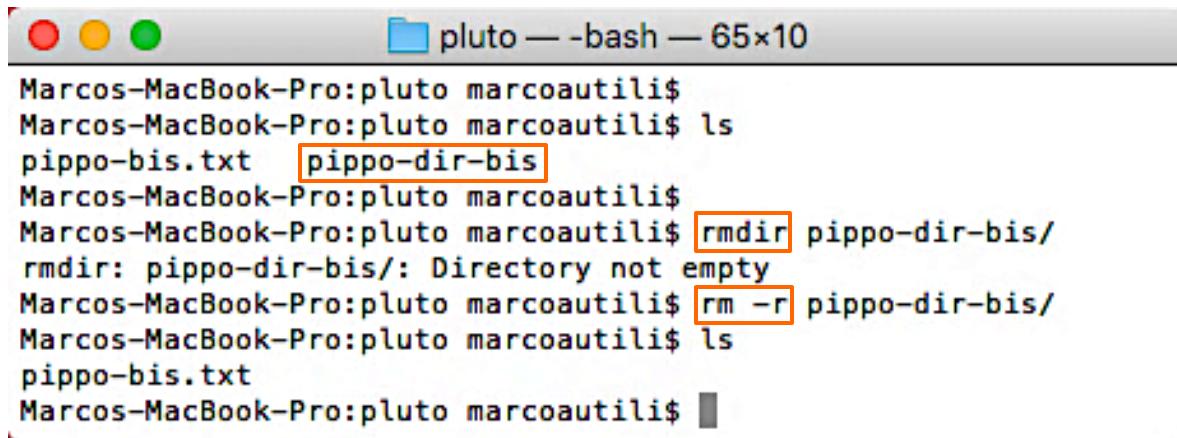
Removing not empty directories



rmdir versus rm

rm [options] <file>

- many many options... see the dedicated man page
- similar to cp for what concerns the -r option



```
Marcos-MacBook-Pro:pluto marcoautili$  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
pippo-bis.txt pippo-dir-bis  
Marcos-MacBook-Pro:pluto marcoautili$ rmdir pippo-dir-bis/  
rmdir: pippo-dir-bis/: Directory not empty  
Marcos-MacBook-Pro:pluto marcoautili$ rm -r pippo-dir-bis/  
Marcos-MacBook-Pro:pluto marcoautili$ ls  
pippo-bis.txt  
Marcos-MacBook-Pro:pluto marcoautili$
```

Removing

- Be careful when using `rm -r *`
 - a good practice would be to use the option `i` (i.e., interactive) together with `r`
 - This option will prompt you before removing each file and directory and give you the option to

More on
wildcards later

```
Marcos-MacBook-Pro:pluto marcoautili$ ls
```

```
pippo-bis.txt pippo-dir-bis
```

```
Marcos-MacBook-Pro:pluto marcoautili$ ls pippo-dir-bis/  
paperino.txt
```

```
Marcos-MacBook-Pro:pluto marcoautili$ rm -ri pippo-dir-bis/  
examine files in directory pippo-dir-bis/? y
```

```
remove pippo-dir-bis//paperino.txt? n
```

```
remove pippo-dir-bis/? y
```

```
rm: pippo-dir-bis/: Directory not empty
```

```
Marcos-MacBook-Pro:pluto marcoautili$
```

To mitigate the "destructive" nature of
"rm", the system can be configured in
such a way that "rm" is aliased to "rm -i"

See next slide...

```
Marcos-MacBook-Pro:pluto marcoautili$ ls
```

```
pippo-bis.txt pippo-dir-bis
```

```
Marcos-MacBook-Pro:pluto marcoautili$
```

```
Marcos-MacBook-Pro:pluto marcoautili$
```

```
Marcos-MacBook-Pro:pluto marcoautili$ rm -ri pippo-dir-bis/  
examine files in directory pippo-dir-bis/? y
```

```
remove pippo-dir-bis//paperino.txt? y
```

```
remove pippo-dir-bis/? y
```

```
Marcos-MacBook-Pro:pluto marcoautili$
```

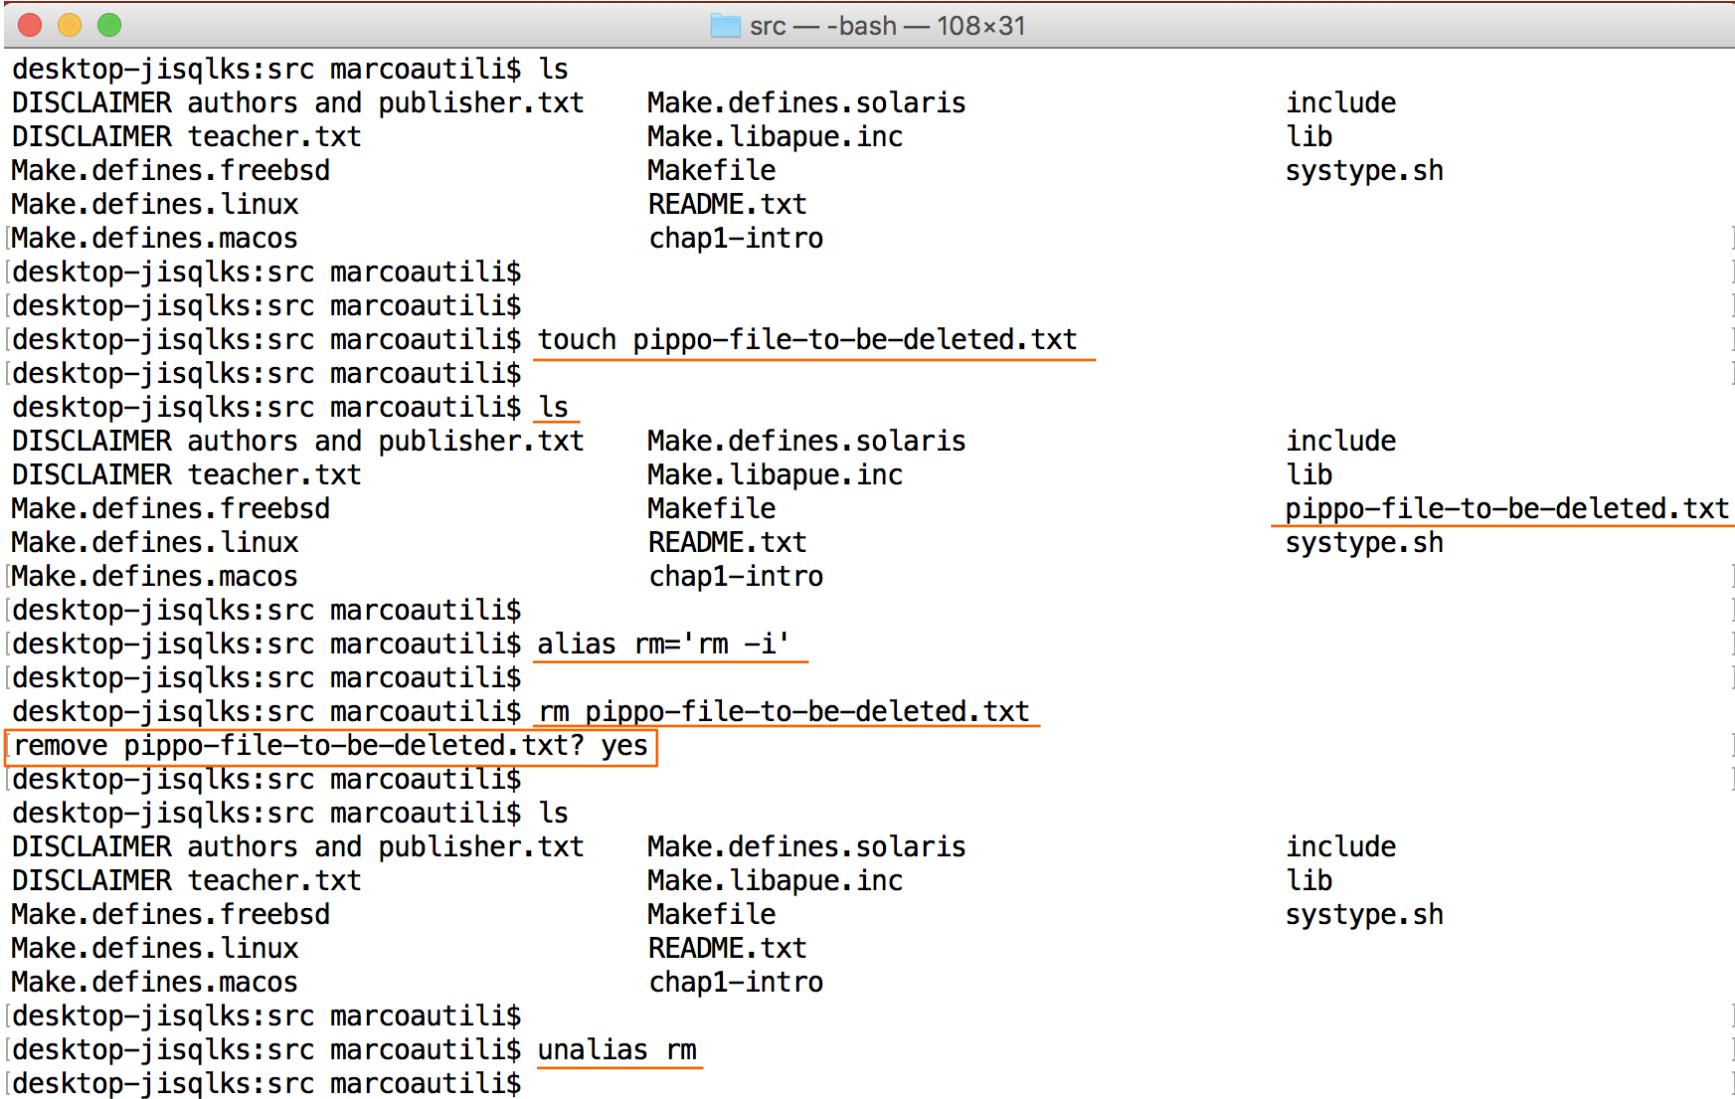
```
Marcos-MacBook-Pro:pluto marcoautili$
```

```
Marcos-MacBook-Pro:pluto marcoautili$ ls
```

```
pippo-bis.txt
```

```
Marcos-MacBook-Pro:pluto marcoautili$
```

Command aliasing

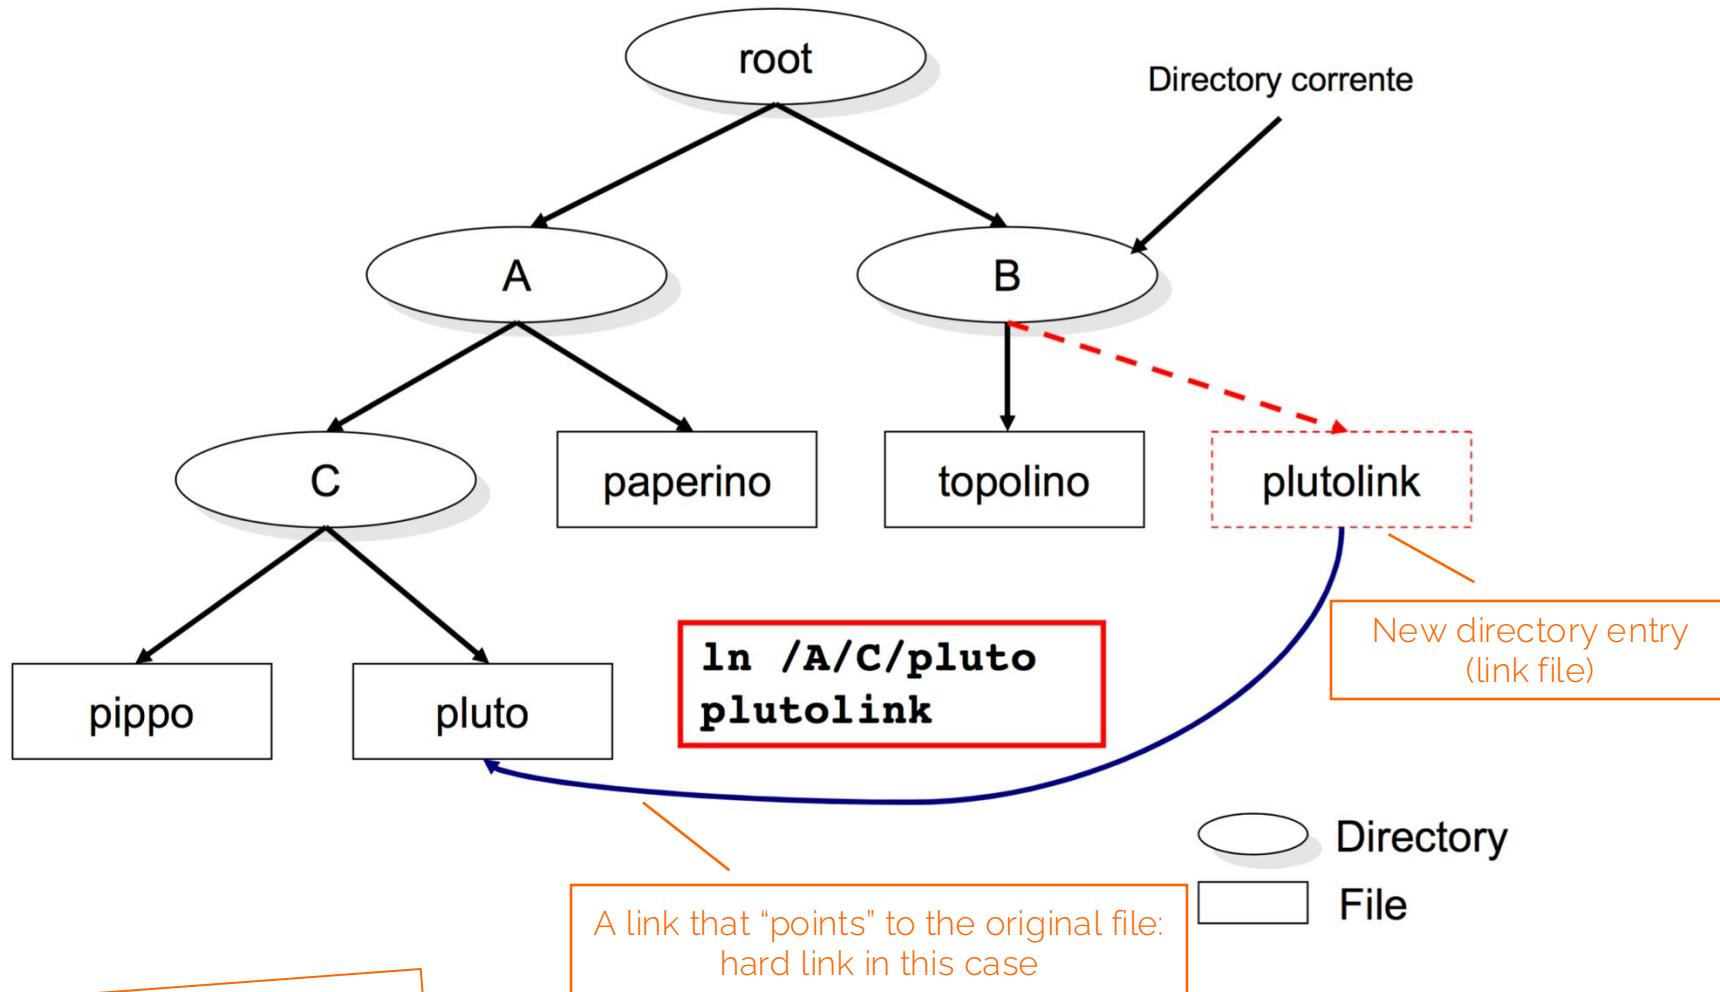


The screenshot shows a terminal window titled "src — -bash — 108x31". The terminal displays a series of commands demonstrating command aliasing:

```
desktop-jisqlks:src marcoautili$ ls
DISCLAIMER authors and publisher.txt
DISCLAIMER teacher.txt
Make.defines.freebsd
Make.defines.linux
Make.defines.macos
[desktop-jisqlks:src marcoautili$
[desktop-jisqlks:src marcoautili$
[desktop-jisqlks:src marcoautili$ touch pippo-file-to-be-deleted.txt
[desktop-jisqlks:src marcoautili$ ls
DISCLAIMER authors and publisher.txt
DISCLAIMER teacher.txt
Make.defines.freebsd
Make.defines.linux
Make.defines.macos
[desktop-jisqlks:src marcoautili$
[desktop-jisqlks:src marcoautili$ alias rm='rm -i'
[desktop-jisqlks:src marcoautili$ rm pippo-file-to-be-deleted.txt
remove pippo-file-to-be-deleted.txt? yes
[desktop-jisqlks:src marcoautili$
[desktop-jisqlks:src marcoautili$ ls
DISCLAIMER authors and publisher.txt
DISCLAIMER teacher.txt
Make.defines.freebsd
Make.defines.linux
Make.defines.macos
[desktop-jisqlks:src marcoautili$
[desktop-jisqlks:src marcoautili$ unalias rm
[desktop-jisqlks:src marcoautili$
```

The terminal shows several examples of command aliasing. The first example creates an alias "rm" which is equivalent to "rm -i". The second example demonstrates the use of this alias by attempting to delete the file "pippo-file-to-be-deleted.txt". A confirmation dialog is shown asking if the user wants to remove the file. The third example shows the removal of the alias "rm". The terminal also lists several files in the current directory, including "Make.defines.solaris", "Make.libapue.inc", "Makefile", "README.txt", and "chap1-intro".

(Hard) Link



Cat and Less commands

- `cat <file>` (stands for concatenate)
 - Its main purpose if to concatenate files together
 - It can be conveniently used for viewing files
 - However, it is appropriate to view small files
- `less <file>`
 - more appropriate for larger files
 - arrow keys can be used to move up and down within a file
 - `SpaceBar` to move forward a whole page
 - `b` to move back a whole page
 - `q` quit the file
- see also `more <file>`

Homework Activities

- Use the commands we have learnt so far to create, delete, copy, move, rename files and directories
- Navigate the file system and check the created directories and files
- Be careful with commands without undo features
 - i.e., be careful when performing destructive actions... command aliasing can be of help!

Editing files

Editing a file is to modify the content of a file

- Text editor
 - Enter and modify text in a text file
- Word processor
 - Enter, modify and format text in a document
- Line editor
 - Edit file one line at a time
 - Unix examples: ex, ed and sed
- Full screen editor
 - Shows a whole screen of text at a time

The following slides on editors will not be part of the exam

Editor features

- enter text
 - search and replace
 - copy, cut and paste
 - undo and redo
-
- importing and exporting text
 - save and cancel

Text files

- As you already know, Unix file name **does not require file extension**
 - Unix file system **does not consider the extension** when treating files
- However, some extensions are commonly used
 - Program source code: **.c .cc .cpp .f .f77 .f95**
 - Compiled object code: **.o .a .so .sa**
 - Compressed files: **.z .gz .zip**
 - Archive files: **.tar .tz**
 - Web site source code: **.html .shtml .php**
 - Notably, executable files typically have no extension

Unix Text editors

- vi
 - vim (vi improved)
 - emacs
 - pico
 - nano
-
- GUI editors
 - emacs
 - mousepad
 - Xedit

History on wikipedia...

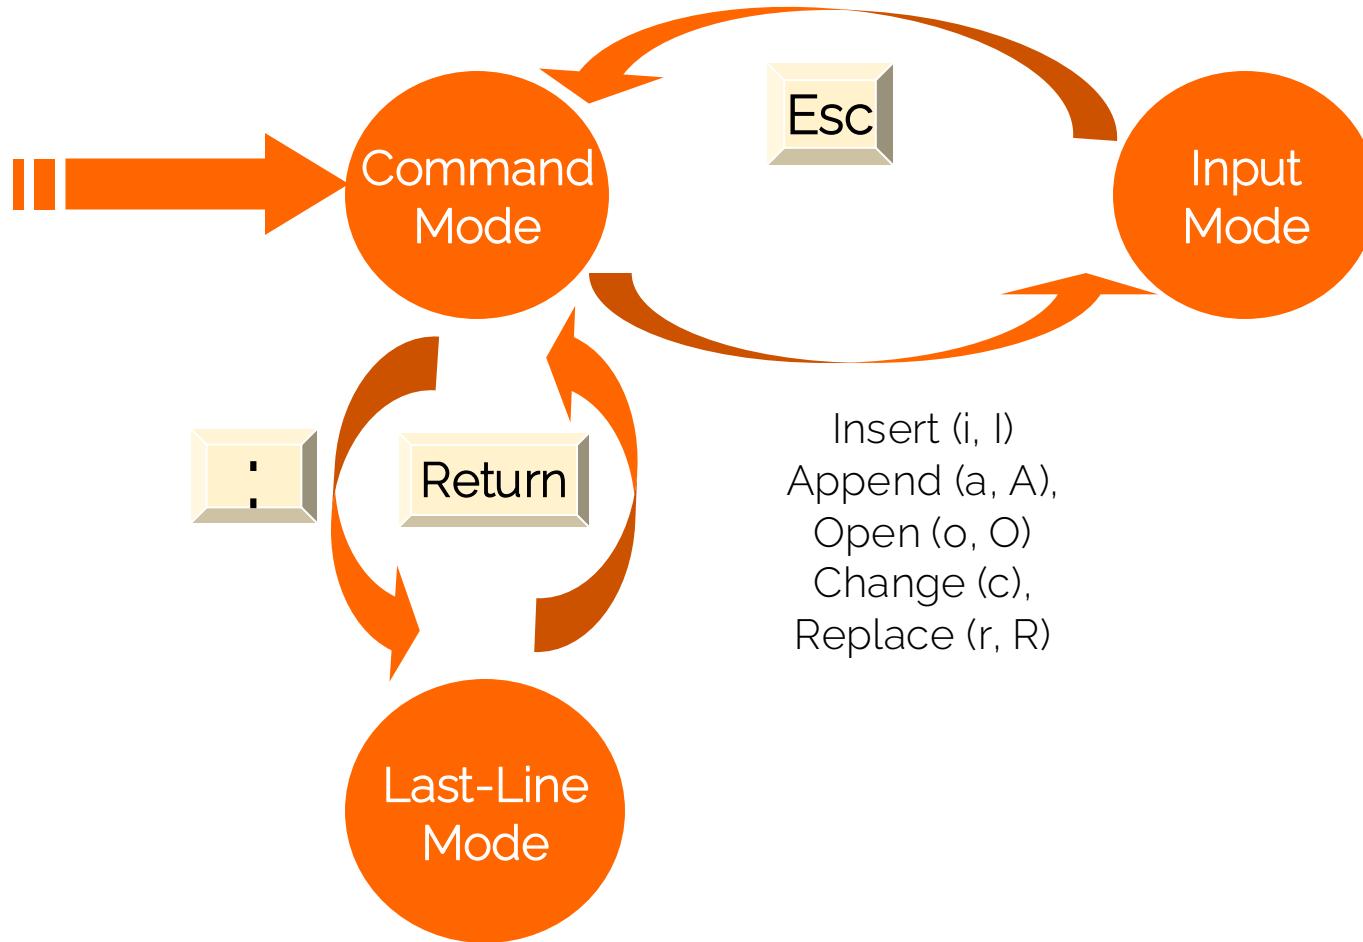
vi Text Editor

- A **command line** text editor
- A **plain text** editor
- Not a word processor
- Learning it might be "**a bother**"
- **Useful** and **powerful**
 - (more than Notepad on Window or Textedit on Mac)
- No GUI... forget the mouse

- vi has multiple modes of operation:
 - input mode, command mode, last-line mode

Let's start with the basics ...

vi Editing modes



vi commands

- `:w` write the current buffer (**save**)
- `:q!` quit without saving
- `:wq` write out (save) and quit
- `:x` write out (save) and quit
- `ZZ` write out (save) and quit
- `F1` help
- `:help`
- `:help <command>`
 - `:q` to exit help window

vi commands

- Insert characters
 - **i** converts to insert mode
 - then type characters
 - **<esc>** to exit insert mode
- Delete characters
 - **x** deletes character under the cursor
- Insert lines
 - **o** open line below cursor
 - **O** open line above cursor
 - **<esc>** to exit insert mode
- Append characters
 - **a** converts to insert mode after the cursor
 - **A** converts to insert mode at the end of a line
 - **<esc>** to exit insert mode

vi commands

- Using motions for movement
 - **sx, dx, up, down** motions to move
 - **<n>w** move cursor n words forward
 - **2w** move cursor two words forward
 - **3w** move cursor four words forward
- Deletion
 - **d\$** delete until the end of line
 - **dw** delete until the beginning of next word
 - **de** delete until the end of current word
 - **d + motions** (sx, dx, up, down)
 - **d<n> + motions**
- Using repetition as part of deletion
 - **2dw** delete next two words (better to try)
- Deleting a line
 - **dd** delete line
 - **2dd** delete two lines
- Undo
 - **u** undo one command
 - **U** restore a line
 - **ctrl-R** redo a command

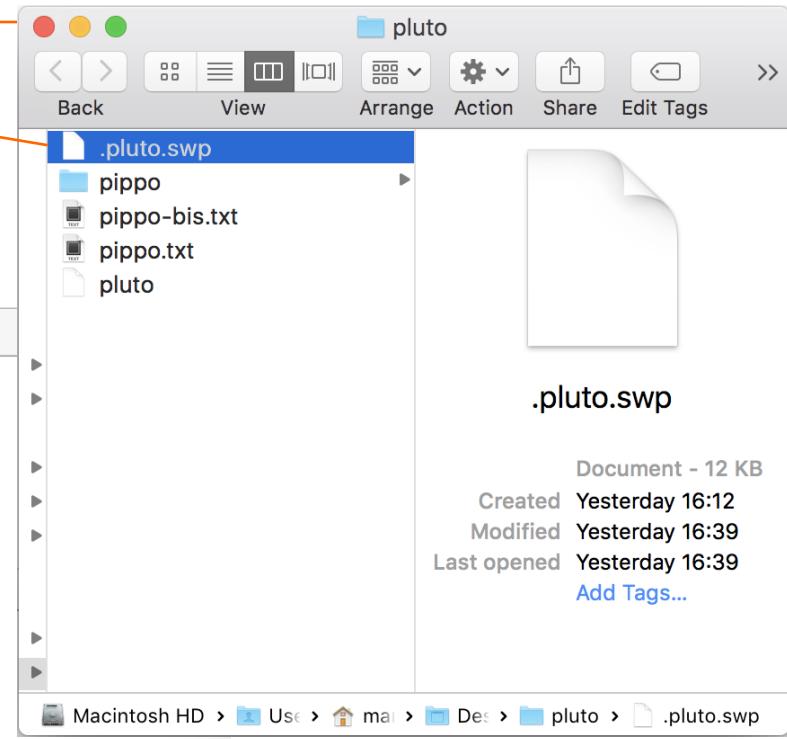
State files and recovery

```
E325: ATTENTION
Found a swap file by the name ".pluto.swp"
    owned by: marcoautili    dated: Tue Sep 13 16:39:17 2016
    file name: ~marcoautili/Desktop/pluto/pluto
    modified: YES
    user name: marcoautili    host name: iMac-2.local
    process ID: 1129
While opening file "pluto"
    dated: Tue Sep 13 16:07:39 2016

(1) Another program may be editing the same file.
If this is the case, be careful not to end up with two
different instances of the same file when making changes.
Quit, or continue with caution.

(2) An edit session for this file crashed.
If this is the case, use ":recover" or "vim -r pluto"
to recover the changes (see ":help recovery").
If you did this already, delete the swap file ".pluto.swp"
to avoid this message.

Swap file ".pluto.swp" already exists!
[O]pen Read-Only, (E)dit anyway, (R)ecover, (D)elete it, (Q)uit, (A)bort:
```



Read next slide
and try yourself ...

More on state files and recovery

- When mentioning a “swap file”, the first thing that comes into your mind might be the kind of file that you would create to increase the swap space on a Unix-like system
- The `.swp` file in previous slide is not a swap file in the OS sense, it is a **state file** that keeps your changes since the last save (except the last 200 characters), buffers that you have saved, unsaved macros and the undo structure
- Read more in VIM's help: **`vim +help\ swap-file`**
- If there is a crash (power failure, OS crash, etc.), then you can recover your changes using this swap-file. After saving the changes from the swap file to the original file, you will need to exit vim and remove the swap file yourself

<http://www.computerworld.com/article/2931534/it-management/what-are-unix-swap-swp-files.html>

vi commands

- **p** put back the deleted text (in new place)
 - one of the **delete command above + put = cut-and-paste**
- More general cut-and-paste
 - **v** start visual mode per character
 - move the cursor to highlight and select
 - **V** start visual mode line wise
 - **y** yank (copy to buffer then **p** put in new place)
- Search
 - **/<phrase>** search
 - **/<phrase>\c** ignore case
 - **?<phrase>** search backwards
 - **n** repeat search
 - **N** repeat search in the other direction
- Screen movement
 - **ctl-b** move one screen backward
 - **ctl-f** move one screen forward
 - **ctl-u** move half a screen backward
 - **ctl-d** move half a screen forward

vi commands

- Search for matching parentheses
 - simply put the cursor on (, [, { or], },)
 - % switch between matching left and matching right
- Substitute (replace)
 - :s/thee/the changes first one
 - :s/thee/the/g changes all (global change)
 - :s/thee/the/gc change all with query
 - :#, #s/thee/the/g only change within that line range
- Setting options for search and viewing
 - :set nu enable line numbers
 - :set ic ignore case
 - :set hlsearch highlight matches
 - :set icsearch incremental search
 - :set noic, etc. turn the option off

File navigation commands

- **arrow keys** move the cursor around
- **j, k, h, l** move the cursor down, up, left and right (similar to the arrow keys)
- **^** the caret move the cursor to beginning of current line
- **\$** move the cursor to end of the current line
- **<n>G** move to the **nth** line (e.g., 5G moves to 5th line)
- **G** move to the last line
- **gg** go to top of file
- **w** move to the beginning of the next word
- **<n>w** move forward n word (e.g., 2w moves two words forwards)
- **b** move to the beginning of the previous word
- **<n>b** move back n word
- **[** move backward one paragraph
- **]** move forward one paragraph
- **ctrl-g** show position in file

Running Commands

- The vi has the capability to run commands from within the editor
- To run a command, you only need to go into command mode and type :!<any command>
- For example, if you want to check whether a file exists before you try to save your file to that filename, you can type :!ls and you will see the output of ls on the screen

The emacs editor

- Originally started as editor macros in 1976
- Gosling Emacs available for Unix in 1981
- GNU Emacs created by Richard Stallman in 1984
 - very popular editor on Unix until recently
 - history: editor war: emacs vs. vi
- Uses lisp-like macro language for powerful features and extensions:
 - programming language sensitive editing
 - email client
 - news reader
- Has built-in tutorial: ^h-t

The Pico and Nano editors

- Part of the popular pine mail utility on UNIX
 - developed by the University of Washington
- pico = pine email composer
- nano is improved open source of pico available for GNU/Linux
 - very intuitive operation
 - on-screen guide and help

Homework Activities

- Play with all the vi commands
- Search the web for more
- Play with the :help command

Wildcards

- Allow the definition of patterns (also known as **globbing patterns** or **globbing pathnames**)
 - a means to deal with a set of files (or directories) at once
- Standard wildcards
 - ***** match zero or more characters
 - **?** match a single character
 - **[]** match a set or a range of characters
 - **[agd]** - the character is one of those included within the square brackets
 - **[a-d]** - the character is either **a**, **b**, **c** or **d**
 - **[^a-d]** or **![a-d]** - reverse a range, the character is none of **a**, **b**, **c**, **d**
 - **{}** terms are separated by commas and each term must be the name of something or a wildcard (an "or" relationship, one or the other)
 - **cp {*.doc,* .pdf,pippo.txt,mydirectory}** - this wildcard will copy anything that matches either wildcard(s) ***.doc** or *** .pdf**, or the exact name(s) **pippo.txt** or **mydirectory**
 - **[a-d]** is equivalent to **[a,b,c,d]**
 - **** (backslash) is used as an "escape" character, i.e., to protect a subsequent special character, thus **\\"** matches a backslash ****
 - Note that you may need to use quotation marks and backslash(es) **"\\\"\\"**

Wildcards

- `ls ?ip*`
- `ls ??pp*`
- `ls ?ip*.txt`
- `ls *.???`
- `ls ?ip*.txt`
- `ls [psa]*.txt`
- `ls [0-9]` using the hyphen character it ranges over the set {0, 1, 2, ..., 9}
- `ls [^a-k]*` reverse a range
- `ls -lh /home/*/*.txt`
- `ls -lh ./*/*.txt` can be used anywhere within a path

IMPORTANT

wildcards are translated by the shell not by commands
`ls *.txt` will be translated into `ls pippo-bis.txt pippo.txt`

```
iMac-2:pluto marcoautili$  
[iMac-2:pluto marcoautili$ ls  
[pippo      pippo-bis.txt  pippo.txt      pluto      pluto copy      pluto copy 2  
[iMac-2:pluto marcoautili$  
[iMac-2:pluto marcoautili$  
[iMac-2:pluto marcoautili$ ls p*  
pippo-bis.txt  pippo.txt      pluto      pluto copy      pluto copy 2  
  
pippo:  
prova.txt  
[iMac-2:pluto marcoautili$ ls *.txt  
pippo-bis.txt  pippo.txt  
[iMac-2:pluto marcoautili$
```

Wildcards

- file ./pluto/*

The terminal window shows the command `file ./pluto/*` being run, and the output indicates the following file types:

- ./pluto/pippo: directory
- ./pluto/pippo-bis.txt: empty
- ./pluto/pippo.txt: ASCII text
- ./pluto/pluto: ASCII text
- ./pluto/pluto copy: ASCII text
- ./pluto/pluto copy 2: ASCII text

The Finder window shows a folder named "pluto" containing the following files:

- pippo
- pippo-bis.txt
- pippo.txt
- pluto
- pluto copy
- pluto copy 2

The "pippo-bis.txt" file is selected in the Finder.

- mv p*[y2] ./pippo

The terminal window shows the command `ls` being run, and the output lists the following files:

- pippo
- pippo-bis.txt
- pippo.txt
- pluto
- pluto copy
- pluto copy 2

The terminal then shows the command `mv p*[y2] ./pippo` being run, followed by another `ls` command which shows the contents of the "pippo" directory:

- pippo
- pippo-bis.txt
- pippo.txt
- pluto

The "pippo-bis.txt" file is selected in the Finder.

- etc etc ... try yourself ☺

More on globbing

- Long ago, in UNIX V6, there was a program `/etc/glob` to expand wildcard patterns
- Soon afterward, this became a shell built-in
 - On my Mac, just type `man 1 glob` to go to the builtin commands page
 - There is also a C Library Routine `glob(3)` that will perform this function for a user program
`(#include <glob.h>)`
- According to the standard POSIX.2, 3.13, the rules for wildcard matching are as follows:
 - A string is a wildcard pattern if it contains one of the characters '`?`', '`*`' or '`[`'
- That said, `globbing` can be defined as the operation that expands a wildcard pattern into the list of pathnames matching the pattern

Let's now open a parenthesis
on builtin commands...

Builtin commands

ppp — less — man builtin — 102x47

BUILTIN(1) BSD General Commands Manual BUILTIN(1)

NAME

`builtin`, `!`, `%`, `.`, `:`, `@`, `{`, `}`, `alias`, `alloc`, `bg`, `bind`, `bindkey`, `break`, `breaksw`, `builtins`, `case`, `cd`, `chdir`, `command`, `complete`, `continue`, `default`, `dirs`, `do`, `done`, `echo`, `echotc`, `elif`, `else`, `end`, `endif`, `endsw`, `esac`, `eval`, `exec`, `exit`, `export`, `false`, `fc`, `fg`, `filetest`, `fi`, `for`, `foreach`, `getopts`, `glob`, `goto`, `hash`, `hashstat`, `history`, `hup`, `if`, `jobid`, `jobs`, `kill`, `limit`, `local`, `log`, `login`, `logout`, `ls-F`, `nice`, `nohup`, `notify`, `onintr`, `popd`, `printenv`, `pushd`, `pwd`, `read`, `readonly`, `rehash`, `repeat`, `return`, `sched`, `set`, `setenv`, `settc`, `setty`, `setvar`, `shift`, `source`, `stop`, `suspend`, `switch`, `telltc`, `test`, `then`, `time`, `times`, `trap`, `true`, `type`, `ulimit`, `umask`, `unalias`, `uncomplete`, `unhash`, `unlimit`, `unset`, `unsetenv`, `until`, `wait`, `where`, `which`, `while` — shell built-in commands

SYNOPSIS

`builtin [-options] [args ...]`

DESCRIPTION

Shell builtin commands are commands that can be executed within the running shell's process. Note that, in the case of csh(1) builtin commands, the command is executed in a subshell if it occurs as any component of a pipeline except the last.

If a command specified to the shell contains a slash `'/'`, the shell will not execute a builtin command, even if the last component of the specified command matches the name of a builtin command. Thus, while specifying `'echo'` causes a builtin command to be executed under shells that support the `echo` builtin command, specifying `'/bin/echo'` or `'./echo'` does not.

While some builtin commands may exist in more than one shell, their operation may be different under each shell which supports them. Below is a table which lists shell builtin commands, the standard shells that support them and whether they exist as standalone utilities.

Only builtin commands for the csh(1) and sh(1) shells are listed here. Consult a shell's manual page for details on the operation of its builtin commands. Beware that the sh(1) manual page, at least, calls some of these commands `'built-in commands'` and some of them `'reserved words'`. Users of other shells may need to consult an info(1) page or other sources of documentation.

Builtin commands

Commands marked ``No*'') under External do exist externally, but are implemented as scripts using a builtin command of the same name.

Command	External	csh(1)	sh(1)
!	No	No	Yes
%	No	Yes	No
.	No	No	Yes
:	No	Yes	Yes
@	No	Yes	Yes
{	No	No	Yes
}	No	No	Yes
alias	No**	Yes	Yes
alloc	No	Yes	No
bg	No**	Yes	Yes
bind	No	No	Yes
bindkey	No	Yes	No
break	No	Yes	Yes
breaksw	No	Yes	No
builtin	No	No	Yes
builtins	No	Yes	No
case	No	Yes	Yes
cd	No**	Yes	Yes
chdir	No	Yes	Yes
command	No**	No	Yes
complete	No	Yes	No
continue	No	Yes	Yes
default	No	Yes	No
dirs	No	Yes	No
do	No	No	Yes
done	No	No	Yes
echo	Yes	Yes	Yes
echotc	No	Yes	No
elif	No	No	Yes

Note that here I'm running zsh
and these are predefined aliases

```
marcoautili@iMac ~ %
marcoautili@iMac ~ % alias
run-help=man
which-command=whence
marcoautili@iMac ~ %
marcoautili@iMac ~ % run-help ls
marcoautili@iMac ~ %

.
.

.
.

[marcoautili@iMac ~ %
[marcoautili@iMac ~ % whence ls
/bin/ls
[marcoautili@iMac ~ % whence cd
cd
[marcoautili@iMac ~ % whence alias
alias
```

man page for builtin commands

man bash (then, search for, e.g., /BUILTIN COMMANDS)

```
event line. An a may be used as a synonym for g.  
G Apply the following 's' modifier once to each word in the event line.  
  
SHELL BUILTIN COMMANDS  
Unless otherwise noted, each builtin command documented in this section as accepting options preceded by - accepts -- to signify the end of the options. For example, the :, true, false, and test builtins do not accept options.  
: [arguments]  
No effect; the command does nothing beyond expanding arguments and performing any specified redirections. A zero exit code is returned.  
  
. filename [arguments]  
source filename [arguments]  
Read and execute commands from filename in the current shell environment and return the exit status of the last command executed from filename. If filename does not contain a slash, file names in PATH are used to find the directory containing filename. The file searched for in PATH need not be executable. When bash is not in posix mode, the current directory is searched if no file is found in PATH. If the sourcepath option to the shopt builtin command is turned off, the PATH is not searched. If any arguments are supplied, they become the positional parameters when filename is executed. Otherwise the positional parameters are unchanged. The return status is the status of the last command exited within the script (0 if no commands are executed), and false if filename is not found or cannot be read.  
  
alias [-p] [name[=value] ...]  
Alias with no arguments or with the -p option prints the list of aliases in the form alias name=value on standard output. When arguments are supplied, an alias is defined for each name whose value is given. A trailing space in value causes the next word to be checked for alias substitution when the alias is expanded. For each name in the argument list for which no value is supplied, the name and value of the alias is printed. Alias returns true unless a name is given for which no alias has been defined.
```

/BUILTIN COMMANDS

Builtin commands

- Shell builtins work significantly faster than external programs, because there is no program loading overhead
- However, their code is inherently present in the shell, and thus modifying or updating them requires modifications to the shell
 - for these reason, shell builtins are usually used for simple, almost trivial, functions, such as text output
- Because of the nature of some operating systems, some functions of the systems must necessarily be implemented as shell builtins
 - The most notable example is the cd command, which changes the working directory of the shell program itself
 - Since each executable program runs in a separate process, and working directories are specific to each process, loading cd as an external program would not affect the working directory of the shell that loaded it!
- Use the type or command to find out if given command is an internal builtin or external
 - type -a <command-name>
 - command -V <command-name>

https://en.wikipedia.org/wiki/Shell_builtin

Builtin commands

iMac-di-User:~ marcoautili\$

iMac-di-User:~ marcoautili\$ type -a top top -- display and update sorted information about processes (q to quit)

top is /usr/bin/top

iMac-di-User:~ marcoautili\$ type -a alias alias -- see previous slides

alias is a shell builtin

alias is /usr/bin/alias

iMac-di-User:~ marcoautili\$ command -V ls

ls is /bin/ls

iMac-di-User:~ marcoautili\$ type -a ls

ls is /bin/ls

iMac-di-User:~ marcoautili\$ commamd -V date

-bash: commamd: command not found

iMac-di-User:~ marcoautili\$ type -a date

date is /bin/date

iMac-di-User:~ marcoautili\$ type -a vim

vim is /usr/bin/vim

iMac-di-User:~ marcoautili\$ command -V vim

vim is /usr/bin/vim

iMac-di-User:~ marcoautili\$

iMac-di-User:~ marcoautili\$ man builtin

iMac-di-User:~ marcoautili\$

iMac-di-User:~ marcoautili\$

iMac-di-User:~ marcoautili\$

iMac-di-User:~ marcoautili\$

iMac-di-User:~ marcoautili\$



marcoautili — bash —

● ○ ■

● ○ ■

marcoautili — less < man builtin — 97x51

BUILTIN(1)

BSD General Commands Manual

BUILTIN(1)

NAME

builtin, !, %, ., :, @, {, }, alias, alloc, bg, bind, bindkey, break, breaksw, builtins, case, cd, chdir, command, complete, continue, default, dirs, do, done, echo, echotc, elif, else, end, endif, endsw, esac, eval, exec, exit, export, false, fc, fg, filetest, fi, for, foreach, getopt, glob, goto, hash, hashtstat, history, hup, if, jobid, jobs, kill, limit, local, log, login, logout, ls-F, nice, nohup, notify, onintr, popd, printenv, pushd, pwd, read, readonly, rehash, repeat, return, sched, set, setenv, settc, setty, setvar, shift, source, stop, suspend, switch, telltc, test, then, time, times, trap, true, type, ulimit, umask, unalias, uncomplete, unhash, unlimit, unset, unsetenv, until, wait, where, which, while -- shell built-in commands

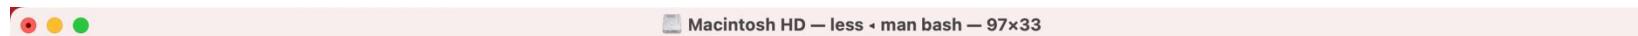
SYNOPSIS

builtin [-options] [args ...]

DESCRIPTION

Shell builtin commands are commands that can be executed within the running shell's process. Note that, in the case of csh(1) builtin commands, the command is executed in a subshell if it occurs as any component of a pipeline except the last.

Back to globbing...



Pathname Expansion

After word splitting, unless the **-f** option has been set, **bash** scans each word for the characters *****, **?**, and **[**. If one of these characters appears, then the word is regarded as a pattern, and replaced with an alphabetically sorted list of file names matching the pattern. If no matching file names are found, and the shell option **nullglob** is disabled, the word is left unchanged. If the **nullglob** option is set, and no matches are found, the word is removed. If the **failglob** shell option is set, and no matches are found, an error message is printed and the command is not executed. If the shell option **nocaseglob** is enabled, the match is performed without regard to the case of alphabetic characters. When a pattern is used for pathname expansion, the character **``..''** at the start of a name or immediately following a slash must be matched explicitly, unless the shell option **dotglob** is set. When matching a pathname, the slash character must always be matched explicitly. In other cases, the **``..''** character is not treated specially. See the description of **shopt** below under **SHELL BUILTIN COMMANDS** for a description of the **nocaseglob**, **nullglob**, **failglob**, and **dotglob** shell options.

The **GLOBIGNORE** shell variable may be used to restrict the set of file names matching a pattern. If **GLOBIGNORE** is set, each matching file name that also matches one of the patterns in **GLOBIGNORE** is removed from the list of matches. The file names **``..''** and **``...''** are always ignored when **GLOBIGNORE** is set and not null. However, setting **GLOBIGNORE** to a non-null value has the effect of enabling the **dotglob** shell option, so all other file names beginning with a **``..''** will match. To get the old behavior of ignoring file names beginning with a **``..''**, make **``.*''** one of the patterns in **GLOBIGNORE**. The **dotglob** option is disabled when **GLOBIGNORE** is unset.

Pattern Matching

Any character that appears in a pattern, other than the special pattern characters described below, matches itself. The NUL character may not occur in a pattern. A backslash escapes the following character; the escaping backslash is discarded when

Homework Activities

- Play with wildcards and with different commands
- Define increasingly complex patterns also using them in the middle of paths
- For more information on standard wildcards (globbing pathnames, globbing patterns) refer to the manual page:
 - on my Mac **man 1 glob** (to go to the builtin commands page)
 - on my Mac **man 3 glob** (to go to the C Library Functions page)
 - on Linux you should have **man 7 glob**
<http://man7.org/linux/man-pages/man7/glob.7.html>
- On Mac, for **zsh**, try also **man zshbuiltins**

File Permissions

- Allow for **assigning access rights** to specific users and groups of users
- Control the ability of users to view, change, navigate, and/or execute the contents of the file system
- Unix-like and otherwise POSIX-compliant systems, including Linux-based systems and all macOS versions, have a simple system for managing individual file permissions

File Permissions

- **r read** - the contents of the file can be viewed
- **w write** - the contents of the file can be changed
- **x execute** - the program or script can be executed or run
- **owner** - a single person who owns the file
 - (typically, the person who created the file, but ownership may be granted to some one else by certain users)
- **group** - every file belongs to a single group
- **others** - everyone else who is not in the group or the owner

https://en.wikipedia.org/wiki/File_system_permissions

File Permissions

- **chmod [permissions] [path]**

chmod stands for change file mode

- Permissions can be changed for

- user (or owner) -> **u**
- group -> **g**
- others -> **o**
- all -> **a**

- Permission can be

- granted -> **+**
- revoked -> **-**

Multiple permissions, e.g., **wx**,
can be assigned at once to, e.g.,
both user and group **ug**



```
iMac-2:pluto marcoautili$ chmod u-w pippo.txt
iMac-2:pluto marcoautili$ chmod u+wx pippo.txt
iMac-2:pluto marcoautili$ chmod ug+wx pippo.txt
```

Shorthand to set permissions

Octal	Binary
0	000
1	001
2	010
3	011
4	100
5	101
6	110
7	111

```
[iMac-2:pluto marcoautili$ ls -l
total 16
-rw-r--r--@ 1 marcoautili  staff    0 Sep 12 16:26 Pippo-bis.txt
drwxr-xr-x@ 5 marcoautili  staff  170 Sep 14 15:07 pippo
-rw-r--r--@ 1 marcoautili  staff   38 Sep 13 12:26 pippo.txt
-rw-r--r--@ 1 marcoautili  staff  176 Sep 14 11:48 pluto
[iMac-2:pluto marcoautili$ chmod 777 pippo.txt
[iMac-2:pluto marcoautili$ ls -l
total 16
-rw-r--r--@ 1 marcoautili  staff    0 Sep 12 16:26 Pippo-bis.txt
drwxr-xr-x@ 5 marcoautili  staff  170 Sep 14 15:07 pippo
-rwxrwxrwx@ 1 marcoautili  staff   38 Sep 13 12:26 pippo.txt
-rw-r--r--@ 1 marcoautili  staff  176 Sep 14 11:48 pluto
[iMac-2:pluto marcoautili$ chmod 644 pippo.txt
[iMac-2:pluto marcoautili$ ls -l
total 16
-rw-r--r--@ 1 marcoautili  staff    0 Sep 12 16:26 Pippo-bis.txt
drwxr-xr-x@ 5 marcoautili  staff  170 Sep 14 15:07 pippo
-rw-r--r--@ 1 marcoautili  staff   38 Sep 13 12:26 pippo.txt
-rw-r--r--@ 1 marcoautili  staff  176 Sep 14 11:48 pluto
iMac-2:pluto marcoautili$ ]
```

Directory Permissions

- Different meaning wrt files
- **r** - the contents of the directory can be read
 - (i.e., the `ls` command)
- **w** - the directory can be written
 - (i.e., it is possible to create files and directories in it)
- **x** - the directory can be entered
 - (i.e., `cd`)

```
iMac-2:pluto marcoautili$ ls -l
total 16
-rw-r--r--@ 1 marcoautili  staff   0 Sep 12 16:26 Pippo-bis.txt
-rw-r--r--@ 1 marcoautili  staff   38 Sep 13 12:26 pippo.txt
drwxr-xr-x@ 5 marcoautili  staff  170 Sep 14 15:07 pippoDIR
-rw-r--r--@ 1 marcoautili  staff  176 Sep 14 11:48 pluto
iMac-2:pluto marcoautili$
iMac-2:pluto marcoautili$ chmod 400 pippoDIR/
iMac-2:pluto marcoautili$
iMac-2:pluto marcoautili$ ls -l
total 16
-rw-r--r--@ 1 marcoautili  staff   0 Sep 12 16:26 Pippo-bis.txt
-rw-r--r--@ 1 marcoautili  staff   38 Sep 13 12:26 pippo.txt
dr-----@ 5 marcoautili  staff  170 Sep 14 15:07 pippoDIR
-rw-r--r--@ 1 marcoautili  staff  176 Sep 14 11:48 pluto
iMac-2:pluto marcoautili$ ls -ld pippoDIR/
dr-----@ 5 marcoautili  staff  170 Sep 14 15:07 pippoDIR/
iMac-2:pluto marcoautili$ cd pippoDIR/
-bash: cd: pippoDIR/: Permission denied
```

read permission granted

execute permission denied

Note that the **-d option** (standing for directory) permits to **list a directory** as a **regular file**, rather than its content (as usual). Thus, the directory is not searched recursively.

Extended attributes and ACL

Extended attributes (metadata)

- Extended file attributes are file system features that enable users to associate files with **metadata not interpreted by the filesystem**, whereas regular attributes have a purpose strictly defined by the filesystem (such as permissions or records of creation and modification times)
- Typical uses include storing the author of a document, the character encoding of a plain-text document, or a checksum, cryptographic hash or digital certificate, and discretionary access control information

Access Control List (security)

- An access-control list (ACL), with respect to a computer file system, is a list of permissions attached to **an object**
- An ACL **specifies which users or system processes are granted access to objects**, as well as **what operations are allowed on given objects**
- Each entry in a typical ACL specifies a subject and an operation. For instance, if a file object has an ACL that contains (**Alice: read,write; Bob: read**), this would give Alice permission to read and write the file and Bob to only read it

https://en.wikipedia.org/wiki/Extended_file_attributes

https://en.wikipedia.org/wiki/Access-control_list

Extended attributes and ACLs

xattr -l <some file>

xattr -c <some file>

Used to display, modify or remove the **extended attributes** of one or more files, including directories and symbolic links

chmod -N <some file>

Used to remove file ACL(s)

```
iMac-2:/ marcoautili$ pwd
/
iMac-2:/ marcoautili$ ls -l
total 45
drwxrwxr-x+ 76 root admin 2584 Sep  5 12:58 Applications
drwxr-xr-x+ 67 root wheel 2278 Dec 18 2015 Library
drwxr-xr-x@ 2 root wheel 68 Aug 24 2015 Network
drwxr-xr-x@ 4 root wheel 136 Sep  5 13:01 System
drwxr-xr-x  7 root admin 238 Dec 17 2015 Users
drwxrwxrwt@ 4 root admin 136 Sep  8 16:18 Volumes
drwxr-xr-x@ 39 root wheel 1326 Aug 25 02:03 bin
drwxrwxr-t@ 2 root admin 68 Aug 24 2015 cores
dr-xr-xr-x  3 root wheel 4173 Sep  6 10:17 dev
lrwxr-xr-x@ 1 root wheel 11 Oct 18 2015 etc -> private/etc
dr-xr-xr-x  2 root wheel 1 Sep  9 14:49 home
-rw-r--r--@ 1 root wheel 313 Aug 23 2015 installer.failurerequests
dr-xr-xr-x  2 root wheel 1 Sep  9 14:49 net
drwxr-xr-x  3 root wheel 102 Dec  2 2015 opt
drwxr-xr-x@ 6 root wheel 204 Oct 18 2015 private
drwxr-xr-x@ 59 root wheel 2006 Aug 25 02:03 sbin
lrwxr-xr-x@ 1 root wheel 11 Oct 18 2015 tmp -> private/tmp
drwxr-xr-x@ 12 root wheel 408 Dec 23 2015 usr
lrwxr-xr-x@ 1 root wheel 11 Oct 18 2015 var -> private/var
iMac-2:/ marcoautili$
```

More on that in next slides...

https://en.wikipedia.org/wiki/Extended_file_attributes

https://en.wikipedia.org/wiki/Access-control_list

Extended attributes for comments

myMovie.mov Info

myMovie.mov 3,8 MB
Modified: 7 March 2018 at 21:26

Add Tags...

General:

- Kind: QuickTime movie
- Size: 3.845.518 bytes (4,2 MB on disk)
- Where: Macintosh HD ▸ Users ▸ marcoautili
- Created: 7 March 2018
- Modified: 7 March 2018
- Stationery pad
- Locked

More Info:

- Dimensions: 1080x720
- Codecs: AAC, H.264
- Colour profile: HD (1-1)
- Duration: 00:08
- Audio channels: Mono

Name & Extension:

myMovie.mov

Hide extension

Comments:

[REDACTED]

Open with: QuickTime Player.app

Use this application to open all documents like this one.

Change All...

Preview:

Sharing & Permissions:

Name	Privilege
marcoautili (Me)	Read & Write
staff	Read only
everyone	Read only

ls -ls myMovie.mov

```
iMac:~ marcoautili$ ls -ls myMovie.mov
8216 -rw-r--r--@1 marcoautili staff 3845518 Mar 7 2018 myMovie.mov
iMac:~ marcoautili$
```

**BEFORE:
NO extended attribute**

myMovie.mov Info

myMovie.mov 3,8 MB
Modified: 7 March 2018 at 21:26

Add Tags...

General:

- Kind: QuickTime movie
- Size: 3.845.518 bytes (4,2 MB on disk)
- Where: Macintosh HD ▸ Users ▸ marcoautili
- Created: 7 March 2018 at 21:26
- Modified: 7 March 2018 at 21:26
- Stationery pad
- Locked

More Info:

- Dimensions: 1080x720
- Codecs: AAC, H.264
- Colour profile: HD (1-1)
- Duration: 00:08
- Audio channels: Mono

Name & Extension:

myMovie.mov

Hide extension

Comments:

This text will go into extended attribute!

Open with: QuickTime Player.app (default)

Use this application to open all documents like this one.

Change All...

Preview:

Sharing & Permissions:

Name	Privilege
marcoautili (Me)	Read & Write
staff	Read only
everyone	Read only

ls -ls myMovie.mov

```
iMac:~ marcoautili$ ls -ls myMovie.mov
8216 -rw-r--r--@1 marcoautili staff 3845518 Mar 7 2018 myMovie.mov
iMac:~ marcoautili$
```

**AFTER:
extended attribute**

How to modify file ACLs

```
1 chmod ACL file ...
```

The **chmod** command can be used to modify the Access Control Lists (ACLs) associated with files and directories.

The **chmod** command can also be used to modify the file mode of files and directories: [chmod -- change file modes](#).

If a file/directory has an ACL, the sign "+" will be printed when using the command ["ls -l"](#).

Each file/directory has one ACL, containing an ordered list of entries.

Each entry refers to a user or group, and grants (["allow"](#)) or denies (["deny"](#)) a set of permissions.

In cases where a user and a group exist with the same name, the user or the group name can be prefixed with ["user:"](#) or ["group:"](#).

If the user or group name contains spaces you can use ':' as the delimiter between name and permission.

-e: print the ACL associated with the file, if present, in long (-l) output

I did not test the following examples on my Mac
Try yourself on your system and adjust if needed

Examples

- Add a new ACL entry that grants the permission " `read` ", " `write` ", " `append` ", and " `execute` " to "user1" on the file "file1".

```
1 $ ls -le file1
2 -rwx-----+ 1 mtitek mtitek 5 7 Feb 07:45 file1
3
4 #change ACL permissions
5 $ chmod +a "user:user1 allow read,write,append,execute" file1
6
7 $ ls -le file1
8 -rwx-----+ 1 mtitek mtitek 5 7 Feb 07:45 file1
9 0: user:user1 allow read,write,execute,append
```

The `+a` mode parses a new ACL entry from the next argument on the command line and inserts it into the canonical location in the ACL.

If the supplied entry refers to an identity already listed, the two entries are combined.

- Add a new ACL entry that denies the permission " `write` ", " `append` ", and " `execute` " to "user2" on the file "file1".

```
1 #change ACL permissions
2 $ chmod +a "user:user2 deny write,append,execute" file1
3
4 $ ls -le
5 -rwx-----+ 1 mtitek mtitek 5 7 Feb 07:45 file1
6 0: user:user2 deny write,execute,append
7 1: user:user1 allow read,write,execute,append
```

- Add a new ACL entry, in a specific location, that grants the permission " `read` " to "user3" on the file "file1".

```
1 #change ACL permissions
2 $ chmod +a# 2 "user:user3 allow read" file1
3
4 $ ls -le
5 -rwx-----+ 1 mtitek mtitek 5 7 Feb 07:45 file1
6 0: user:user2 deny write,execute,append
7 1: user:user1 allow read,write,execute,append
8 2: user:user3 allow read
```

When a specific ordering is required, the exact location at which an entry will be inserted is specified with the `+a#` mode.

- Delete an ACL entry on the file "file1".

- Delete deny write permission for "user2" on the file "file1":

```
1 #change ACL permissions
2 $ chmod -a "user:user2 deny write" file1
3
4 #Note that only the "deny write" permission is deleted for "user2"
5 $ ls -le
6 -rwx-----+ 1 mtitek mtitek 4 7 Feb 07:45 file1
7 0: user:user2 deny execute,append
8 1: user:user1 allow read,write,execute,append
9 2: user:user3 allow read
```

The -a mode is used to delete ACL entries. All entries exactly matching the supplied entry will be deleted.

- Delete deny execute and append permissions for "user2" on the file "file1":

```
1 $ chmod -a "user:user2 deny execute,append" file1
2
3 #Note that "user2" has now no ACL entry
4 $ ls -le
5 -rwx-----+ 1 mtitek mtitek 4 7 Feb 07:45 file1
6 0: user:user1 allow read,write,execute,append
7 1: user:user3 allow read
```

If the entry lists a subset of rights granted by an entry, only the rights listed are removed. Entries may also be deleted by index using the -a# mode.

- Delete an ACL entry by its index:

```
1 $ chmod -a# 0 file1
2
3 #Note that "user1" has now no ACL entry
4 $ ls -le
5 -rwx-----+ 1 mtitek mtitek 4 7 Feb 07:45 file1
6 0: user:user3 allow read
```

http://www.mtitek.com/tutorials/linux/fileperm_chmod_acl.php

- Modify an ACL entry by its index.

```
1 #change ACL permissions
2 $ chmod =a# 0 "group:group1 deny write,execute,append" file1
3
4 $ ls -le
5 -rwx-----+ 1 mtitek mtitek 5 7 Feb 07:45 file1
6 0: group:group1 deny write,execute,append
```

- Set the permission "list","search","add_file","add_subdirectory",and "delete_child" to "user1" on the directory "folder1".

```
1 $ mkdir folder1
2
3 $ ls -le
4 drwxr-xr-x 2 mtitek mtitek 68 7 Feb 07:52 folder1
5
6 #change ACL permissions
7 $ chmod +a "user:user1 allow list,search,add_file,add_subdirectory,delete_child" folder1
8
9 $ ls -le
10 drwxr-xr-x+ 2 mtitek mtitek 68 7 Feb 07:52 folder1
11 0: user:user1 allow list,add_file,search,add_subdirectory,delete_child
```

See man chmod

http://www.mtitek.com/tutorials/linux/fileperm_chmod_acl.php

Filters, Redirection and Pipelines

- A filter is a computer program (or subroutine) to process a stream, producing another stream
- In Unix and Unix-like operating systems, a filter is a program that gets its data from its standard input (the main input stream) or files and writes its main results to its standard output (the main output stream)
- Filters do not modify input files
- Filters may be strung together into a pipeline with the pipe operator ("|"). This operator signifies that the main output of the command to the left is passed as main input to the command on the right... more on that later

[https://en.wikipedia.org/wiki/Filter_\(software\)#Unix](https://en.wikipedia.org/wiki/Filter_(software)#Unix)

Head filter

Print the beginning of a text file (or piped data)

- head [options] <file_name>

The screenshot shows a terminal window titled "FiltersTestDIR — bash — 64x21". The window contains the following text:

```
[iMac-2:FiltersTestDIR marcoautili$]  
[iMac-2:FiltersTestDIR marcoautili$ head persone.txt]  
Nome Cognome Età  
  
Mario Rossi 35  
Mario Bianchi 36  
Paolo Rossi 20  
Pippo Verdi 44  
Pluto Verdone 21  
Minnie Rosa 28
```

An orange callout box points from the word "head" in the first command to the text "10 lines are shown by default".

```
[iMac-2:FiltersTestDIR marcoautili$ head -n 4 persone.txt]  
Nome Cognome Età  
  
Mario Rossi 35  
Mario Bianchi 36
```

An orange callout box points from the "-n 4" option in the command to the text "The first 4 lines are shown".

```
[iMac-2:FiltersTestDIR marcoautili$ head -c 12 persone.txt]  
Nome Cognome  
[iMac-2:FiltersTestDIR marcoautili$]  
[iMac-2:FiltersTestDIR marcoautili$]
```

An orange callout box points from the "-c 12" option in the command to the text "The first 12 bytes are shown".

Tail filter

Print the tail of a text file (or piped data)

- `tail [options] <file_name>`

```
iMac-2:FiltersTestDIR marcoautili$  
iMac-2:FiltersTestDIR marcoautili$  
iMac-2:FiltersTestDIR marcoautili$ tail -3 persone.txt
```

Pippo Verdi 44
Pluto Verdone 21
Minnie Rosa 28

```
iMac-2:FiltersTestDIR marcoautili$ tail -c 3 persone.txt
```

28

```
iMac-2:FiltersTestDIR marcoautili$ tail -c 3 persone.txt
```

```
iMac-2:FiltersTestDIR marcoautili$
```

There is no
End of Line

```
personne.txt  
1 Nome Cognome Età  
2  
3 Mario Rossi 35  
4 Mario Bianchi 36  
5 Paolo Rossi 20  
6 Pippo Verdi 44  
7 Pluto Verdone 21  
8 Minnie Rosa 28
```

There is
End of Line

No need to write
`-n 3`

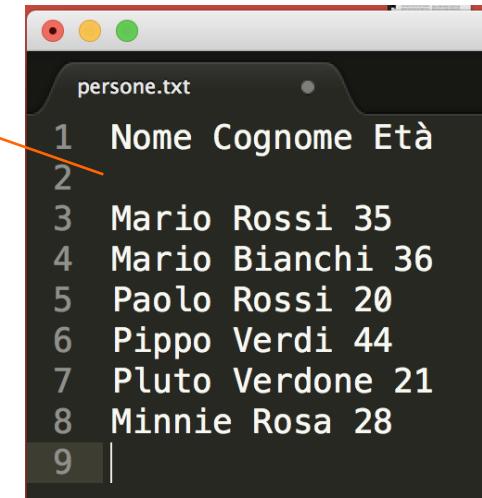
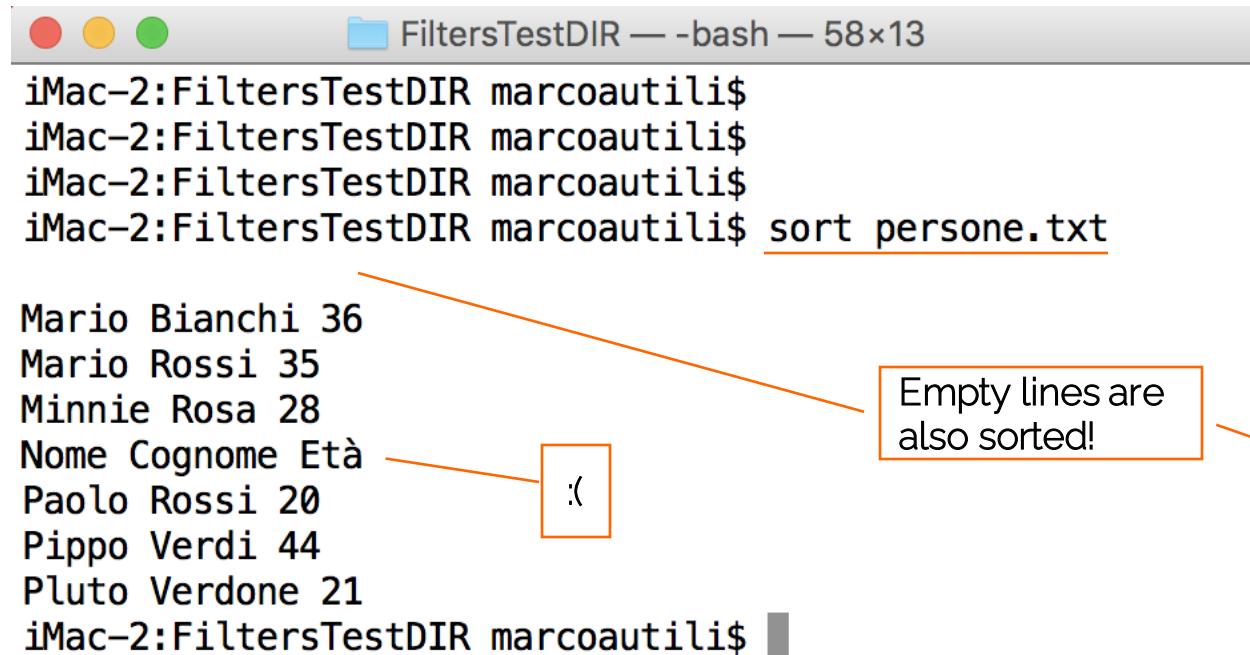
```
personne.txt  
1 Nome Cognome Età  
2  
3 Mario Rossi 35  
4 Mario Bianchi 36  
5 Paolo Rossi 20  
6 Pippo Verdi 44  
7 Pluto Verdone 21  
8 Minnie Rosa 28  
9 |
```

Sort filters

The sort utility sorts **text** and **binary** files by lines (by default it will sort alphabetically - see man :-)

- sort [-options] [path]

```
iMac-2:FiltersTestDIR marcoautili$  
iMac-2:FiltersTestDIR marcoautili$  
iMac-2:FiltersTestDIR marcoautili$  
iMac-2:FiltersTestDIR marcoautili$ sort persone.txt  
  
Mario Bianchi 36  
Mario Rossi 35  
Minnie Rosa 28  
Nome Cognome Età  
Paolo Rossi 20  
Pippo Verdi 44  
Pluto Verdone 21  
iMac-2:FiltersTestDIR marcoautili$
```



Numero	Nome Cognome	Età
1	Mario Rossi	35
2	Mario Bianchi	36
3	Paolo Rossi	20
4	Pippo Verdi	44
5	Pluto Verdone	21
6	Minnie Rosa	28

Number Lines Filter

Add line numbers

- `nl [-options] [path]`

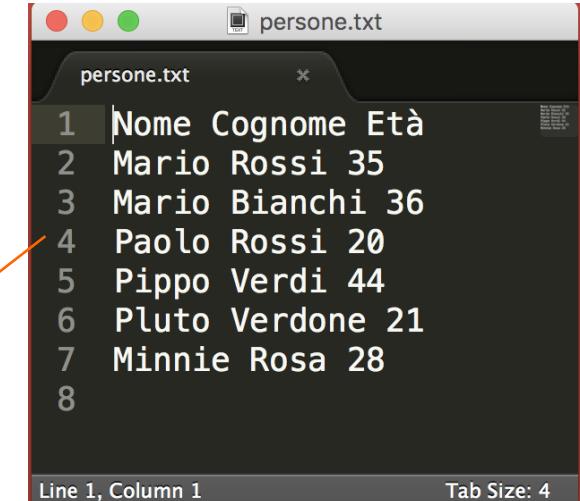
```
iMac-2:FiltersTestDIR marcoautili$ nl persone.txt
1 Nome Cognome Età
2 Mario Rossi 35
3 Mario Bianchi 36
4 Paolo Rossi 20
5 Pippo Verdi 44
6 Pluto Verdone 21
7 Minnie Rosa 28
iMac-2:FiltersTestDIR marcoautili$ nl -s '.' ' -w 10 persone.txt
1. Nome Cognome Età
2. Mario Rossi 35
3. Mario Bianchi 36
4. Paolo Rossi 20
5. Pippo Verdi 44
6. Pluto Verdone 21
7. Minnie Rosa 28
iMac-2:FiltersTestDIR marcoautili$
```

Basic formatting

width

Custom
formatting

Note that these
numbers are not
part of the file
content!



Word Count Filter

Count lines, words, and characters

- `wc [-options] [path]`

```
iMac:PART 2 - Filters-Test-DIR marcoautili$  
iMac:PART 2 - Filters-Test-DIR marcoautili$ wc persone.txt  
     8      21    129 persone.txt  
iMac:PART 2 - Filters-Test-DIR marcoautili$ wc -c persone.txt  
    129 persone.txt  
iMac:PART 2 - Filters-Test-DIR marcoautili$ wc -m persone.txt  
    128 persone.txt  
iMac:PART 2 - Filters-Test-DIR marcoautili$ wc -l persone.txt  
     8 persone.txt  
iMac:PART 2 - Filters-Test-DIR marcoautili$ wc -w persone.txt  
    21 persone.txt  
iMac:PART 2 - Filters-Test-DIR marcoautili$
```

Lines, words, and bytes by default

Bytes only

Chars only

Lines only

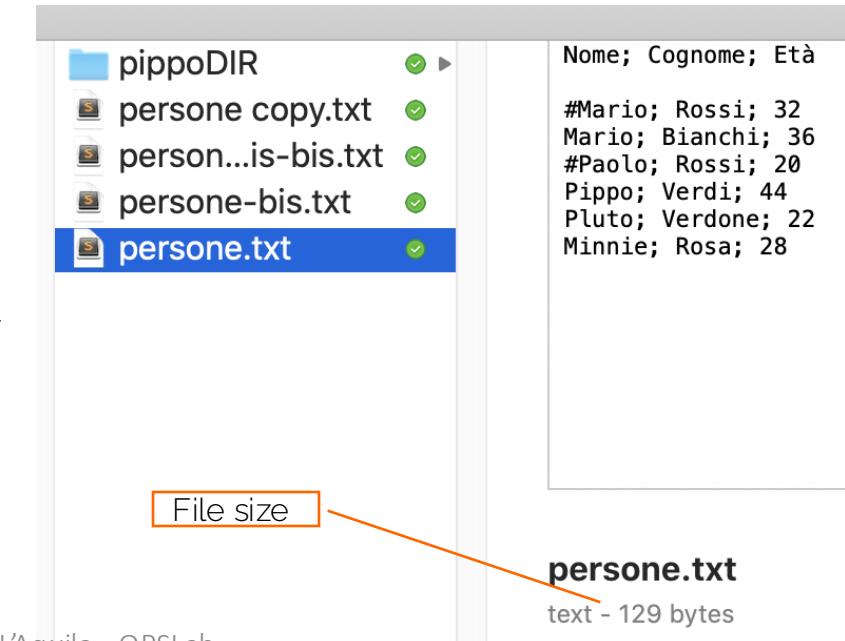
Words only

-c The number of bytes in each input file is written to the standard output. This will cancel out any prior usage of the **-m** option.

-m The number of characters in each input file is written to the standard output. If the current locale does not support multibyte characters, this is equivalent to the **-c** option. This will cancel out any prior usage of the **-c** option.

-l The number of lines in each input file is written to the standard output.

-w The number of words in each input file is written to the standard output.



Cut Filter

Print only certain fields (if the contents of the file is organized into fields)

- `cut [-options] [path]`

```
iMac-2:FiltersTestDIR marcoautili$ cut -f 1 -d ' ' persone.txt
Nome
Mario
Mario
Paolo
Pippo
Pluto
Minnie
[iMac-2:FiltersTestDIR marcoautili$ cut -f 1 -d ';' persone.txt
Nome Cognome Età
Mario Rossi 35
Mario Bianchi 36
Paolo Rossi 20
Pippo Verdi 44
Pluto Verdone 21
Minnie Rosa 28
[iMac-2:FiltersTestDIR marcoautili$ cut -f 1 -d ';' persone.txt
Nome
Mario
Mario
Paolo
Pippo
Pluto
Minnie
iMac-2:FiltersTestDIR marcoautili$
```

"Single Space"
separator.
TAB by default

"Semicolon"
separator: it
does not
work on this
file!

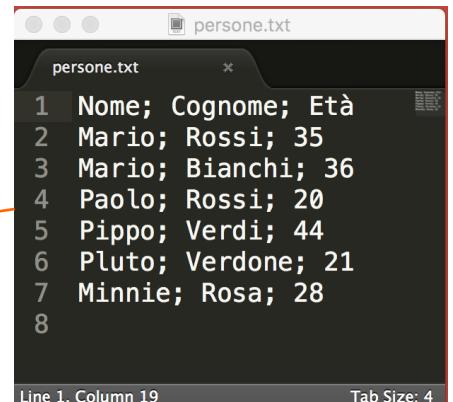
"Semicolon"
separator



personne.txt

Line	Nome	Cognome	Età
1	Mario	Rossi	35
2	Mario	Bianchi	36
3	Paolo	Rossi	20
4	Pippo	Verdi	44
5	Pluto	Verdone	21
6	Minnie	Rosa	28
7			
8			

Line 1, Column 1 Tab Size: 4



personne.txt

Line	Nome	Cognome	Età
1	Mario	Rossi	35
2	Mario	Bianchi	36
3	Paolo	Rossi	20
4	Pippo	Verdi	44
5	Pluto	Verdone	21
6	Minnie	Rosa	28
7			
8			

Line 1, Column 19 Tab Size: 4

Unique Filter

Remove duplicate adjacent lines

- `uniq [options] [path]`

```
iMac-2:FiltersTestDIR marcoautili$ uniq persone.txt
Nome; Cognome; Età
Mario; Rossi; 35
Mario; Bianchi; 36
Paolo; Rossi; 20
Pippo; Verdi; 44
Pluto; Verdone; 21
Minnie; Rosa; 28
iMac-2:FiltersTestDIR marcoautili$ uniq persone.txt
Nome; Cognome; Età
Mario; Rossi; 35
Mario; Rossi; 35
Mario; Bianchi; 36
Paolo; Rossi; 20
Paolo; Rossi; 20
Pippo; Verdi; 44
Pluto; Verdone; 21
Minnie; Rosa; 28
iMac-2:FiltersTestDIR marcoautili$
```

Adjacent lines: it works!

```
1 Nome; Cognome; Età
2 Mario; Rossi; 35
3 Mario; Rossi; 35
4 Mario; Bianchi; 36
5 Paolo; Rossi; 20
6 Paolo; Rossi; 20
7 Paolo; Rossi; 20
8 Pippo; Verdi; 44
9 Pluto; Verdone; 21
10 Minnie; Rosa; 28
11
```

```
1 Nome; Cognome; Età
2 Mario; Rossi; 35
3
4 Mario; Rossi; 35
5 Mario; Bianchi; 36
6 Paolo; Rossi; 20
7 Paolo; Rossi; 20
8
9 Paolo; Rossi; 20
10 Pippo; Verdi; 44
11 Pluto; Verdone; 21
12 Minnie; Rosa; 28
```

Not adjacent lines: it does not work!

Later on we will see how to fix this ...

Tac Command

Reverse cat: print last line first

- tac [path]

```
iMac-2:FiltersTestDIR marcoautili$ uniq persone.txt  
Nome; Cognome; Età  
Mario; Rossi; 35  
  
Mario; Rossi; 35  
Mario; Bianchi; 36  
Paolo; Rossi; 20  
  
Paolo; Rossi; 20  
Pippo; Verdi; 44  
Pluto; Verdone; 21  
Minnie; Rosa; 28  
[iMac-2:FiltersTestDIR marcoautili$ tac persone.txt  
-bash: tac: command not found  
[iMac-2:FiltersTestDIR marcoautili$  
[iMac-2:FiltersTestDIR marcoautili$  
[iMac-2:FiltersTestDIR marcoautili$ man tac  
No manual entry for tac  
[iMac-2:FiltersTestDIR marcoautili$  
[iMac-2:FiltersTestDIR marcoautili$  
[iMac-2:FiltersTestDIR marcoautili$
```

```
personne.txt  
1 Nome; Cognome; Età  
2 Mario; Rossi; 35  
3 Mario; Bianchi; 36  
4 Paolo; Rossi; 20  
5 Pippo; Verdi; 44  
6 Pluto; Verdone; 21  
7 Minnie; Rosa; 28  
8  
Line 2, Column 17  
Tab Size: 102
```

I do not have it!
What about you?...
just try!

Diff Filter

- `diff` compare two files (**line by line**) and print the lines that are different
- It outputs a set of instructions for how to change the first file to make it identical to the second file
- It does not modify the files
- **It is quite powerful...**
 - For instance, the `-e` option tells diff to output a script, which can be used by the editing programs ed or ex, that contains a sequence of commands
 - ... we will make basic use of it as follows

<http://www.computerhope.com/unix/udiff.htm>

Diff Filter

The first line of the **diff** output will contain:

- line numbers corresponding to the first file
- a letter (a for *add*, c for *change*, or d for *delete*)
- line numbers corresponding to the second file.

In the example, "2,6c2,6" means:

- "Lines 2 through 6 in the first file need to be changed to match lines 2 through 6 in the second file."

It then tells us what those lines are in each file:

- Lines preceded by a < are lines from the **first file**
- lines preceded by > are lines from the **second file**
- the three dashes ("---") merely separate the lines of file 1 and file 2.

personne.txt

1 Nome; Cognome; Età
2 # Mario; Rossi; 35
3 Mario; Bianchi; 36
4 # Paolo; Rossi; 20
5 Pippo; Verdi; 44
6 Pluto; Verdone; 21
7 Minnie; Rosa; 28
8

Line 2, Column 3 Tab Size: 4

personne-bis.txt

1 Nome; Cognome; Età
2 # Mario; Rossi; 40
3 Mario; Bianchi; 50
4 # Paolo; Rossi; 60
5 Pippo; Verdi; 60
6 Paolo; Verdone; 21
7 Minnie; Rosa; 28
8

Line 7, Column 6 Tab Size: 4

iMac-2:FiltersTestDIR marcoautili\$ diff persone.txt persone-bis.txt

2,6c2,6

< # Mario; Rossi; 35

< Mario; Bianchi; 36

< # Paolo; Rossi; 20

< Pippo; Verdi; 44

< Pluto; Verdone; 21

> # Mario; Rossi; 40

> Mario; Bianchi; 50

> # Paolo; Rossi; 60

> Pippo; Verdi; 60

> Paolo; Verdone; 21

iMac-2:FiltersTestDIR marcoautili\$

These instructions tell how to change the first file to make it match the second file

Diff Filter



Filters-Test-DIR — -bash — 84x10

```
> Pippo; Verdi; 60
[desktop-jisqlks:Filters-Test-DIR marcoautili$]
[desktop-jisqlks:Filters-Test-DIR marcoautili$]
[desktop-jisqlks:Filters-Test-DIR marcoautili$] diff persone.txt persone-bis-bis.txt
5d4
< #Paolo; Rossi; 20
[desktop-jisqlks:Filters-Test-DIR marcoautili$]
[desktop-jisqlks:Filters-Test-DIR marcoautili$]
[desktop-jisqlks:Filters-Test-DIR marcoautili$]
[desktop-jisqlks:Filters-Test-DIR marcoautili$]
```

- You need to **delete** line **5** in the first file so that both files sync up at line **4**
- It then shows us the contents of the line that needs to be deleted.

The screenshot shows two terminal windows and their corresponding file contents. The left terminal window shows the command 'diff persone.txt persone-bis-bis.txt' and its output, which includes line 5 from 'personne.txt' followed by a delete character (indicated by a backspace) and line 4 from 'personne-bis-bis.txt'. The right terminal window shows the contents of 'personne.txt' and 'personne-bis-bis.txt'. Both files contain the same list of names and ages, except for line 5 which is present in 'personne.txt' but absent in 'personne-bis-bis.txt'.

File Contents:

- personne.txt:**

```
1 Nome; Cognome; Età
2
3 #Mario; Rossi; 35
4 Mario; Bianchi; 36
5 #Paolo; Rossi; 20
6 Pippo; Verdi; 44
7 Pluto; Verdone; 22
8 Minnie; Rosa; 28
9
```
- personne-bis-bis.txt:**

```
1 Nome; Cognome; Età
2
3 #Mario; Rossi; 35
4 Mario; Bianchi; 36
5 Pippo; Verdi; 44
6 Pluto; Verdone; 22
7 Minnie; Rosa; 28
8
```



Filters-Test-DIR — -bash — 85x7

```
desktop-jisqlks:Filters-Test-DIR marcoautili$ diff persone.txt persone-bis-bis.txt
3a4
> Mario; Bianchi; 36
[desktop-jisqlks:Filters-Test-DIR marcoautili$]
[desktop-jisqlks:Filters-Test-DIR marcoautili$]
[desktop-jisqlks:Filters-Test-DIR marcoautili$]
[desktop-jisqlks:Filters-Test-DIR marcoautili$]
```

- After line **3** in the first file, a line needs to be **added**: the line to be added is line **4** from the second file
- It then shows us what that line to be added is

The screenshot shows two terminal windows and their corresponding file contents. The left terminal window shows the command 'diff persone.txt persone-bis-bis.txt' and its output, which includes line 3 from 'personne.txt' followed by a new line character and line 4 from 'personne-bis-bis.txt'. The right terminal window shows the contents of 'personne.txt' and 'personne-bis-bis.txt'. Both files contain the same list of names and ages, except for line 3 which is present in 'personne.txt' but absent in 'personne-bis-bis.txt'.

File Contents:

- personne.txt:**

```
1 Nome; Cognome; Età
2
3 #Mario; Rossi; 35
4 #Paolo; Rossi; 20
5 Pippo; Verdi; 44
6 Pluto; Verdone; 22
7 Minnie; Rosa; 28
8
```
- personne-bis-bis.txt:**

```
1 Nome; Cognome; Età
2
3 #Mario; Rossi; 35
4 Mario; Bianchi; 36
5 #Paolo; Rossi; 20
6 Pippo; Verdi; 44
7 Pluto; Verdone; 22
8 Minnie; Rosa; 28
9
```

Regular Expressions

Regular expressions are a type of globbing pattern used when working with **text**

- Similarly to wildcards for paths, regular expressions allow for creating patterns for text files, although more powerful, e.g., identify all strings in a file that represent URLs, comment lines, email addresses, etc.

IMPORTANT NOTE

- As you know, **POSIX** (Portable Operating System Interface for uniX) is a **collection of standards** that define some of the functionality that a (UNIX) operating system should support
- One of these standards defines **two flavors** of regular expressions:
 - Basic Regular Expressions (BRE)
 - Extended Regular Expressions (ERE)
 - Commands using regular expressions, such as **grep** and **egrep**, implement these flavors on POSIX-compliant UNIX systems

Egrep and Regular Expressions

The grep utility searches any given input files, selecting **lines** that match one or more patterns (egrep stands for **extended grep**)

- egrep [command line options] <pattern> [path]

```
iMac-2:FiltersTestDIR marcoautili$ man egrep
iMac-2:FiltersTestDIR marcoautili$
iMac-2:FiltersTestDIR marcoautili$ egrep Mario persone.txt
# Mario; Rossi; 35
Mario; Bianchi; 36
iMac-2:FiltersTestDIR marcoautili$ egrep 'o; V' persone.txt
Pippo; Verdi; 44
[Pluto; Verdone; 21
iMac-2:FiltersTestDIR marcoautili$ egrep -n 'o; V' persone.txt
5:Pippo; Verdi; 44
[6:Pluto; Verdone; 21
iMac-2:FiltersTestDIR marcoautili$ egrep -c 'o; V' persone.txt
2
iMac-2:FiltersTestDIR marcoautili$
```

```
personne.txt
1 Nome; Cognome; Età
2 # Mario; Rossi; 35
3 Mario; Bianchi; 36
4 # Paolo; Rossi; 20
5 Pippo; Verdi; 44
6 Pluto; Verdone; 21
7 Minnie; Rosa; 28
8

Line 2, Column 3          Tab Size: 4
```

Regular Expressions Format

```
iMac:~ marcoautili$ man 7 re_format  
Or simply: man re_format
```

```
RE_FORMAT(7)      BSD Miscellaneous Information Manual      RE_FORMAT(7)
```

NAME

re_format -- POSIX 1003.2 regular expressions

DESCRIPTION

Regular expressions ('`REs'''), as defined in IEEE Std 1003.2 ('`POSIX.2'''), come in two forms: modern REs (roughly those of egrep(1); 1003.2 calls these ``extended'' REs) and obsolete REs (roughly those of ed(1); 1003.2 ``basic'' REs). Obsolete REs mostly exist for backward compatibility in some old programs; they will be discussed at the end. IEEE Std 1003.2 ('`POSIX.2''') leaves some aspects of RE syntax and semantics open; '=' marks decisions on these aspects that may not be fully portable to other IEEE Std 1003.2 ('`POSIX.2''') implementations.



A (modern) RE is one= or more non-empty= branches, separated by '|'. It matches anything that matches one of the branches.

A branch is one= or more pieces, concatenated. It matches a match for the first, followed by a match for the second, etc.

A piece is an atom possibly followed by a single= `*', `+', `?', or bound. An atom followed by '*' matches a sequence of 0 or more matches of the atom. An atom followed by '+' matches a sequence of 1 or more matches of the atom. An atom followed by '?' matches a sequence of 0 or 1 matches of the atom.

Regular Expressions Format

marcoautili — less - man 7 re_format — 95x36

A bound is `{' followed by an unsigned decimal integer, possibly followed by `,' possibly followed by another unsigned decimal integer, always followed by `}'.

The integers must lie between 0 and RE_DUP_MAX (255=) inclusive, and if there are two of them, the first may not exceed the second. An atom followed by a bound containing one integer i and no comma matches a sequence of exactly i matches of the atom. An atom followed by a bound containing one integer i and a comma matches a sequence of i or more matches of the atom. An atom followed by a bound containing two integers i and j matches a sequence of i through j (inclusive) matches of the atom.

An atom is a regular expression enclosed in `()' (matching a match for the regular expression), an empty set of `()' (matching the null string)=, a bracket expression (see below), `.' (matching any single character), `^' (matching the null string at the beginning of a line), `'\$' (matching the null string at the end of a line), a `\' followed by one of the characters `^.[\$()|*+?{\\' (matching that character taken as an ordinary character), a `\' followed by any other character= (matching that character taken as an ordinary character, as if the `\' had not been present=), or a single character with no other significance (matching that character). A `{' followed by a character other than a digit is an ordinary character, not the beginning of a bound=. It is illegal to end an RE with `\'.

Regular Expressions Format

marcoautili — less ◀ man 7 re_format — 95x36

A bracket expression is a list of characters enclosed in `[]'. It normally matches any single character from the list (but see below). If the list begins with `^', it matches any single character (but see below) not from the rest of the list. If two characters in the list are separated by `-', this is shorthand for the full range of characters between those two (inclusive) in the collating sequence, e.g. `'[0-9]' in ASCII matches any decimal digit. It is illegal for two ranges to share an endpoint, e.g. `'[a-c-e]'. Ranges are very collating-sequence-dependent, and portable programs should avoid relying on them.

To include a literal `]' in the list, make it the first character (following a possible `^'). To include a literal `-', make it the first or last character, or the second endpoint of a range. To use a literal `-' as the first endpoint of a range, enclose it in `[' and `']' to make it a collating element (see below). With the exception of these and some combinations using `[' (see next paragraphs), all other special characters, including `\'', lose their special significance within a bracket expression.

Within a bracket expression, a collating element (a character, a multi-character sequence that collates as if it were a single character, or a collating-sequence name for either) enclosed in `[' and `']' stands for the sequence of characters of that collating element. The sequence is a single element of the bracket expression's list. A bracket expression containing a multi-character collating element can thus match more than one character, e.g. if the collating sequence includes a `ch' collating element, then the RE `'[.ch.]>*c' matches the first five characters of `chchcc'.

....

Regular Expressions made easier for you 😊

A (modern) **regular expression** is one or more **non-empty branches** separated by a pipe |
It matches anything that matches **one of** the branches

A **branch** is one or more **pieces**, concatenated. It matches a match for the first, followed by a match for the second, etc.

A **piece** is an **atom** possibly followed by a single *, +, ?, or **bound**

- * - the preceding atom matches 0 or more times
- + - the preceding atom matches 1 or more times
- ? - the preceding atom matches 0 or 1 times only
- [n] - the preceding atom matches exactly n times
- [n,m] -the preceding atom matches at least n and not more than m times ($n \leq m$)
- [n,] -the preceding atom matches at least n times (there is no upper bound)
 - A single [or a [followed by a character other than a digit is an ordinary character, not the beginning of a bound. Thus, e.g., a single [matches [and [a matches [a
 - It is illegal to end a regular expression with \

Regular Expressions made easier for you 😊

An **atom** is

- **(reg exp)** - a regular expression enclosed in () matches a match for the regular expression
- **()** - matches the empty string
- **.** - a dot matches any single character
- **^** - matches the null string at the beginning of a line
- **\$** - matches the null string at the end of a line
- **\ followed by one of the characters ^ . [\$ () | * + ? { ** - matches that character taken as an ordinary character
- **\ followed by any other character** - matches that character taken as an ordinary character, as if the \ had not been present
- **A single character with no other significance** - matches that character

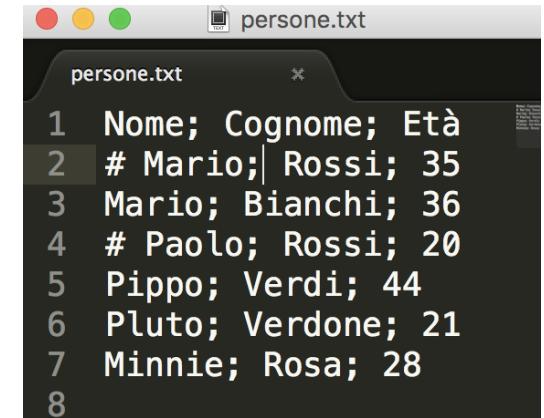
Regular Expressions made easier for you 😊

An atom is also a **bracket expression**

A **bracket expression** is a **list of characters enclosed in []**

- **[agd]** – matches a single character included within []
 - **[^agd]** – matches any single character different from those included within []
 - **[c-f]** - the dash within the square brackets operates as a range (inclusive). In this case it means either the letters c, d, e or f
-
- It is illegal for two ranges to share an endpoint, e.g., **[a-c-e]**
However, **[h-j4-9]** is legal as well as **[a-dh-j]** and **[5-81-3]**
 - To include the character **|** in the list, make it the first character of the list, e.g., **[\abhg]** or **[\|]**
 - Note that **\|** can be used to escape any of the characters **^ . [\$ () * + ? { \ \ , e.g., [\\$, [?], [*]**
Thus, **[\\$], [?], [*]** are equivalent to the escaping **\\$, \?, ***
 - Note that **[\|]** is equivalent to **\|** and indeed it matches **\|** since the bracket expression **\|** matches **|** and the single **|** matches **|**

Egrep: simple examples



Content of persone.txt:

```
1 Nome; Cognome; Età
2 # Mario; Rossi; 35
3 Mario; Bianchi; 36
4 # Paolo; Rossi; 20
5 Pippo; Verdi; 44
6 Pluto; Verdone; 21
7 Minnie; Rosa; 28
8
```

```
iMac-2:FiltersTestDIR marcoautili$ egrep '35|36' persone.txt
# Mario; Rossi; 35
Mario; Bianchi; 36
iMac-2:FiltersTestDIR marcoautili$ egrep 8$ persone.txt
Minnie; Rosa; 28
iMac-2:FiltersTestDIR marcoautili$ egrep '[aeiou]{2,}' persone.txt
# Mario; Rossi; 35
Mario; Bianchi; 36
# Paolo; Rossi; 20
Minnie; Rosa; 28
iMac-2:FiltersTestDIR marcoautili$ egrep '3.' persone.txt
# Mario; Rossi; 35
Mario; Bianchi; 36
iMac-2:FiltersTestDIR marcoautili$
```

Annotations from top to bottom:

- Persons 35 or 36 years old
- Lines ending with 8
- Lines containing at least two adjacent vowels
- Persons 35 or 36 years old (as before)

Egrep: simple examples

Filters-Test-DIR — bash — 82x14

```
desktop-jisqlks:Filters-Test-DIR marcoautili$ egrep '(ss)|(pp)' persone.txt
#Mario; Rossi; 3
#Paolo; Rossi; 20
Pippo; Verdi; 44
desktop-jisqlks:Filters-Test-DIR marcoautili$ egrep '^Ma' persone-bis-bis.txt
Mario; Bianchi; 36
desktop-jisqlks:Filters-Test-DIR marcoautili$ egrep '(ss) | (pp)' persone.txt
desktop-jisqlks:Filters-Test-DIR marcoautili$
```

All the lines starting with "Ma"

All the lines containing "ss" or "pp"

No space in between!

Line 4, Column 19 Tab Siz

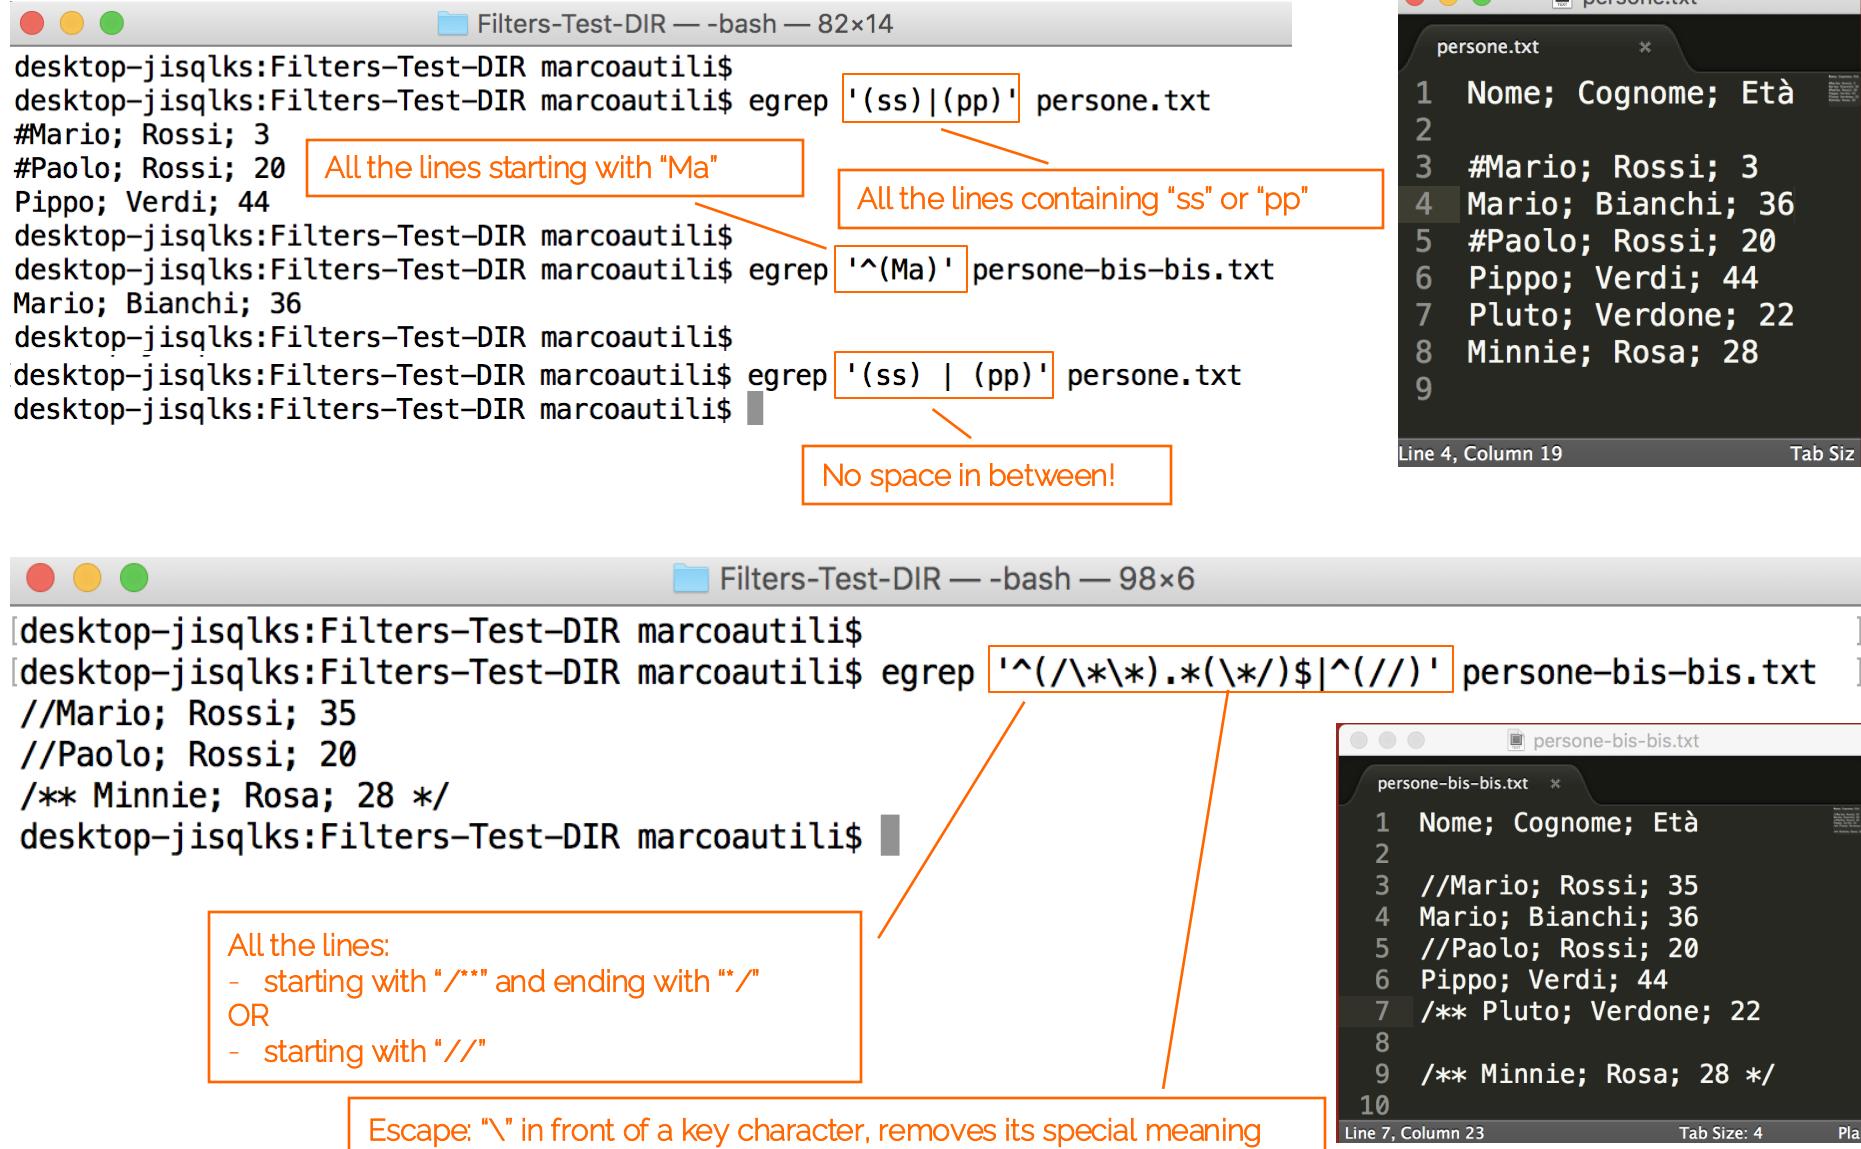
Filters-Test-DIR — bash — 98x6

```
desktop-jisqlks:Filters-Test-DIR marcoautili$ egrep '^(\/*\*).*(\*/)$|^\/\/' persone-bis-bis.txt
//Mario; Rossi; 35
//Paolo; Rossi; 20
/** Minnie; Rosa; 28 */
desktop-jisqlks:Filters-Test-DIR marcoautili$
```

All the lines:
- starting with "/*" and ending with "*/"
OR
- starting with "//"

Escape: "\ in front of a key character, removes its special meaning

Line 7, Column 23 Tab Size: 4 Plai



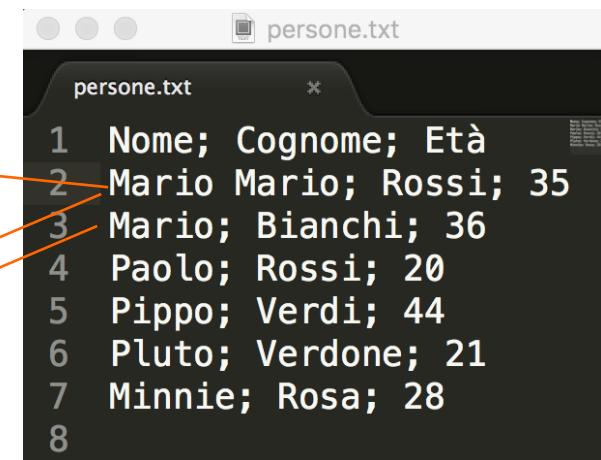
Stream editor Filter

- Very powerful... some basics follow
- Type **CTRL+c** in case of undesired behavior 😊

`sed <command> [file]`

- Search and replace
 - the basic expression format is `s/search/replace/g`:
 - "s" stands for substitute and "g" stands for global (it is optional). There are others ...
 - `s/search/replace` -> will replace only the first instance on each line

```
iMac-2:FiltersTestDIR marcoautili$ sed s/Mario/Pasquale/g persone.txt
Nome; Cognome; Età
Pasquale Pasquale; Rossi; 35
Pasquale; Bianchi; 36
Paolo; Rossi; 20
Pippo; Verdi; 44
Pluto; Verdone; 21
Minnie; Rosa; 28
iMac-2:FiltersTestDIR marcoautili$ sed s/Mario/Pasquale/ persone.txt
Nome; Cognome; Età
Pasquale Mario; Rossi; 35
Pasquale; Bianchi; 36
Paolo; Rossi; 20
Pippo; Verdi; 44
Pluto; Verdone; 21
Minnie; Rosa; 28
iMac-2:FiltersTestDIR marcoautili$
```



Stream Editor Filter

```
iMac-2:FiltersTestDIR marcoautili$  
iMac-2:FiltersTestDIR marcoautili$ sed -e '1,2d' persone.txt  
Mario; Bianchi; 36  
# Paolo; Rossi; 20  
Pippo; Verdi; 44  
Pluto; Verdone; 21  
Minnie; Rosa; 28  
iMac-2:FiltersTestDIR marcoautili$ sed -e '/^#/d' persone.txt  
Nome; Cognome; Età  
Mario; Bianchi; 36  
Pippo; Verdi; 44  
Pluto; Verdone; 21  
Minnie; Rosa; 28  
iMac-2:FiltersTestDIR marcoautili$
```

Delete the first two lines

Delete #-commented lines

```
personne.txt  
1 Nome; Cognome; Età  
2 # Mario; Rossi; 35  
3 Mario; Bianchi; 36  
4 # Paolo; Rossi; 20  
5 Pippo; Verdi; 44  
6 Pluto; Verdone; 21  
7 Minnie; Rosa; 28  
8  
Line 2, Column 3 Tab Size: 4
```

One method of combining multiple commands is to use a **-e** before each command:

sed -e 's/a/A/' -e 's/b/B/' persone.txt

The "**-e**" isn't needed in the case of a single command, e.g., **sed 's/a/A/' persone.txt**

The long argument version should be:

'sed --expression='s/a/A/' --expression='s/b/B/' (NOT WORKING ON MY macOS... Try on your machine)

Sed and regular expression



SED(1)

BSD General Commands Manual

SED(1)

NAME

sed -- stream editor

SYNOPSIS

```
sed [-Ealn] command [file ...]
sed [-Ealn] [-e command] [-f command_file] [-i extension] [file ...]
```

DESCRIPTION

The **sed** utility reads the specified files, or the standard input if no files are specified, modifying the input as specified by a list of commands. The input is then written to the standard output.

A single command may be specified as the first argument to **sed**. Multiple commands may be specified by using the **-e** or **-f** options. All commands are applied to the input in the order they are specified regardless of their origin.

The following options are available:

-E Interpret regular expressions as extended (modern) regular expressions rather than basic regular expressions (BRE's). The *re_format(7)* manual page fully describes both formats.

:

Example in
next slide

Sed: BRE VS ERE

```
[aaca]
[a2cbb]
[a2bc]
```

```
bash-3.2$ sed -e '/^\\[a{2,}\\]/d' es2.txt
[aaca]
[a2cbb]
[a2bc]
bash-3.2$ sed -E '/^\\[a{2,}\\]/d' es2.txt
[a2cbb]
[a2bc]
```

BRE

ERE

Awk Filter

- AWK is a **language** for processing text files
- A **file** is treated as a sequence of **records**, and by default each line is a record
- Each **line** is broken up into a sequence of **fields**, so we can think of the first word in a line as the first field, the second word as the second field, and so on
- AWK reads the input a line at a time
- An **AWK program** is a sequence of pattern-action statements
 - A line is scanned for each pattern in the program, and for each pattern that matches, the associated action is executed

[https://en.wikipedia.org/wiki/Filter_\(software\)#Unix](https://en.wikipedia.org/wiki/Filter_(software)#Unix)

Awk Filter

- Special-purpose language for line-oriented pattern processing
 - pattern {action}

- action =
 - if (conditional) *statement* else *statement*
 - while (conditional) statement
 - break
 - continue
 - variable=expression
 - print expression-list

We will not treat AWK

* [https://en.wikipedia.org/wiki/Filter_\(software\)#Unix](https://en.wikipedia.org/wiki/Filter_(software)#Unix)

Homework activities

- Play with all the filters
- Have a look at man egrep and play with its options
- Define increasingly complex regular expressions and test them using egrep on a file you have aptly defined

Pipeline and Redirection

Let's introduce standard data streams first

- Standard data streams are pre-connected communication channels
 - Every program/command executed on the command line automatically has the following standard streams connected to it

STDIN (0)

- Standard input (data fed into the program, from the terminal through the keyboard by default)

STDOUT (1)

- Standard output (data printed by the program, defaults to the terminal)

STDERR (2)

- Standard error (for error messages, also defaults to the terminal)

The three streams have numbers associated with them



Redirecting STDOUT to a file

The operator ">" (or ">>") permits to redirect the output of a program or a command to a file, rather than printing the output on the screen (as by default)

The screenshot shows a macOS desktop environment. On the left, a terminal window titled "OPSLab-Test-DIR — bash — 81x42" displays a series of commands and their outputs. On the right, a Finder window titled "OPSLab-Test-DIR" shows the directory structure.

Terminal Output:

```
iMac-2:OPSLab-Test-DIR marcoautili$ pwd  
/Users/marcoautili/Desktop/OPSLab-Test-DIR  
iMac-2:OPSLab-Test-DIR marcoautili$ ls -la > output-file.txt  
iMac-2:OPSLab-Test-DIR marcoautili$ cat output-file.txt  
total 32  
drwxr-xr-x@ 7 marcoautili staff 238 Sep 20 15:32 .  
drwx-----@ 34 marcoautili staff 1156 Sep 20 15:32 ..  
-rw-r--r--@ 1 marcoautili staff 6148 Sep 20 15:30 .DS_Store  
drwxr-xr-x@ 6 marcoautili staff 204 Sep 20 08:53 Filters-Test-DIR  
-rw-r--r-- 1 marcoautili staff 0 Sep 20 15:32 output-file.txt  
-rw-r--r--@ 1 marcoautili staff 14 Sep 20 15:26 pippo.txt  
-rw-r--r--@ 1 marcoautili staff 14 Sep 20 15:29 pluto.txt  
iMac-2:OPSLab-Test-DIR marcoautili$  
iMac-2:OPSLab-Test-DIR marcoautili$ diff pippo.txt pluto.txt >> output-file.txt  
iMac-2:OPSLab-Test-DIR marcoautili$ cat output-file.txt  
total 32  
drwxr-xr-x@ 7 marcoautili staff 238 Sep 20 15:32 .  
drwx-----@ 34 marcoautili staff 1156 Sep 20 15:32 ..  
-rw-r--r--@ 1 marcoautili staff 6148 Sep 20 15:30 .DS_Store  
drwxr-xr-x@ 6 marcoautili staff 204 Sep 20 08:53 Filters-Test-DIR  
-rw-r--r-- 1 marcoautili staff 0 Sep 20 15:32 output-file.txt  
-rw-r--r--@ 1 marcoautili staff 14 Sep 20 15:26 pippo.txt  
-rw-r--r--@ 1 marcoautili staff 14 Sep 20 15:29 pluto.txt  
1c1  
< pippo-file.txt  
\ No newline at end of file  
---  
> pluto-file.txt  
\ No newline at end of file  
iMac-2:OPSLab-Test-DIR marcoautili$ diff pippo.txt pluto.txt > output-file.txt  
iMac-2:OPSLab-Test-DIR marcoautili$ cat output-file.txt  
1c1  
< pippo-file.txt  
\ No newline at end of file  
---  
> pluto-file.txt  
\ No newline at end of file  
iMac-2:OPSLab-Test-DIR marcoautili$
```

Finder Window:

- Left pane: COP-2008.pdf, EmailSe...-1.0.dmg, Evaluati...ndard.pdf, Files_Na...-OK.scpt, IMPIANTI, inglese-inf, maillist.xml, OPSLab-Test-DIR (selected).
- Right pane: Filters-Test-DIR, output-file.txt (selected), pippo.txt, pluto.txt.

Redirecting STDIN/STDOUT from/to a file

```
iMac-2:OPSLab-Test-DIR marcoautili$ cat output-file.txt  
1c1  
< pippo-file.txt  
\ No newline at end of file  
---  
> pluto-file.txt  
\ No newline at end of file  
iMac-2:OPSLab-Test-DIR marcoautili$ wc -l < output-file.txt  
6  
iMac-2:OPSLab-Test-DIR marcoautili$ wc -l < output-file.txt > output-file-bis.txt  
iMac-2:OPSLab-Test-DIR marcoautili$ cat output-file-bis.txt  
6  
iMac-2:OPSLab-Test-DIR marcoautili$
```

In PART 3, we will see another type of redirection,
called **Here Strings <<<**

Redirecting STDERR



OPSLab-Test-DIR — -bash — 88x21

```
ls: minnie.txt: No such file or directory  
[iMac-2:OPSLab-Test-DIR marcoautili$  
[iMac-2:OPSLab-Test-DIR marcoautili$  
[iMac-2:OPSLab-Test-DIR marcoautili$  
[iMac-2:OPSLab-Test-DIR marcoautili$  
[iMac-2:OPSLab-Test-DIR marcoautili$ ls -l pluto.txt minnie.txt  
ls: minnie.txt: No such file or directory  
-rw-r--r--@ 1 marcoautili staff 14 Sep 20 15:29 pluto.txt  
[iMac-2:OPSLab-Test-DIR marcoautili$  
[iMac-2:OPSLab-Test-DIR marcoautili$ ls -l pluto.txt minnie.txt 2> errors-txt  
-rw-r--r--@ 1 marcoautili staff 14 Sep 20 15:29 pluto.txt  
[iMac-2:OPSLab-Test-DIR marcoautili$  
[iMac-2:OPSLab-Test-DIR marcoautili$ cat errors-txt  
ls: minnie.txt: No such file or directory  
[iMac-2:OPSLab-Test-DIR marcoautili$  
[iMac-2:OPSLab-Test-DIR marcoautili$ ls -l pluto.txt minnie.txt > output-file.txt 2>&1  
[iMac-2:OPSLab-Test-DIR marcoautili$  
[iMac-2:OPSLab-Test-DIR marcoautili$ cat output-file.txt  
ls: minnie.txt: No such file or directory  
-rw-r--r--@ 1 marcoautili staff 14 Sep 20 15:29 pluto.txt  
[iMac-2:OPSLab-Test-DIR marcoautili$
```

STDERR, by default redirected to STDOUT

STDOUT

2> errors-txt

STDERR is redirected to a file, thus ...

... only STDOUT is displayed on the screen

> output-file.txt 2>&1

Both STDOUT and STDERR are redirected to the file output-file.txt

More on 2>&1 in next slide...

If we do not specify any number with >, then it
redirects stream 1 by default, i.e., 1>

Duplicating File Descriptors

The redirection operator

`[n]<&word`

is used to **duplicate input file descriptors**

- If word expands to one or more digits, the file descriptor denoted by n is made to be a copy of that file descriptor. If the digits in word do not specify a file descriptor open for input, a redirection error occurs.
- If word evaluates to '-', file descriptor n is closed. If n is not specified, the standard input (file descriptor 0) is used

The operator

`[n]>&word`

is used similarly to **duplicate output file descriptors**

- If n is not specified, the standard output (file descriptor 1) is used. If the digits in word do not specify a file descriptor open for output, a redirection error occurs. If word evaluates to '-', file descriptor n is closed. As a special case, if n is omitted, and word does not expand to one or more digits or '-', the standard output and standard error are redirected as described previously

Run commands in sequence

Semicolons “;”

- run each command, regardless if the preceding ones fail
- `$ com1;com2;com3`

Double-ampersands “&&”

- run the next command if the preceding command succeeded, and all the commands correctly set exit codes
- `$ com1 && com2 && com3`

```
iMac:~ marcoautili$  
iMac:~ marcoautili$ echo ciao 1 ; cat pippo ; echo ciao 2  
ciao 1  
cat: pippo: No such file or directory  
ciao 2  
iMac:~ marcoautili$  
iMac:~ marcoautili$ echo ciao 1 && cat pippo && echo ciao 2  
ciao 1  
cat: pippo: No such file or directory  
iMac:~ marcoautili$  
iMac:~ marcoautili$
```

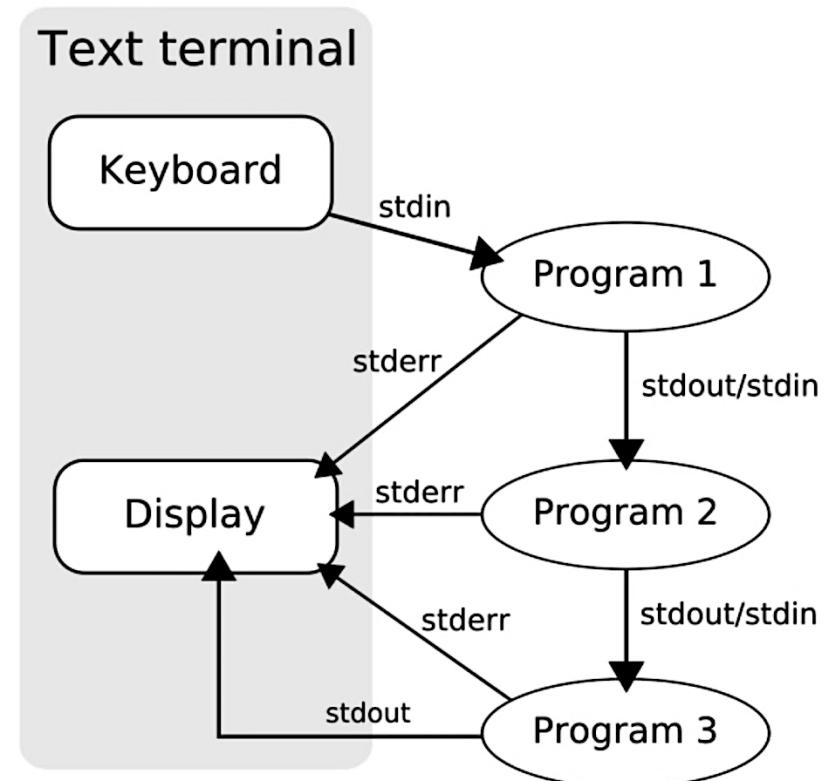
The third command is not executed

Pipeline

In Unix-like computer operating systems, a **pipeline** is a sequence of processes chained together by their standard streams, so that the output of each process (**stdout**) feeds directly as input (**stdin**) to the next one

The **exit status** of a pipeline is the exit status of the **last process** in the pipeline

The **shell** waits for the termination of **all the processes** in the pipeline before returning the prompt

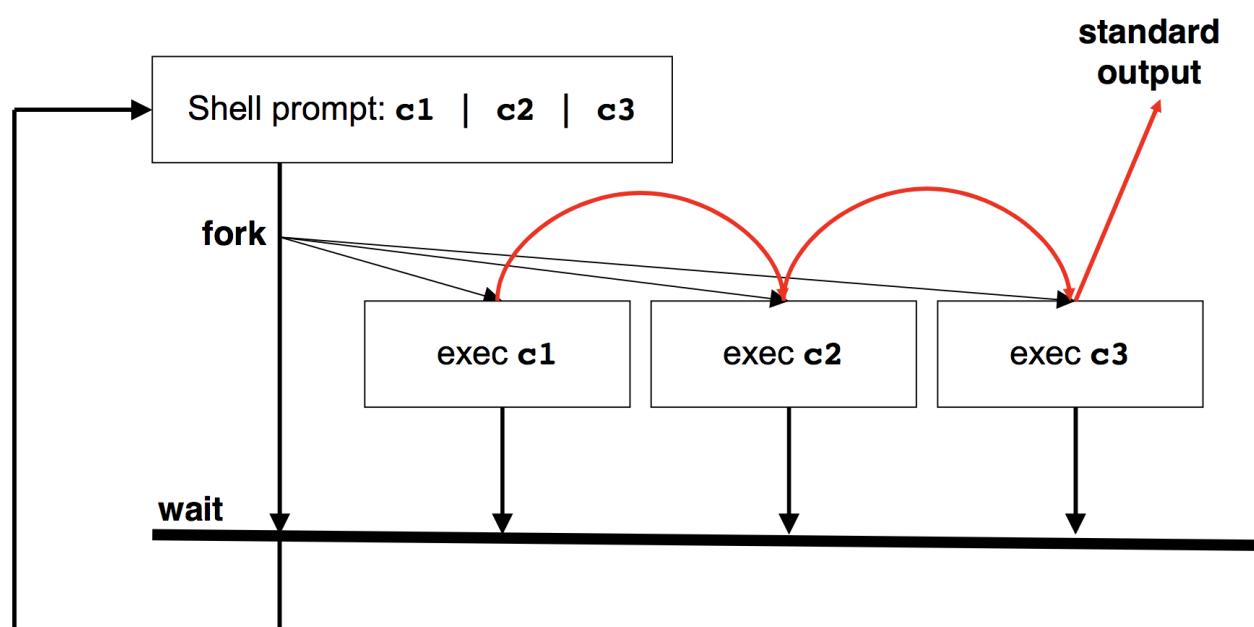


* [https://en.wikipedia.org/wiki/Pipeline_\(Unix\)](https://en.wikipedia.org/wiki/Pipeline_(Unix))

Pipeline execution order

- The order the processes are run does not matter and is not guaranteed
- The shell first creates the pipe, the conduit for the data that will flow between the processes, and then creates the processes with the ends of the pipe connected to them
- For example, the first process that is run may block waiting for input from the second process, or block waiting for the second process to start reading data from the pipe (these waits can be arbitrarily long and do not matter)

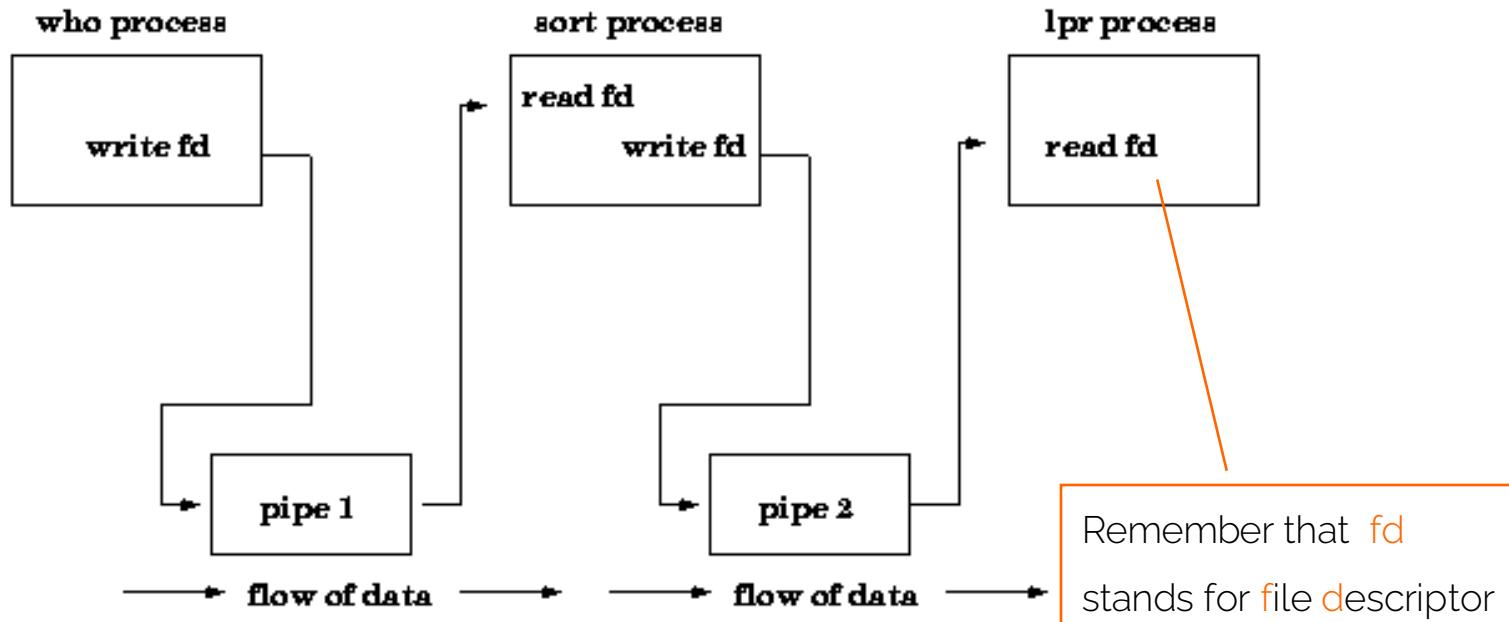
- Whichever order the processes are run, the data eventually gets transferred and “everything works”



Pipeline

Example: `who | sort | lpr`

A pipe to send to the default printer (`lpr`) the ordered (`sort`) list of all users currently logged on (`who`)



Pipeline

```
iMac-2:Filters-Test-DIR marcoautili$ cat persone.txt
Nome; Cognome; Età
# Mario; Rossi; 35
Mario; Bianchi; 36
# Paolo; Rossi; 20
Pippo; Verdi; 44
Pluto; Verdone; 21
Minnie; Rosa; 28
iMac-2:Filters-Test-DIR marcoautili$
iMac-2:Filters-Test-DIR marcoautili$
iMac-2:Filters-Test-DIR marcoautili$ cat persone.txt | egrep '^#'
# Mario; Rossi; 35
# Paolo; Rossi; 20
iMac-2:Filters-Test-DIR marcoautili$
```

List only #-commented lines

Pipeline

```
iMac-2:usr marcoautili$ pwd  
/usr  
iMac-2:usr marcoautili$  
iMac-2:usr marcoautili$ ls -la  
total 8  
drwxr-xr-x@ 12 root wheel 408 Dec 23 2015 .  
drwxr-xr-x 32 root wheel 1156 Sep 20 08:52 ..  
drwxr-xr-x 5 root wheel 170 Aug 23 2015 X11  
lrwxr-xr-x 1 root wheel 3 Oct 18 2015 X11R6 -> X11  
drwxr-xr-x 3 root wheel 102 Sep 20 2015 adic  
drwxr-xr-x 1055 root wheel 35870 Sep 5 13:00 bin  
drwxr-xr-x 267 root wheel 9078 Sep 5 13:01 lib  
drwxr-xr-x 186 root wheel 6324 Sep 5 13:01 libexec  
drwxr-xr-x 5 root wheel 170 Dec 23 2015 local  
drwxr-xr-x 243 root wheel 8262 Sep 5 13:00 sbin  
drwxr-xr-x 45 root wheel 1530 Oct 18 2015 share  
drwxr-xr-x 4 root wheel 136 Oct 18 2015 standalone  
iMac-2:usr marcoautili$  
iMac-2:usr marcoautili$ ls -la | egrep 'Sep'  
drwxr-xr-x 32 root wheel 1156 Sep 20 08:52 ..  
drwxr-xr-x 3 root wheel 102 Sep 20 2015 adic  
drwxr-xr-x 1055 root wheel 35870 Sep 5 13:00 bin  
drwxr-xr-x 267 root wheel 9078 Sep 5 13:01 lib  
drwxr-xr-x 186 root wheel 6324 Sep 5 13:01 libexec  
drwxr-xr-x 243 root wheel 8262 Sep 5 13:00 sbin  
iMac-2:usr marcoautili$  
iMac-2:usr marcoautili$ ls -la | egrep 'Sep' | sort  
drwxr-xr-x 3 root wheel 102 Sep 20 2015 adic  
drwxr-xr-x 32 root wheel 1156 Sep 20 08:52 ..  
drwxr-xr-x 186 root wheel 6324 Sep 5 13:01 libexec  
drwxr-xr-x 243 root wheel 8262 Sep 5 13:00 sbin  
drwxr-xr-x 267 root wheel 9078 Sep 5 13:01 lib  
drwxr-xr-x 1055 root wheel 35870 Sep 5 13:00 bin  
iMac-2:usr marcoautili$  
iMac-2:usr marcoautili$
```

List the directories and files
"modified in September"...

Note that it is just an example... it is not always correct. Why?

List the directories and files
"modified in September", then sort the output lines

Note that, it sorts alphabetically the whole line... and not by date!

Redirecting VS Pipelining

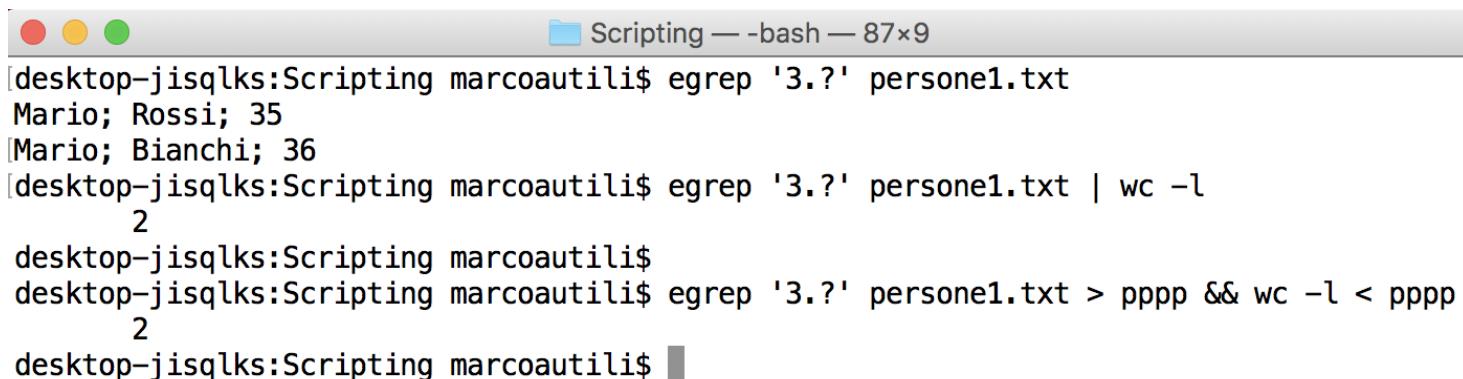
- Roughly speaking, both "do the same basic thing"
 - they redirect the file descriptors processes are connected to; the difference lies in the how
 - A pipe connects the stdout of one process to the stdin of another
 - A redirection redirects from/to a file (> from stdout to a file, < from a file to stdin)

program > file

- You are running program
- The output of program will be placed in file

program1 | program2

- You may write program1 > temp_file && program2 < temp_file
 - Run program1 and save the output into temp_file
 - Run program2 and pass the contents of temp_file to it

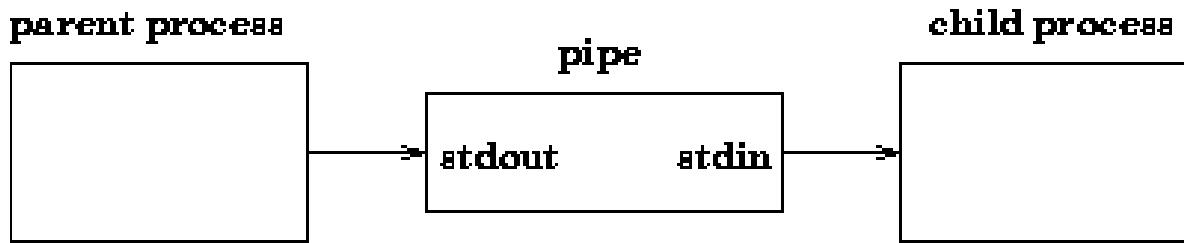


```
[desktop-jisqlks:Scripting marcoautili$ egrep '3.?' persone1.txt
 Mario; Rossi; 35
 Mario; Bianchi; 36
[desktop-jisqlks:Scripting marcoautili$ egrep '3.?' persone1.txt | wc -l
      2
desktop-jisqlks:Scripting marcoautili$
desktop-jisqlks:Scripting marcoautili$ egrep '3.?' persone1.txt > pppp && wc -l < pppp
      2
desktop-jisqlks:Scripting marcoautili$
```

Summing up

When a pipe is used in a Unix-like command line

- the first process is assumed to be writing to stdout
- the second is assumed to be reading from stdin



```
[desktop-jisqlks:~ marcoautili$  
[desktop-jisqlks:~ marcoautili$ ps aux | grep Sublime | grep -v grep | awk '{print $2}'  
995  
[desktop-jisqlks:~ marcoautili$  
[desktop-jisqlks:~ marcoautili$  
[desktop-jisqlks:~ marcoautili$  
[desktop-jisqlks:~ marcoautili$  
[desktop-jisqlks:~ marcoautili$ ps aux | grep Sublime | grep -v grep | awk '{print $2}'  
desktop-jisqlks:~ marcoautili$
```

Here, "in between", I closed my text editor Sublime or killed it... [more on that later](#)

Note that depending on what version of UNIX-like system you're using, certain programs may behave a little differently (for example, ps might output the PID in column \$3)... Again, it is just an example ☺

Activities

- Play with all you have learned so far
- For instance, create files similar to, yet more complex of, persone.txt and combine wildcards, filters, regular expressions, pipelines to perform increasingly complex operations/transformations on it

IMPORTANT

the practical examples I have provided you with so far are NOT enough to pass the exam.

You need to exercise a lot... yourself



This applies to the coming arguments as well



Programs and Processes (from PART I)

Program

- an executable file residing on disk in a directory
 - A program is read into memory and is executed by the kernel as a result of one of the seven exec functions

Processes and Process ID

- An executing instance of a program is called a **process**,
 - Sometimes the term **task** is used to refer to a program that is being executed 😊
- The UNIX System guarantees that every process has a unique numeric identifier called the **process ID**
 - the process ID is **always a non-negative integer**

Process Control

- There are three primary functions for process control: **fork**, **exec**, and **waitpid**
- The **exec function has seven variants**, but we often refer to them collectively as simply the **exec function**

Multitasking operating systems

- Many processes or tasks might be running “at the same time”
- At the same time does not necessarily mean that processes are executing at exactly the same time
 - it does mean that more than one processes can be part-way through execution at the same time, and that more than one process is advancing over a given period of time
- new processes “can” interrupt already started ones before they finish, instead of waiting for them to end. As a result, a computer executes segments of multiple processes in an interleaved manner, while the processes share common processing resources such as central processing units (CPUs) and main memory
- Additional terminologies should be mentioned (and compared) here:
 - simultaneously, concurrently, parallel execution
- From the command line, the `top` program displays and updates sorted information about processes directly on the terminal (press `q` to exit)

Processes

ps [aux] (stand for process status)

- show only the processes running in your current terminal (usually, should not be too many)
- passing the argument aux a complete system view is shown

```
iMac-2:~ marcoautili$  
iMac-2:~ marcoautili$  
iMac-2:~ marcoautili$ ps  
 PID TTY      TIME CMD  
 320 ttys000  0:00.04 -bash  
iMac-2:~ marcoautili$  
iMac-2:~ marcoautili$  
iMac-2:~ marcoautili$  
iMac-2:~ marcoautili$ ps aux  
USER        PID %CPU %MEM      VSZ   RSS TT STAT STARTED      TIME COMMAND  
marcoautili    255  4.1  0.5 2958444 182976 ?? R  1:17PM 0:14.54 /Applications/Utilities/Terminal.app/Contents/MacOS/Terminal -psn_0_40970  
marcoautili    584  0.2  0.2 3781628  82448 ?? S  1:26PM 0:11.85 /System/Library/CoreServices/Dock.app/Contents/Resources/DashboardClient.app/  
Contents/MacOS/DashboardClient  
_windowserver  184  0.2  0.5 6670308 163776 ?? Ss  1:17PM 1:31.71 /System/Library/Frameworks/ApplicationServices.framework/Frameworks/CoreGraph  
ics.framework/Resources/WindowServer -daemon  
marcoautili    387  0.1  0.8 4352980 266204 ?? S  1:18PM 3:53.22 /Applications/Dropbox.app/Contents/MacOS/Dropbox  
marcoautili    771  0.0  0.1 2611364  31860 ?? S  2:18PM 0:00.46 /Applications/Microsoft PowerPoint.app/Contents/SharedSupport/Office365Servic  
eV2.app/Contents/MacOS/Office365ServiceV2  
root       766  0.0  0.0 2508128  4464 ?? Ss  2:17PM 0:00.01 /usr/sbin/ocspd  
marcoautili    761  0.0  0.1 3629928 20912 ?? Ss  2:16PM 0:00.09 /System/Library/Frameworks/WebKit.framework/Versions/A/XPCServices/com.apple.  
WebKit.WebContent.xpc/Contents/MacOS/com.apple  
marcoautili    717  0.0  0.4 4134396 135988 ?? Ss  2:08PM 0:01.12 /System/Library/StagedFrameworks/Safari/WebKit.framework/Versions/A/XPCServic  
es/com.apple.WebKit.WebContent.xpc/Contents/  
marcoautili    708  0.0  0.0 2507884  4912 ?? Ss  2:06PM 0:00.02 /System/Library/Frameworks/AudioToolbox.framework/XPCServices/com.apple.audio  
.ComponentHelper.xpc/Contents/MacOS/com.appl  
marcoautili    707  0.0  0.0 2472520  6356 ?? Ss  2:06PM 0:00.01 /System/Library/Frameworks/AudioToolbox.framework/XPCServices/com.apple.audio  
.SandboxHelper.xpc/Contents/MacOS/com.apple.  
marcoautili    706  0.0  0.0 3646884 15020 ?? Ss  2:06PM 0:00.03 /System/Library/StagedFrameworks/Safari/WebKit.framework/Versions/A/XPCServic  
es/com.apple.WebKit.Databases.xpc/Contents/M  
arcoautili    690  0.0  0.0 2542084  9380 ?? S  1:54PM 0:00.08 /System/Library/PrivateFrameworks/AOSKit.framework/Versions/A/XPCServices/com  
.apple.iCloudHelper.xpc/Contents/MacOS/com.a  
marcoautili    666  0.0  0.0 2520472 13620 ?? S  1:44PM 0:00.03 /System/Library/Frameworks/CoreServices.framework/Frameworks/Metadata.framewo  
rk/Versions/A/Support/mdworker -s mdworker -  
marcoautili    619  0.0  0.1 2614612 30920 ?? Ss  1:31PM 0:01.37 /System/Library/Frameworks/Foundation.framework/Versions/C/XPCServices/Lookup
```

Kill command

`kill [signal] <PIIDs>`

- The `kill` utility sends the signal `[signal]` to the processes specified by the `<PIIDs>` operands
- A user can kill her own processes
- Only the super-user may send signals to other users' processes
- The PID `-1` have special meanings:
 - if superuser, broadcast the signal to all processes
 - otherwise, broadcast to all processes belonging to the user

EXIT STATUS

- The `kill` utility exits `0` on success, and `> 0` if an error occurs
- Some shells may provide a builtin `kill` command which is similar or identical to this utility
(see previous slides on builtins)

Killing a process

```
iMac-2:~ marcoautili$  
iMac-2:~ marcoautili$ ps aux | egrep 'Sublime'  
marcoautili 1102 0.0 0.0 2472728 880 s000 S+ 6:06PM 0:00.00 egrep Sublime  
marcoautili 1092 0.0 0.5 2900864 162476 ?? S 6:05PM 0:02.40 /Applications/Sublime Text 2.app/Contents/MacOS/Sublime Text 2  
iMac-2:~ marcoautili$  
iMac-2:~ marcoautili$ kill 1092  
iMac-2:~ marcoautili$  
iMac-2:~ marcoautili$ ps aux | egrep 'Sublime'  
marcoautili 1121 0.0 0.0 2447128 848 s000 S+ 6:08PM 0:00.00 egrep Sublime  
iMac-2:~ marcoautili$
```

- It can happen that, by simply sending the signal "-15", the process does not terminate
- In this case, sending the signal "-9" might be of help: it forces the termination
- Pay attention when killing processes... You should know what you are doing ☺

If not specified, the signal "15" (**SIGTERM**) is send by default to the process,
i.e., it is equivalent to

kill -15 1092
kill -TERM 1092
kill -s TERM 1092

Some of the more commonly used signals

- 1 HUP (hang up) (see <https://en.wikipedia.org/wiki/SIGHUP>)
- 2 INT (interrupt)
- 3 QUIT (quit)
- 6 ABRT (abort)
- 9 KILL (non-catchable, non-ignorable kill)
- 14 ALRM (alarm clock)
- 15 TERM (software termination signal)

Killing a process

Killing multiple processes

- `kill -9 <PID1> <PID2> ... <PIDn>`

Getting the name of a signal

- `kill -l <signal number>`

Sending signal by name

- `kill -s <signal name> <PID>`

```
iMac-di-User:~ marcoautili$ kill -l
 1) SIGHUP      2) SIGINT      3) SIGQUIT      4) SIGILL
 5) SIGTRAP     6) SIGABRT     7) SIGEMT      8) SIGFPE
 9) SIGKILL     10) SIGBUS     11) SIGSEGV     12) SIGSYS
13) SIGPIPE     14) SIGALRM     15) SIGTERM     16) SIGURG
17) SIGSTOP     18) SIGTSTP     19) SIGCONT     20) SIGCHLD
21) SIGTTIN     22) SIGTTOU     23) SIGIO      24) SIGXCPU
25) SIGXFSZ     26) SIGVTALRM   27) SIGPROF     28) SIGWINCH
29) SIGINFO     30) SIGUSR1    31) SIGUSR2

iMac-di-User:~ marcoautili$
iMac-di-User:~ marcoautili$
iMac-di-User:~ marcoautili$ kill -l 1 4 9
HUP
ILL
KILL
iMac-di-User:~ marcoautili$
```

See also [man killall](#)
(kill processes by name)

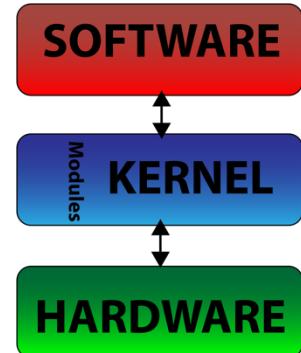
```
iMac-di-User:~ marcoautili$ ps aux | egrep 'Sublime'
marcoautili      50818  0.0  0.6 6120312 212024 ??  S    12:23PM  1:00.67 /Applications/Sublime Text 2.app/Content
marcoautili      51046  0.0  0.0 4267980   652 s000  R+  12:49PM  0:00.00 egrep Sublime
iMac-di-User:~ marcoautili$ kill -s KILL 50818
iMac-di-User:~ marcoautili$
iMac-di-User:~ marcoautili$ ps aux | egrep 'Sublime'
marcoautili      51052  0.0  0.0 4399052    836 s000  S+  12:49PM  0:00.00 egrep Sublime
iMac-di-User:~ marcoautili$
iMac-di-User:~ marcoautili$
```

Back to the pipelining example

```
marcoautili — bash — 79x23
bash-3.2$ ps aux | grep Sublime
marcoautili      79190  0.0  0.0  4399420      760 s000  S+   10:23AM  0:00.00
grep Sublime
marcoautili      79186  0.0  0.5  6077024 156372    ??  S   10:23AM  0:02.07
/Applications/Sublime Text 2.app/Contents/MacOS/Sublime Text 2
bash-3.2$
bash-3.2$
bash-3.2$ ps aux | grep Sublime | grep -v grep
marcoautili      79186  0.0  0.5  6076500 156364    ??  S   10:23AM  0:02.07
/Applications/Sublime Text 2.app/Contents/MacOS/Sublime Text 2
bash-3.2$
bash-3.2$
bash-3.2$ ps aux | grep Sublime | grep -v grep | awk '{print $2}'
79186
bash-3.2$
bash-3.2$
bash-3.2$ ps aux | grep Sublime | grep -v grep | awk '{print $2}' | xargs kill
bash-3.2$
bash-3.2$
```

Signals (more on)

- Signals are a limited form of **inter-process communication (IPC)**, typically used in Unix, Unix-like, and other POSIX-compliant operating systems
- A signal is an **asynchronous notification** sent to a process (or to a specific thread within the same process) in order to notify it of **an event that occurred**
- When a signal is sent, the operating system interrupts the target process' normal flow of execution to deliver the signal. Execution **can be interrupted during any non-atomic instruction**
 - If the process has previously **registered a signal handler**, that routine is executed. Otherwise, the **default signal handler** is executed
- Signals are **similar to interrupts**, the difference being that
 - interrupts are **mediated by the CPU** and **handled by the kernel**
 - signals are **mediated by the kernel** (possibly via system calls) and **handled by processes**



[https://en.wikipedia.org/wiki/Signal_\(IPC\)](https://en.wikipedia.org/wiki/Signal_(IPC))
<https://en.wikipedia.org/wiki/Interrupt>

Signals (more on)

- The kernel may pass an “interrupt” as a “signal” to the process that caused it (typical examples are SIGSEGV, SIGBUS, SIGILL and SIGFPE)
- A process's execution may result in the generation of a hardware exception, for instance, if the process attempts to divide by zero or incurs a page fault
 - In fact, if a process attempts to divide an integer by zero, a Floating-Point Exception is generated, and the kernel sends the SIGFPE signal to the process
 - Also, if process makes an invalid virtual memory reference, or segmentation fault, i.e., it performs a SEGmentation Violation, the kernel sends the SIGSEGV signal to the process
- The exact mapping between signal names and exceptions is obviously dependent upon the CPU, since exception types differ between architectures

[https://en.wikipedia.org/wiki/Signal_\(IPC\)](https://en.wikipedia.org/wiki/Signal_(IPC))

Additional material

- Tutorial: Master The Command Line

<https://ryanstutorials.net/linuxtutorial/>

- Tutorial: The Unix Shell

<https://swcarpentry.github.io/shell-novice/>

- Command Line Cheatsheet

<http://cht.sh>

- Explain Shell

<https://explainshell.com>

References

- Reference textbook, [APUE3ed] Advanced Programming in the UNIX Environment, 3rd Ed. Richard Stevens, Steven Rago
- <http://ryanstutorials.net/linuxtutorial/>
- <http://faculty.cs.niu.edu/~freedman/>
- <https://www.scribd.com/document/387330980/amministrare-gnu-linux-v4-0-web-cover-bis-pdf>